

The AU-PAC Model: A Semi-Structural Macroeconomic Model for Australia

David Stephan

Abstract

This paper presents AU-PAC, a semi-structural macroeconomic model for Australia adapting the FR-BDF framework of Lemoine et al. (2019, BdF WP #736), with supply-block calibration following the Dubois et al. (2026, BdF WP #1044) update. Following the FRB/US modelling tradition, the model combines Polynomial Adjustment Costs (PAC) with explicit expectations and a supply block based on CES technology, with substitution elasticity $\sigma = 0.54$ estimated from the labour FOC under a two-break trend-efficiency specification. AU-PAC contains 180 endogenous variables, 53 exogenous shocks, and approximately 351 parameters, with five PAC behavioural equations governing value-added prices, consumption, business investment, household investment, and employment, and a FR-BDF error-correction trade block with long-run income and real-exchange-rate elasticities. Expectations are formed as closed-form policy functions obtained from inversion of the estimated core Expectation Satellite VAR (E-SAT), supplemented by per-block auxiliary regressions following wp1044 §3.2.3. The model is estimated on Australian quarterly data from 1994Q1 to 2024Q4. **Phase L2 wp1044-faithful partial replication** (described in L2_REPLICATION_REPORT.md) rebuilt all five PAC blocks under the full wp1044 functional form — PAC expectation imposed at coef = +1, block-specific auxiliary VARs, derived growth-neutrality terms, exact χ from depth-m characteristic polynomial, COVID/period dummies — using an iterative-OLS pipeline against the per-block `data/pac_blocks/results_*.mat` artefacts. Four of the five blocks fit the wp1044 structural form on AU data with R^2 in $[0.41, 0.81]$ and structurally interpretable coefficients; the **consumption block's $\beta_0 = 0.23$ is the closest of any block to wp1044's $\beta_0 = 0.29$ — the one block where AU's ECM speed is comparable to France's rather than several times faster** (re-estimated 2026-05-31 on the deterministically-reproducible SA-corrected data layer, §6.14; was 0.27 on the pre-SA layer). The fifth block, business investment, structurally rejects the wp1044 PAC restriction: across 11 specification variants explored in Phase L2 P1c (documented in PAC_BI_AU_EXPLORATION.md), the free-estimated $PV(\Delta\hat{q})$ coefficient comes out at -5.0 instead of the structural $+1$, and every strict-PAC variant gives $R^2 < 0$ on raw `dln_ib`. The paper adopts a **hybrid calibration (Option 1)**: business-investment deep parameters ($\beta_0=0.096$, $\beta_1=0.33$, $\beta_2=0.11$, $\beta_3=0.69$, $\omega=0.35$) are imported from wp1044 Table 3.5.13, while the other four blocks carry their AU L2 estimates. The wage Phillips curve has standard New Keynesian structure under genuine ABS Wage Price Index data: under the BK-constrained estimate (§6.7/§6.14) $\lambda_w = 0.09$ own-lag persis-

tence, $\gamma_w = 0.86$ CPI passthrough, and a positive unemployment-gap slope $\kappa_w = 0.34$. A **Phase L2 OLS audit** (2026-05-28, §6.6) replaced remaining wp1044 (France) calibrations in the CPI Phillips, trade-quantity, and deflator blocks with AU single-equation OLS point estimates from ABS 5206 + RBA I02/F11 data — 25 of 35 audited coefficients are now AU point estimates; the wage Phillips was re-estimated under a BK-stability constraint $\lambda_w + \gamma_w \leq \mathbf{0.95}$ that delivers a substantial Phillips slope $\kappa_w = 0.343$ with near-unit CPI passthrough $\gamma_w = 0.863$ (§6.7, re-estimated on SA unemployment in §6.14). A follow-up **Phase L2.A architectural audit** (§6.8) identified that three model state variables (`yhat_au`, `pi_au`, `u_gap`) were being defined by E-SAT-style reduced-form equations instead of the FR-BDF structural identities (CES potential + demand-side aggregation for `yhat_au`; food/energy/core deflator aggregation for `pi_au`; labour-market accounting for `u_gap`). Per wp736 §2.3, §4.3.3, §4.5.1 and §4.7, the E-SAT equations are forecasting tools for agents, not defining equations for the simulator’s actual state variables. Refactoring to the structural identities yields a 30–70% amplification in real-activity IRFs and unlocks the now-identified wage Phillips slope. The CPI Phillips audit confirms the **flat AU Phillips** stylised fact: only the lag coefficient is identified ($\lambda_\pi = 0.259$, $t = 2.76$, $R^2 = 0.063$); output-gap, VA-price, and import-price channels are all insignificant. Under the post-L2.A calibration, with the trade and labour blocks corrected to use ABS *seasonally-adjusted* series (§6.13–§6.14 — the earlier use of the smoothed *Trend* export/import series had produced a spurious ~11Q IRF oscillation), the response to a 100 bp annualized cash-rate tightening produces a real-GDP trough of -0.144% at Q11 (on the current calibration, after the 2026-05-31 employment- χ / housing / consumption-SA ECM-speed writebacks amplified the real-activity channels; was -0.098% under the prior calibration with the long-run-neutrality fix, and -0.073% before that under supply-side hysteresis), an output-gap trough of -0.086% at Q11, headline CPI inflation falling only -0.004 pp quarterly (≈ -0.015 pp year-ended) at Q3 (food up, energy down, core flat — matching the flat-Phillips empirical reality), wage inflation falling -0.008 pp at Q3 (now structurally responsive via κ_w), the AUD appreciating 0.99% at Q8 via a forward-looking UIP NPV channel, and the 10-year yield rising 0.105 pp on impact; the IRFs are clean hump shapes (1–2 turning points, no oscillation) over the Q40 reporting window. A cross-block finding emerges from the Phase L2 estimates: **AU ECM speeds are 4–8× faster than France’s** in four of five blocks (consumption is the exception); plausible drivers are AU’s more flexible price/wage setting (mining cycle, commodity-pass-through, variable-rate mortgages). Benchmarking against the RBA’s model suite (Beckers 2020 VAR, MARTIN, DINGO, Murphy — Mulqueeney, Ballantyne and Hambur 2025) places AU-PAC at the conservative end of the AU range, with magnitudes that reflect the greater capital–labour substitutability of the CES specification ($\sigma = 0.54$) and the small estimated PAC adjustment-cost coefficients. The forward-guidance puzzle is absent: AU-PAC tracks the linear reference response to within 16% even at $N = 12$ successive rate-cut announcements, while standard New Keynesian models saturate near $1.8\times$. Australia-specific features include an endogenous Taylor rule for the RBA,

variable-rate mortgage transmission, commodity-price channels, the US as the foreign bloc, and a trade ECM calibrated to Australia’s secular openness rise. The model is implemented in Dynare 7.0-arm64 (MATLAB R2026a); production parameters are estimated by equation-by-equation OLS following wp1044 §2.2.

Keywords: Semi-structural model, polynomial adjustment costs, expectations, monetary policy transmission, Australian economy

JEL codes: C51, C54, E17, E37, E52

Errata / current-state note (2026-05-30). This paper was drafted across the model’s MCMC era and the subsequent equation-by-equation-OLS pivot, and some passages have not been fully reconciled. For the avoidance of doubt: - **Estimation.** Production parameters come from **equation-by-equation OLS / calibration** (wp1044 §2.2). The full-system Bayesian MCMC pipeline (`au_pac_bayesian.mod`) and the `aggregate()/cherrypick()/pac.print()` source tree (`dynare/simulation/, phaseW_recherrypick.m`) were **removed** in cleanup 7995ce7; `au_pac.mod` is now self-contained. Any Laplace/MHM/LMD figure (esp. §6.2 E-SAT MHM and §6.4 Table 5.6) is **historical and not reproducible** from the current tree. Companion files referenced but since deleted — `au_pac_var.mod`, `au_pac_mce.mod`, `PAC_BI_AU_EXPLORATION.md`, `mcmc_posterior_table.md` — should be read as historical. - **Dimensions.** 180 endogenous variables, 53 exogenous shocks, 351 parameters (Table 1.1’s 175/49/288 are superseded). - **Toolchain.** Dynare **7.0-arm64** + MATLAB **R2026a** (not 6.5 / R2020a). Line references into `au_pac.mod` (e.g. the BI calibration “lines 1746–1771”) refer to an earlier revision; the BI block now sits near lines 2073–2099. - **Wave 1 corrections (done 2026-05-30).** (i) After the §6.13 SA-data fix, the `h_pac_n_*/h_pac_ih_*` policy-function vectors were re-derived via `pac.print()` and found **bit-identical** — the PAC h-vector depends only on the auxiliary VAR + discount factor, not on the ECM speeds the SA fix changed, so no inconsistency existed. (ii) A new switch `lambda_hyst = 0` restores **long-run money-neutrality**: it gates the transitory-gap-into-trend channels (`σ_ces_rw_gap` in `dln_n_star_bar`, which feeds potential output `ln_QN/ln_Q`; and the mortgage/house-price/output-gap terms in the reporting-only `dln_ih_star_bar/dln_ib_star_bar`). A temporary 100 bp tightening now leaves `ln_Q`, `ln_N`, reported `ln_IH` and `ln_IB` reverting to ≈ 0 by Q200 (previously `ln_Q` +0.054%, `ln_N` +0.093%, `ln_IH` $\approx -15\%$, all wrong-signed/permanent). BK preserved; the real-GDP trough deepens to -0.098% at Q12 under that fix alone (and further to -0.144% at Q11 on the current calibration, after the

2026-05-31 employment- χ / housing / consumption-SA ECM-speed writebacks; set `lambda_hyst = 1` to recover the prior supply-side hysteresis).

Table of Contents

1. Introduction
 2. Bird's-Eye View of the Model
 3. Expectation Formation, PAC Framework, and Phase L2 Methodology
 4. Model Specification
 5. Cross-block findings
 6. Estimation — incl. **§6.6 Phase L2 OLS audit + §6.7 BK-constrained wage Phillips + §6.8 Phase L2.A architectural fixes (2026-05-28)**
 7. Model Properties
 8. Conclusion
 9. Australia-Specific Features
 - Appendix A: Complete Variable List
 - Appendix B: Complete Shock List
 - Appendix C: Growth Neutrality Proofs
 - Appendix D: h-Vector Decomposition
 - Appendix E: Complete Parameter List
 - Appendix F: Identification Analysis
 - Appendix G: Business Investment Exploration (the wp1044 PAC rejection on AU data)
 - References
-

1. Introduction

AU-PAC adapts the FR-BDF semi-structural model of Lemoine et al. (2019, Banque de France Working Paper #736) — updated in Dubois et al. (2026, WP #1044) — to Australian data. The framework belongs to the FRB/US family of semi-structural models used at central banks worldwide, including MUSE at the Bank of Canada, LENS at the Bank of England, and ECB-Base at the European Central Bank. These models occupy a middle ground between reduced-form VARs (which lack structural interpretation) and full DSGE models (which impose strong theoretical restrictions): behavioural equations are derived from optimisation with adjustment costs, while the supply block and financial channels are specified semi-structurally.

Three features are central to the modeling approach:

1. **Explicit expectations.** Agents form expectations about future economic

conditions using either a backward-looking satellite VAR model (E-SAT) or model-consistent forward-looking expectations. The choice of expectation regime has first-order effects on monetary policy transmission.

2. **Polynomial adjustment costs.** Non-financial behavioral equations are derived from agents minimizing costs of deviating from long-run targets subject to polynomial adjustment costs, yielding error-correction equations augmented with the present value of expected future target changes.
3. **Well-defined supply block.** A CES production function with labor-augmenting technical progress determines long-run output and provides consistent targets for employment, investment, and value-added prices through factor demand conditions.

AU-PAC modifies the FR-BDF framework along four substantive dimensions to reflect Australian conditions: (i) commodity-price channels into export volumes, the export deflator, and the import deflator; (ii) an endogenous Taylor rule for the RBA in place of an exogenous short rate; (iii) variable-rate mortgage transmission; and (iv) the US as the foreign bloc in place of the euro area.

Phase L2 wp1044-faithful partial replication — what this paper documents

This version of the paper documents the model’s state at the end of **Phase L2 P1c** (2026-05-26): five PAC blocks rebuilt under the full wp1044 functional form (Eqs 16, 30, 35, 37, 46), with the iterative-OLS pipeline detailed in §3.5 and the block-by-block results catalogued in §4.3–§4.6, summarised cross-block in §5, and used as Bayesian priors in §6. The methodological gap that motivated this rebuild is catalogued in `PAC_EQUATIONS_AUDIT.md`: the earlier `estimate_all_pac_iterative.m` shortcut used a single 9-variable VAR shared across all five blocks, HP-gap proxies for the wp1044 structural targets, and an approximate $\chi = \Sigma\beta + \omega$ formula. The Phase L2 rebuild closes those gaps with five block-specific VARs (6–9 state variables each), block-specific structural target constructions (`p*_Q` – `p_Q` for VA-price; `n*_S` – `n_S` for employment; `c* – c` from permanent-income for consumption; `I*_H/I_H` with price-spread regressor for housing inv; `I*_B/I_B` with synthetic-demand-gap for business inv), and the exact χ from the depth-m characteristic polynomial.

Four of five blocks fit the wp1044 structural form on AU data (R^2 in [0.41, 0.81]), with the **consumption block’s $\beta_0 = 0.23$ the closest to wp1044’s 0.29 of any block — the only one where AU’s ECM speed is comparable to France’s rather than 4–8× faster** (§6.14 re-estimate on the SA-corrected layer; 0.27 pre-SA). The exception is business investment: across 11 specification variants tested in Phase L2 P1c (baseline strict PAC; AU dummies; pre-residualised dummies; PV regressors free; combined PV at coef=+1; terms-of-trade + piecewise trends; commodity-augmented target; wp736 2019 simpler form; depth-2 PAC with simplified regressors; etc. — full table in Appendix G), every strict-PAC variant produces $R^2 < 0$ on raw `dln_ib`, and when

PV terms are free-estimated the PV($\Delta\hat{q}$) coefficient comes out at -5.03 instead of the structural $+1$. **AU business investment data structurally rejects the wp1044 PAC restriction.**

The hybrid calibration in one paragraph

Following the path forward selected in PAC_BI_AU_EXPLORATION.md §7 (Option 1), the production model `dynare/au_pac.mod` imports the business-investment deep parameters from wp1044 Table 3.5.13 ($\beta_0=0.096$, $\beta_1=0.33$, $\beta_2=0.11$, $\beta_3=0.69$, $\omega=0.35$, $\sigma_{CES}=0.5366$) and carries AU L2 estimates for the other four PAC blocks. The Bayesian estimation file `dynare/au_pac_bayesian.mod` removes business-investment parameters from `estimated_params` (treats them as calibrated). The rationale is the standard small-open-economy approach when local data identifies poorly: import deep parameters from a larger / better-identified economy and let AU-specific channels (E-SAT VAR, the other four PAC blocks, AU-sized shocks, AU-specific commodity / housing / mortgage dynamics) inject the country-specific dynamics through the simulation. The hybrid calibration is locked in `dynare/au_pac.mod` lines 1746–1771 with cross-references to this paper and to PAC_BI_AU_EXPLORATION.md. §3.6 elaborates the methodological rationale; §4.6.2 documents the BI block specifically; §5.3 frames the rejection finding cross-block.

Table 1.1: Model dimensions

Dimension	AU-PAC
Endogenous variables	180 declared (≈ 195 after Dynare auxiliary-variable expansion)
Exogenous shocks	53 (incl. 2 COVID dummies)
Parameters	351
PAC equations	5 (4 AU-estimated + 1 wp1044-calibrated; see §3.6)
Trade ECM equations	2 (exports, imports — proper LR + SR specification)
var_model equations	12
Forward-looking objects (Phase T)	5 (pv_i, pv_i_uip, pv_u_gap, pv_yh, pv_r_lh_gap)
HICP reporting components (Round 1.1)	6 (food, energy, core, tradeables, non-tradeables, trimmed mean)
Round 4-8 extensions (2026-05-20)	11 new endogenous vars (i_us, ibar_us, dln_pop_bar, 3 tax gaps, 2 branch decomp, 3 nowcasters)
Estimation sample (Bayesian)	1994Q3–2024Q4 (122 quarters)

Dimension	AU-PAC
Phase L2 PAC sample (per block)	VA-price 108 obs, employment 124, consumption 102, housing inv 70, business inv 122
Phase L2 PAC blocks fitting wp1044 form	4 of 5 (VA-price, employment, consumption, housing inv); business inv structurally rejected
BI calibration source	wp1044 Table 3.5.13 (imported via Option 1, see §3.6)
Bayesian observables	9 (yhat_au, pi_au, i_au, yhat_us, pi_us, pi_w, dln_c, dln_ib, i_10y)

Software and implementation

The model is implemented in Dynare 7.0-arm64 (MATLAB R2026a). The production simulation model `au_pac.mod` is now **self-contained**: the per-PAC-block auxiliary files in `dynare/aux/` emit their closed-form policy-function PAC expectation formulas via `pac.print()`, and those `h_pac_*` coefficients are written into `au_pac.mod` directly. (The earlier `aggregate()/cherrypick()` source tree and the companion `au_pac_var.mod/au_pac_mce.mod/au_pac_bayesian.mod` files were removed in cleanup 7995ce7; see the errata note above.) Source code is available at <https://github.com/DavidASStephan/AUSPAC>. The production calibration uses **equation-by-equation OLS following the wp1044 §2.2 methodology**; per-block results live in `data/pac_blocks/results_*.mat,txt`. (Historical: a full-system Bayesian MCMC estimation was previously run as a robustness check via Dynare’s `estimation()` command with `diffuse_filter`, `csmmwel` mode search, and 20,000-draw two-chain Metropolis-Hastings; that pipeline has been retired and its LMD figures are not reproducible from the current tree.)

Paper outline

Section 2 provides a bird’s-eye view of the model structure and transmission channels. Section 3 develops the expectation formation framework, the PAC microfoundations (with emphasis on the coefficient=+1 structural restriction on PV terms in §3.3), the Phase L2 iterative-OLS estimation pipeline (§3.5), and the rationale for hybrid calibration of one block (§3.6). Section 4 specifies all model equations block by block, with coefficient tables that report Phase L2 AU iterative-OLS estimates side-by-side with wp1044 FR-BDF values for each PAC block. Section 5 — new in this version — collects the **cross-block findings** that emerge from comparing AU L2 against FR-BDF: AU ECM speeds 4–8× faster than France’s (§5.1); the wp1044 PAC framework validates for four of five blocks (§5.2); business investment’s structural rejection plus the Option

1 implementation (§5.3); implications for IRF analysis (§5.4). Section 6 describes the estimation pipeline (data, E-SAT Bayesian, PAC structural with the L2 results from §4, Bayesian full-system, pseudo-real-time recursive forecast evaluation). Section 7 analyses model properties including monetary policy transmission, impulse responses to all shocks, conditional forecasting, and the forward guidance experiment, with explicit notes where the post-hybrid IRFs differ from earlier baselines. Section 8 concludes. Section 9 discusses Australia-specific features. Seven appendices provide variable, shock, growth-neutrality, h-vector, parameter, identification, and BI-exploration tables.

2. Bird's-Eye View

2.1 Model structure

AU-PAC is organized into interconnected blocks, with expectations playing a role in most behavioral equations:




Figure 2.1: Model block diagram. Variables directly affected by expectations appear throughout: the 5 PAC equations, the term structure, exchange rate, wage Phillips curve (PV of unemployment), and the consumption target (PV of permanent income).

2.2 Expectation regimes

The model can be solved under three expectation regimes (following wp736 Section 6):

Table 2.1: Three expectation regimes

Regime	Financial expectations	Non-financial expectations	Forward vars	Status
VAR-based	Backward (AR(1) policy functions, incl. UIP NPV)	Backward (PAC h-vectors from var_model)	0	historical (au_pac_var.mod re- moved in 7995ce7)
Hybrid	Forward (pv_i, pv_i_uip, pv_u_gap, pv_yh)	Backward (PAC h-vectors from var_model)	4	production (au_pac.mod)
Full MCE	Forward (all leads)	Forward (pac_expectation expands to leads)	31	historical (au_pac_mce.mod re- moved in 7995ce7)

The production model is the single **Hybrid** regime; the VAR-based and full-MCE variants were comparison artefacts of the retired MCMC line and are no longer in the tree (their §7 figures regenerate from saved `.mat` artefacts). Under the baseline configuration, the five PAC expectations use closed-form policy-function formulas in lagged structural state variables (§4.4.0), while the term structure (pv_i , eq. 48), UIP NPV ($pv_{i,uip}$, §4.10.3 — forward NPV of the policy-rate gap), unemployment PV, and permanent-income PV all use forward-looking recursive forms. Under full model-consistent expectations (MCE), all expectations are forward-looking. The forward-NPV UIP is what produces the FR-BDF-style amplification of the exchange-rate channel documented in §7.2.

2.3 Key transmission channels

Table 2.2: Monetary policy transmission channels

Channel	Mechanism	Key equations	Key parameters
Interest rate	Short rate -> 10Y yield -> WACC -> user cost -> business investment	eqs (48)-(50), (55)	$\kappa_{10}=0.97$, $\sigma_{ces}=0.54$
Exchange rate	Short rate -> forward NPV $p_{v_i uip}$ -> UIP -> AUD appreciation -> ex- ports/imports	§4.10.3, eq (54)	$\beta_{uip}=0.92$, $\rho_s=0.78$, $\alpha_s=0.15$
Mortgage	Short rate -> 10Y yield -> bank lending -> housing investment	eqs (56), (38)	$spread_{lh}=0.40$, $\kappa_{mort}=0.048$
Expectations	Short rate enters E-SAT -> all h-vectors shift	var_model block	All a_{X_Y} auxiliary params
Wage-price	Output gap -> unemployment PV -> wages -> ULC -> VA price	eqs (30), (28), (29)	κ_w from MCMC, $\gamma_{ulc}=0.295$
Permanent income	Output gap -> PV of future income -> consumption target	eqs (33), (34)	$\beta_c=0.95$, $\kappa_{inc}=0.050$

3. Expectation Formation, PAC Framework, and Phase L2 Methodology

3.1 The E-SAT expectation satellite model

The Expectation SATellite model (E-SAT) is a structural VAR with 11 core equations. Agents with limited information form expectations by forecasting from this small model. The core equations relate the Australian output gap, inflation, and interest rate to their US counterparts:

Australian IS curve (eq. 1):

$$\hat{y}_t = \lambda_q \hat{y}_{t-1} + \delta \hat{y}_t^{US} - \sigma_q (i_{t-1} - \bar{i}_{t-1} - \pi_{t-1}^{gap}) + \lambda_{dom} \hat{y}_t^{dom} + \varepsilon_t^q$$

Deviation from FR-BDF wp736: the final term $\lambda_{dom} \hat{y}_t^{dom}$ is AU-PAC-specific and has no counterpart in FR-BDF's IS curve, which uses only $\hat{y}^{US}$ (the foreign-output spillover) and the lagged real interest gap. We include $\hat{y}^{dom}$ — an expenditure-share-weighted aggregate of contemporaneous domestic-demand growth ($w_c \dot{c} + w_{ib} \dot{i}_b + w_{ih} \dot{i}_h + w_g \dot{g} + w_x \dot{x} - w_m \dot{m}$, see §4.10 below) — to close the Keynesian multiplier loop. Without this term the structural feedback from consumption/investment shocks back to the output gap, and hence into the Taylor rule, runs only through the foreign-output channel, which empirically understates aggregate-demand transmission in an open small commodity economy where domestic absorption is the dominant gap-mover. The Bayesian posterior mean for λ_{dom} is 0.40, four times the prior centre of 0.10, indicating the data strongly supports a substantial domestic-demand feedback (§5).

Australian Phillips curve (eq. 2):

$$\pi_t^{gap} = \lambda_\pi \pi_{t-1}^{gap} + \kappa_\pi \hat{y}_{t-1} + \varepsilon_t^\pi$$

Taylor rule (eq. 3):

$$i_t^{gap} = \lambda_i i_{t-1}^{gap} + (1 - \lambda_i)(\alpha_i \pi_{t-1}^{gap} + \beta_i \hat{y}_{t-1}) + \varepsilon_t^i$$

US IS curve (eq. 4): $\hat{y}_t^{US} = \lambda_{q,US} \hat{y}_{t-1}^{US} + \varepsilon_t^{q,US}$

US Phillips curve (eq. 5): $\pi_t^{gap,US} = \lambda_{\pi,US} \pi_{t-1}^{gap,US} + \kappa_{\pi,US} \hat{y}_{t-1}^{US} + \varepsilon_t^{\pi,US}$

Plus three anchor equations (eqs. 6-8) for the long-run interest rate, AU inflation target, and US inflation target, each following AR(1) processes toward steady state.

Bayesian estimation E-SAT was estimated using Bayesian methods (Metropolis-Hastings, 50,000 draws, 2 chains) on quarterly Australian data from 1993Q1 to 2024Q4.

Table 3.1: E-SAT posterior estimates

Parameter	Description	Prior	Posterior mean
delta	AU-US demand spillover	Beta(0.20, 0.10)	0.199
lambda_q	AU output gap persistence	Beta(0.50, 0.15)	0.448
sigma_q	Real rate sensitivity	Gamma(0.20, 0.10)	0.166
lambda_i	Taylor rule inertia	Beta(0.80, 0.05)	0.828

Parameter	Description	Prior	Posterior mean
alpha_i	Taylor rule inflation weight	Normal(0.30, 0.15)	0.279
beta_i	Taylor rule output weight	Normal(0.15, 0.10)	0.135
lambda_pi	AU inflation persistence	Beta(0.30, 0.10)	0.263
kappa_pi	Phillips curve slope	Gamma(0.06, 0.03)	0.058
lambda_q_us	US output gap persistence	Beta(0.80, 0.10)	0.806
lambda_pi_us	US inflation persistence	Beta(0.65, 0.10)	0.653
kappa_pi_us	US Phillips slope	Gamma(0.02, 0.01)	0.013

Convergence confirmed by Geweke tests (all $p > 0.10$) and Brooks-Gelman diagnostics. Acceptance rates: 46.1-46.3%.

Table 3.2: Steady-state values

Variable	SS value (quarterly)	Annual equivalent	Description
i_ss	1.049%	~4.2%	RBA neutral rate
pi_ss_au	0.625%	~2.5%	RBA inflation target midpoint
pi_ss_us	0.500%	~2.0%	Fed inflation target

Companion matrix The E-SAT model can be written in structural VAR form $AZ_t = BZ_{t-1} + \varepsilon_t$, yielding the reduced form $Z_t = HZ_{t-1} + \eta_t$ where $H = A^{-1}B$. The i -step-ahead forecast is $Z_{t+i}^e = H^i Z_t$, enabling computation of discounted present values of any variable.

3.2 Enriched E-SAT: the var_model architecture

AU-PAC enlarges the E-SAT state vector with auxiliary equations for each PAC target variable. Following wp736 Tables 4.4.4, 4.5.7, 4.6.3–4, 4.6.11–12, and 4.6.16, these auxiliary equations capture how the macro state (output gap, interest-rate gap, inflation gap, unemployment gap) feeds into each target variable’s dynamics.

The enriched system is declared as a single `var_model` in Dynare:

```
var_model(model_name = esat_enriched,
    eqtags = ['var_y', 'var_i', 'var_pi',           // 3 E-SAT core
              'var_u', 'var_yus',                 // 2 additional states
              'var_pQ', 'var_n', 'var_yh', 'var_c', // 4 auxiliary gaps
              'var_ib', 'var_rKB', 'var_ih']);      // 3 auxiliary gaps
```

All 5 PAC equations share this enriched `var_model` for h-vector computation.

Table 3.3: Enriched var_model equations (12x12 companion matrix)

#	Tag	Variable	Equation type	wp736 ref
1	var_y	y_gap_var	IS curve (lagged)	E-SAT eq 1
2	var_i	i_gap_var	Taylor rule (lagged)	E-SAT eq 3
3	var_pi	pi_gap_var	Phillips curve (lagged)	E-SAT eq 2
4	var_u	u_gap_var	Okun's law	Table 4.5.2
5	var_yus	yhat_us_var	US IS (AR(1))	E-SAT eq 4
6	var_pQ	piQ_hat	VA price auxiliary	Table 4.4.4
7	var_n	n_hat	Employment auxiliary	Table 4.5.7
8	var_yh	yh_ratio_hat	Income-output ratio	Table 4.6.3
9	var_c	c_hat	Consumption PV^2	Table 4.6.4
10	var_ib	ib_hat	Business inv output gap	Table 4.6.11
11	var_rKB	rKB_hat	Business inv user cost	Table 4.6.12
12	var_ih	ih_hat	Housing inv target	Table 4.6.16

Table 3.4: Auxiliary equation coefficients

Each auxiliary equation takes the form: $\hat{X}_t = \rho_X \hat{X}_{t-1} + a_{X,y} \hat{y}_{t-1} + a_{X,i} \hat{i}_{t-1} + a_{X,\pi} \hat{\pi}_{t-1} + a_{X,u} \hat{u}_{t-1} + \dots$

Values shown are the Phase W Bayesian posterior means from `simulation/identities/calibration.inc` (Phase B re-estimation 2026-05-09; activated 2026-05-17).

Auxiliary	Own lag (rho)	y_gap	i_gap	pi_gap	u_gap	Other	wp736 ref (France)
piQ_hat	0.334	0.043	- 0.021	0.007	- 0.021	a_pQ_w (Phase U)	Table 4.4.4 (FR rho=0.70)
n_hat	0.743	0.094	- 0.031	0.057	- 0.029	—	Table 4.5.7
yh_ratio_hat	0.193	0.12	—	—	-0.07	—	Table 4.6.3
c_hat	0.581	0.058	- 0.043	0.010	- 0.036	a_c_yh	Table 4.6.4
ib_hat	0.694	0.050	—	0.023	0.004	—	Table 4.6.11
rKB_hat	0.162	—	0.242	—	—	—	Table 4.6.12
ih_hat	0.699	0.097	- 0.152	0.042	0.004	—	Table 4.6.16

Persistence parameters and auxiliary coefficients are estimated on AU observable target proxies via Bayesian shrinkage; the wp736 references provide the cross-equation specification.

3.3 PAC microfoundations

The PAC framework derives behavioral equations from agents minimizing a cost function that penalizes deviations from a target y_t^* and m differences of the decision variable y_t :

$$C_t = \sum_{i=0}^{\infty} \beta^i \left[(y_{t+i} - y_{t+i}^*)^2 + \sum_{k=1}^m b_k ((1-L)^k y_{t+i})^2 \right]$$

The first-order condition yields an error-correction equation (eq. 9):

$$\Delta y_t = a_0(y_{t-1}^* - y_{t-1}) + \sum_{k=1}^{m-1} a_k \Delta y_{t-k} + \sum_{i=0}^{\infty} d_i \Delta y_{t+i}^*$$

where a_0 is the error-correction speed, a_k are AR lag coefficients, and $\omega = \sum d_i$ is the share of nonstationary expectations. **Growth neutrality** requires $1 - \sum a_k - \omega = 0$ on the balanced growth path.

The forward-sum coefficients d_i have a closed-form expression $d_i = \frac{c_y}{c_y + c_\Delta} \chi^i$ where χ is the smallest stable root of the depth- m characteristic polynomial $\lambda^{m+1} - (1 + \sum a_k + \omega)\lambda^m + (\sum a_k)\lambda^{m-1} - \dots = 0$. Re-writing the forward sum as a PV operator on the state via the VAR companion matrix Φ , this normalises so that **the coefficient on PV(Δy^*) in the estimated regression is structurally = +1** (Lemoine et al. 2019, §3, Eqs 7–8; reproduced as wp1044 §3.2). This coefficient-of-+1 restriction is the testable structural restriction of the PAC framework: if the estimated coefficient diverges from +1 (or, worse, takes the wrong sign), the underlying agent optimisation is rejected and the equation becomes reduced-form.

For an IRF / counterfactual / policy-analysis model (which AUSAAC is), the coefficient-of-+1 restriction matters for three reasons: (i) the IRFs through the expectation channel are mechanically wrong if coef $\neq$ +1; (ii) the Lucas critique then applies to the equation (the estimated relationship becomes policy-regime dependent); (iii) there is no balanced-growth-path consistency in steady state (the variable doesn't track its target's long-run growth rate). The Phase L2 partial replication tested this restriction explicitly on AU data for all 5 PAC blocks. **Four of five blocks pass** (VA-price, employment, consumption, housing inv — see §4.3–§4.6.1 for per-block coefficient comparisons against wp1044 Tables 3.3.3 / 3.4.9 / 3.5.2 / 3.5.7). **One block fails on every variant tested** (business investment — see §4.6.2 + §5.3 + Appendix G); for that block the paper adopts the **Option 1 hybrid calibration** documented in §3.6.

Under VAR-based expectations, the infinite sum collapses to $k_0 Z_{t-1}$ where k_0 is the h-vector computed from the var_model companion matrix. Dynare's `pac_expectation()` command automates this computation in the production

model. The Phase L2 iterative-OLS pipeline (§3.5) implements the same PV operator in MATLAB for stand-alone block-by-block estimation.

Table 3.5: PAC equation specifications

Equation	Order m	AR lags	Discount beta	h-vector sum	Manual omega	Amplification
VA price	1	1	0.98	0.452	0.46	1.0x
Consumption	1	1	0.98	0.678	0.369	1.84x
Business inv	2	2	0.98	0.501	0.350	1.43x
Household inv	2	2	0.98	0.569	0.300	1.90x
Employment	4	4	0.98	0.446	0.300	1.49x

The 1.4–1.9× amplification across the consumption, business-investment, household-investment, and employment equations is consistent with the wp736 Section 6 finding that forward expectations amplify monetary transmission.

3.4 Expectation architecture (Phase T policy-function workflow)

The current production model uses a **single-regime, closed-form policy-function** workflow following FR-BDF wp1044 §3.2.3 (Dubois et al. 2026) and the Adjemian/Brayton/Zimic semi-structural-Dynare pattern. For each PAC block, Dynare’s `pac.print()` is run on a per-block aux file to obtain a closed-form expansion of `pac_expectation()` as a linear combination of the model state variables (the `h_pac_*` coefficient set). `cherrypick()` extracts simulation-ready equations and `aggregate()` composes them with the structural identities into the production `au_pac.mod`. Forward-looking objects pv_i , $pv_{u,gap}$, pv_{yh} , $pv_{r,lh,gap}$, $pv_{i,uip}$ enter as recursive constraints in the structural identity layer rather than as part of the PAC expansion. This replaces a previous three-regime workflow (VAR / Hybrid / MCE) that ran the same model under alternative `pac_expectation()` settings; the historical comparison artefacts retained for paper §7.2 figures (`saved_irfs_{var,hybrid,mce}.mat`) come from that earlier workflow and are kept as a documentation of the architectural evolution.

3.5 Iterative OLS estimation pipeline (Phase L2)

The Phase L2 partial replication (described in `L2_REPLICATION_REPORT.md`) implements a stand-alone, block-by-block iterative-OLS estimator of the wp1044 PAC equations, used both as a methodology audit of the Dynare PAC machinery and as a source of structurally interpretable AU coefficients that feed into the per-block coefficient tables in §4. The earlier `data/estimate_all_pac_iterative.m` shortcut shared a single 9-variable

VAR across all five blocks and used HP-gap proxies for the wp1044 structural targets; the gap catalogue is in `PAC_EQUATIONS_AUDIT.md`. The Phase L2 pipeline closes those gaps.

Six-step pipeline (`data/estimate_pac_<block>.m` for $\langle \text{block} \rangle \in \{\text{va_price}, \text{employment}, \text{consumption}, \text{housing_inv}, \text{business_inv}\}$):

1. **Block-specific auxiliary VAR.** For each PAC block, build a block-specific VAR(1) on the wp1044 Tables 3.3.4 / 3.4.10 / 3.5.3 / 3.5.8 / 3.5.14 state vectors (6–9 variables per block; see Table 3.6 below). Estimate Φ by lag-1 OLS on the AU sample (`data/pac_helpers/build_block_var.m`).
2. **Initial coefficient seed.** Seed $\beta_0, \beta_1, \dots, \omega$ at the wp1044 reference values from Tables 3.3.3 / 3.4.9 / 3.5.2 / 3.5.7 / 3.5.13. Calibrate ω to wp1044 (not free-estimated). Calibrate $\sigma_{\text{CES}} = 0.5366$ from the AU labour-FOC supply block (§4.2.3).
3. **Solve for χ from the exact characteristic polynomial.** Given the current $(\beta_1, \dots, \beta_{\text{depth}}, \omega)$, solve the depth- m polynomial $\lambda^{m+1} - (1 + \sum_k \beta_k + \omega)\lambda^m + (\sum_k \beta_k)\lambda^{m-1} - \dots = 0$ for the smallest stable root χ (`data/pac_helpers/solve_pac_chi_exact.m`). For depth 1 this matches the simplification $\chi = \Sigma\beta + \omega$ used in the pre-Phase-L2 builder; for depth ≥ 2 the exact solver is needed (employment depth 3, business inv depth 2).
4. **Compute PV operator on the VAR state.** Form $PV_z = (I - \chi\Phi)^{-1}\chi\Phi$ (`data/pac_helpers/compute_pv_term.m`). For each PV regressor in the wp1044 equation, the corresponding row $e^{X'}PV_z$ gives the closed-form linear combination of lagged states that the agent’s discounted-sum expectation collapses to under VAR-based expectations.
5. **OLS on the PV-residualised LHS.** Subtract the structural PV terms (each at coefficient = +1 per wp1044, or at $-\sigma$ for user-cost PV terms in the business-inv equation; subtract the derived growth-neutrality coefficient $(1 - \sum_k \beta_k - \omega) \cdot \Delta \bar{y}_{t-1}$) from the LHS, then OLS for the free coefficients $\beta_0, \beta_1, \dots$ and the COVID/period dummies. Use classical OLS standard errors (`data/pac_helpers/ols_with_se.m`); damping factor $\alpha = 0.5$ and parameter clamps $\beta_k \in [-1, +1]$ to prevent iteration divergence.
6. **Iterate to convergence.** Re-compute χ from the updated β ’s, re-compute the PV operator, re-fit OLS. Convergence criterion $\|\beta^{(k)} - \beta^{(k-1)}\|_\infty < 10^{-4}$. All 4 fitting blocks converged in 3–21 iterations; the business-inv block converged numerically but with coefficients hitting the clamps (the deeper rejection signal documented in §4.6.2 and Appendix G).

Outputs per block: `data/pac_blocks/results_<block>.{mat,txt}` containing $\hat{\beta}, \hat{\sigma}_\beta, \hat{\chi}, R^2, N$, the active VAR state names, and a comparison row against

the wp1044 reference values. The per-block `.txt` artefacts are the source-of-truth for every coefficient quoted in Tables 4.3.2 / 4.4.4 / 4.5.2 / 4.6.1.x / 4.6.2.

Table 3.6: Block-specific VAR states used in Phase L2

Block	wp1044 reference	State vector (AU L2 implementation)	nstate
VA-price (Eq 16)	Table 3.3.4	$[\hat{y}, \tilde{i}_{\text{gap}}, \tilde{\pi}_{\text{Q_gap}}, \hat{y}_{\text{us}}, \tilde{\pi}_{\text{us}}, \hat{u}, \tilde{\pi}_{\text{w_eff}}, \bar{\pi}_{\text{Q_bar}}]$	8
Employment (Eq 30)	Table 3.4.10	$[\hat{y}, \tilde{i}_{\text{gap}}, \tilde{\pi}_{\text{Q_gap}}, \hat{y}_{\text{us}}, \tilde{\pi}_{\text{us}}, \hat{n}_{\text{S}}]$	6
Consumption (Eq 35)	Table 3.5.3	$[\hat{y}, \tilde{i}_{\text{gap}}, \tilde{\pi}_{\text{gap}}, \hat{y}_{\text{us}}, \tilde{\pi}_{\text{us}}, \hat{y}_{\text{H_gap}}, \Delta \text{w_eff}, \hat{u}] + 8$ same-vector lag-1 stacks	16 (with lags)
Housing inv (Eq 37)	Table 3.5.8	$[\hat{y}, \tilde{i}_{\text{gap}}, \tilde{\pi}_{\text{Q_gap}}, \hat{y}_{\text{us}}, \tilde{\pi}_{\text{us}}, \hat{p}_{\text{SH}} - \hat{p}_{\text{IH}} \text{ gap}, \hat{y}_{\text{H_gap}}]$	7
Business inv (Eq 46)	Table 3.5.14	$[\hat{y}, \tilde{i}_{\text{gap}}, \tilde{\pi}_{\text{Q_gap}}, \hat{y}_{\text{us}}, \tilde{\pi}_{\text{us}}, \hat{q}, \hat{i}_{\text{KB}}]$	7

Each block’s auxiliary equations (wp1044 Eqs 17, 31, 33, 36, 45) project the structural target onto these states. Where AU data does not have direct observables for wp1044 state components (notably the wp1044 EA-foreign block — replaced with US — and wp1044’s wage-Phillips Eq 18 + Okun Eq 19 aux structure — proxied by $\tilde{\pi}_{\text{w_eff}}$), the AU state vector substitutes the closest direct AU observable; these substitutions are documented in `PAC_EQUATIONS_AUDIT.md` per block.

3.6 Hybrid calibration: when local data doesn’t identify

For four of five PAC blocks (§4.3–§4.6.1), the Phase L2 iterative-OLS pipeline of §3.5 produces structurally interpretable coefficients on AU data: the PAC expectation coefficient is structurally imposed at +1 per the FOC derivation,

and the free coefficients $\beta_0, \beta_1, \dots$ converge to values that are economically signed, statistically significant, and (in the case of the consumption β_0) quantitatively close to wp1044's French estimate. The model is then used as-estimated: the AU L2 coefficients enter the production `au_pac.mod` as calibration starting values, are refined in §6.4's Bayesian full-system estimation, and produce IRFs in §7.

For the fifth block (business investment), the same pipeline produces a structural rejection. The full diagnostic is in §4.6.2 and Appendix G; the summary is: 11 specification variants tested, none deliver $R^2 > 0$ on raw `dln_ib` under strict-PAC (coef=+1) constraints, and when PV terms are free-estimated the lead $PV(\Delta\hat{q})$ coefficient comes out at -5.03 instead of $+1$. This is not a finite-sample noise issue: the coefficient mismatch is order-of-magnitude, persists across spec changes including replacing the target variable ($q \rightarrow$ terms-of-trade-augmented q_AU) and lag-structure changes, and is documented at length in `PAC_BI_AU_EXPLORATION.md`. Three non-exclusive hypotheses for *why* AU BI rejects the wp1044 PAC restriction are catalogued in §5.3.

Decision rule (small-open-economy modelling tradition): when local data does not identify the deep parameters of a single block and the production model's purpose is IRF / counterfactual / policy analysis (rather than now-casting/forecasting that single series), it is standard to **import the deep parameters from a larger or better-identified economy** and let the country-specific dynamics enter through the rest of the model. Concretely for AUS-PAC: the BI block's six deep parameters ($\beta_0, \beta_1, \beta_2, \beta_3, \omega, \sigma_{CES}$) are set to their wp1044 Table 3.5.13 values; the wp1044 PAC structural form (coef = +1 on $PV(\Delta\hat{q}) + PV(\Delta\bar{q})$, coef = $-\sigma$ on $PV(\Delta\log \hat{r}_{KB}) + PV(\Delta\log \bar{r}_{KB})$) is preserved in the simulation; AU-specific dynamics enter the BI block through (i) the AU-driven E-SAT VAR state that propagates into BI's PAC expectation, (ii) the other 4 PAC blocks' AU-estimated dynamics that feed BI via the GDP identity / demand block, (iii) AU-sized shocks (`stderr eps_ib` = 2.7211 AU-estimated), and (iv) AU-specific commodity and mining-cycle shocks that flow through trade and the deflator block.

Pros and cons:

- *Pros*: Strict PAC preserved; IRFs through the BI expectation channel are structurally correct (coefficient = +1 on PV terms); Lucas-critique-immune in counterfactual analysis; defensible cross-country tradition (the small-open-economy literature, e.g. the BoC MUSE and BoE LENS papers, also import parameters from large-country reference models when local data doesn't identify).
- *Cons*: BI dynamics in the AU model are imported, not estimated; AU mining-cycle dynamics are not directly captured in the BI block's own coefficients (they enter only via the E-SAT VAR's response to commodity shocks); the model fits AU history less well on `dln_ib` than a free-OLS reduced form would (the diagnostic free-OLS reduced form is preserved in `data/pac_blocks/estimate_pac_business_inv_au_v3.m` Variant A and gives $R^2 = 0.53$; it is **not** used in the production model, but is available

for historical-decomposition exercises).

Disclosure: the Bayesian estimation file `au_pac_bayesian.mod` removes the six BI parameters from `estimated_params`, treating them as calibrated. The other four PAC blocks remain in `estimated_params` and are refined from their AU L2 starting values through the joint posterior. This is the design adopted for the production paper, the IRFs in §7, and the cross-block findings in §5.

4. Model Specification

4.1 Notation

Prefix/Suffix	Meaning	Example
<code>dln_</code>	Log difference (quarterly growth rate)	<code>dln_c</code> = consumption growth
<code>_gap</code>	Deviation from target/trend	<code>i_gap</code> = <code>i_au</code> - <code>ibar</code>
<code>_star</code>	Target/desired value	<code>dln_c_star</code> = target consumption growth
<code>_bar</code>	Trend/long-run anchor	<code>pi_bar_au</code> = LR inflation anchor
<code>pi_</code>	Inflation rate	<code>piQ</code> = VA price inflation
<code>s_</code>	Spread	<code>s_COE</code> = equity spread
<code>pv_</code>	Present value of expectations	<code>pv_u_gap</code> = PV of future unemployment gaps
<code>_level</code>	Log-level accumulator for PAC	<code>ln_c_level</code> = cumulated consumption

All variables are in quarterly percentage points unless noted otherwise. At steady state, all gap and growth variables equal zero; inflation and interest rate variables equal their long-run anchors.

4.2 Supply block (FR-BDF Section 4.3)

4.2.1 CES production function (FR-BDF eq 24; Dubois et al. 2026 eq 1) The production technology for market branches is a constant-elasticity-of-substitution (CES) function with labour-augmenting technical progress (eq. 10):

$$Q_t^* = \gamma \left[\alpha K_t^{(\sigma-1)/\sigma} + (1-\alpha) (E_t H_t N_t^*)^{(\sigma-1)/\sigma} \right]^{\sigma/(\sigma-1)}$$

where K_t is the capital stock, N_t^* trend employment, H_t trend hours per worker, and E_t a labour-augmenting efficiency index. The elasticity of substitution between capital and effective labour is σ ; α is the CES distribution parameter; γ is a scale parameter. The Cobb-Douglas form is recovered as the special case $\sigma \rightarrow 1$.

Log-linearisation around a base-year steady state. Differentiating logarithmically, the CES law of motion for potential output is

$$\Delta \ln Q_t^* = s_K \Delta \ln K_t + s_N (\Delta \ln E_t + \Delta \ln H_t + \Delta \ln N_t^*)$$

with state-contingent factor shares

$$s_K = \alpha \gamma^{(\sigma-1)/\sigma} (K_t/Q_t^*)^{(\sigma-1)/\sigma}, \quad s_N = (1-\alpha) \gamma^{(\sigma-1)/\sigma} (E_t H_t N_t^*/Q_t^*)^{(\sigma-1)/\sigma}$$

evaluated at the steady state. Following Dubois et al. (2026, Section 3.1.2), we calibrate γ such that at a base year T^* we have $K_{T^*} = E_{T^*} H_{T^*} N_{T^*}^*$, i.e. the steady-state capital stock equals effective labour input. Under this normalisation the level CES collapses to $Q_{T^*}^* = \gamma K_{T^*}$, and the steady-state shares simplify to $s_K = \alpha$ and $s_N = 1 - \alpha$, *independently of* σ . The first-order log-linearisation around the base-year steady state therefore takes the form used in the model code:

$$\Delta \ln Q_t^* = \alpha \Delta \ln K_t + (1-\alpha) (\Delta \ln E_t + \Delta \ln H_t + \Delta \ln N_t^*) \quad (10)$$

This equation is algebraically identical to a Cobb-Douglas linearisation but is best understood as the first-order CES approximation under the FR-BDF 2026 base-year normalisation. The substitution elasticity σ enters separately through the factor-demand first-order conditions in Sections 4.4.2 (employment target) and 4.6.1 (investment target), which respond to relative-price changes with an elasticity governed by σ rather than by the constant shares α and $1 - \alpha$.

In AUSPAC's notation we summarise the joint contribution of hours and efficiency growth as a labour-augmenting TFP process $\ln TFP_t$ such that $\Delta \ln TFP_t \equiv (1-\alpha)(\Delta \ln E_t + \Delta \ln H_t)$; equivalently, labour productivity grows at rate $\Delta \ln Prod_t = \Delta \ln TFP_t/(1-\alpha)$. Equation (10) is then written in the model code as

$$\Delta \ln Q_t^* = \alpha \Delta \ln K_t + (1-\alpha) \Delta \ln N_t^* + \Delta \ln TFP_t$$

which is the form actually implemented in `eq_dln_y_star` of `au_pac.mod`.

4.2.2 Capital accumulation, TFP process Capital accumulation (FR-BDF eq 32, eq. 11):

$$\Delta \ln K_t = (1 - \delta_K) \Delta \ln K_{t-1} + \delta_K \Delta \ln I_t^B$$

TFP process (eq. 12):

$$\Delta \ln TFP_t = \rho_{TFP} \Delta \ln TFP_{t-1} + \varepsilon_t^{TFP}$$

with $\rho_{TFP} = 0.99$ (near unit root).

4.2.3 CES calibration (FR-BDF 2026 method, Dubois et al. Section 3.1.2) The Dubois et al. (2026) update to FR-BDF introduces a streamlined three-step calibration of the CES production function, replacing the 2019 grid-search procedure:

1. **Scale parameter** γ is calibrated analytically from observed output and capital in a base year (we use 2019, pre-COVID and post-mining-boom): $\gamma = \exp(\ln \overline{Q_{T^*}} - \ln \overline{K_{T^*}})$ where the overline denotes the within-year mean. This uses the $K_{T^*} = E_{T^*} H_{T^*} N_{T^*}^*$ normalisation discussed in §4.2.1.
2. **Substitution elasticity** σ is estimated from the long-run labour first-order condition (FR-BDF 2026 eq 3), not from the 2019-version investment FOC. The labour FOC is robust to the unconventional-monetary-policy era and to the AU mining-boom episode that drove a wedge between real user cost and investment in the 2019-version specification. Unobserved efficiency E_t is proxied by trend labour productivity $\hat{\Phi}_t$ (FR-BDF 2026 eq 6), a deterministic AR(1) trend with slope breaks at 2002Q2 and 2008Q3 plus a calibrated COVID level shift.
3. **Distribution parameter** α and **markup** μ are calibrated from independent AU sources. The FR-BDF cross-restriction approach is not informative on AU data because γ reflects the AU chain-volume convention rather than a structural quantity (see comparison table below). We set $\alpha = 0.45$ to match the AU capital-income share from ABS 5204 Table 48, and $\mu = 1.20$ to the mid-range of RBA RDP markup estimates (Hambur 2018, Andrews & Hambur 2022).

The trend efficiency $\bar{E}_t$ is then estimated as a second AR(1) regression with the same two-break structure (FR-BDF 2026 eq 7), using the Solow residual recovered from the level CES at the calibrated (σ, α, γ) . The estimation driver is `data/estimate_ces_2026.m`; outputs are written to `dynare/ces_2026_calibration.{txt,mat}`.

Table 4.2.1: Supply block — AU calibration (FR-BDF 2026 method)

Parameter	Symbol	AUSPAC value	FR-BDF 2026 (France)	Description
CES elasticity	σ	0.5366 (labour FOC, FD spec, prior $N(0.5, 0.2^2)$)	0.4951	Substitution between K and L
Capital share	α	0.45 (AU capital- income share, ABS 5204 Tab 48)	0.21	CES distribution parameter
Scale parameter	γ	0.0458 (2019 Q_market/K_total mean)	0.2561	CES scale parameter (units- driven)
Markup	μ	1.20 (RBA RDP 2018-09 mid-range)	1.33	Aggregate markup
Depreciation	δ_K	0.025	(varies)	~10% annual
TFP persistence	ρ_{TFP}	0.99	(n/a, exogenous trend)	Near unit root
ULC pass- through	$\gamma_{ULC} = (1 - \alpha)\sigma$	0.2951	0.40	Linearised CES dual (eq 4)
User-cost pass- through	$\gamma_{UCK} = \alpha\sigma$	0.2415	0.10	Linearised CES dual (eq 4)
Trend $\bar{E}$ growth pre- 2002Q2	—	2.54% p.a.	2.40%	Step 2 of FR-BDF 2026
Trend $\bar{E}$ growth 2002Q2– 2008Q3	—	0.87% p.a.	1.40%	Step 2 of FR-BDF 2026
Trend $\bar{E}$ growth post- 2008Q3	—	0.39% p.a.	0.70%	Step 2 of FR-BDF 2026

The calibration follows Dubois et al. (2026, §3.1.2): γ from the base-year Q/K mean under the $K_{T^*} = E_{T^*} H_{T^*} N_{T^*}^*$ normalisation; σ from the long-run labour FOC (eq 3) with trend productivity $\hat{\Phi}_t$ proxying unobserved efficiency; α and μ

from independent AU sources because the FR-BDF cross-restriction approach is not informative on AU data (AU’s $\gamma = 0.046$ reflects ABS chain-volume conventions rather than a structural quantity, so the implied markup is outside any plausible range). The factor-demand elasticity $\sigma = 0.54$, however, is structural and statistically close to the FR-BDF 2026 estimate of 0.50 (FD-spec t-stat 4.38 on 105 quarters). Driver: `data/estimate_ces_2026.m`; outputs in `dynare/ces_2026_calibration.{txt,mat}`.

Why the 2026 method replaced the 2019 grid search on AU data. The 2019-method investment-FOC estimation of σ failed for AU (see `data/estimate_sigma_stage1.m`: “level OLS, FD OLS, and FD without homogeneity all gave wrong-signed or implausible estimates due to mining-boom user-cost endogeneity”). Similarly, the 2019 grid-search calibration of (α, γ) over intercept cross-restrictions failed (see `data/estimate_ces_stage23.m`: “AU min L1 » FR-BDF tolerance, falling back to AU-economic calibration”). Both AU failures are exactly the problems Dubois et al. (2026) document for French data after the unconventional-monetary-policy era — the FR-BDF 2026 methodology innovations were motivated by precisely the data pathologies AUSPAC encountered.

L1.2 block-specific HP-filtered trends. The Phase L2 data layer (`data/prepare_12_data.m`) constructs HP-filtered trend objects for each PAC block’s target variable. The trends feed both the structural targets (the wp1044 ECM regressors of §3.5 and §4.3–§4.7) and the growth-neutrality regressors (the derived-coefficient $(1 - \Sigma\beta - \omega) \cdot \Delta\bar{y}_{\{t-1\}}$ term). Mean trend growth rates over the 1994Q1–2024Q4 sample, with first-eight-quarter and last-eight-quarter means for comparison (annualised %, computed by `data/make_paper_artifacts.py`):

block trend series	mean (% p.a.)	first 8q mean	last 8q mean
real GDP trend (y_bar)	3.02	3.68	2.06
market VA trend (q_bar)	3.07	3.87	1.87
employment trend (n_bar)	1.87	0.51	2.87
housing inv trend (I_H_bar)	3.94	6.74	3.64
user cost trend (r_KB_bar)	−0.96	1.07	26.34
VA-price inflation trend (pi_Q_bar)	0.00	0.02	−0.26

Three structural observations. (1) Real GDP trend growth has decelerated from 3.7% p.a. in the early-1990s to 2.1% p.a. in the last eight quarters of the sample — consistent with the slowing potential-output debate in Australian macro (Cusbert 2017, Bishop and Cassidy 2018). (2) The user-cost trend (r_KB_bar) shows the biggest end-of-sample jump (+26% p.a. in the last 8 quarters) reflecting the 2022-23 RBA tightening cycle pushing the nominal user cost up. (3) The market-VA trend (q_bar , 3.07% p.a. mean) is close to but a touch above the real-GDP trend, consistent with market-VA (excluding non-market services

like government and dwellings) being a slightly faster-growing aggregate. These trend numbers feed directly into the L1.1 trend-efficiency $\bar{E}_t$ calibration (Table 4.2.1) and the growth-neutrality regressors that the iterative-OLS pipeline (§3.5) imposes on every PAC equation.

4.3 Value-added price of market branches (FR-BDF Section 4.4)

The VA price equation is central because this deflator enters all other price equations, enabling expectations to affect price setting economy-wide.

4.3.1 Target equation (Factor price frontier, FR-BDF eq 38) Theoretical motivation (FR-BDF wp736 §4.4.1 / Dubois et al. 2026 §3.5.15). The long-run VA price is anchored to the CES dual cost function: in log-deviation form, the implied target inflation can be written as

$$\pi_t^{Q*} \approx \rho_{pQ}^* \pi_{t-1}^{Q*} + \gamma_{ULC} \Delta \ln ULC_t + \gamma_{UCK} \Delta \ln UC_t^K + (1 - \rho_{pQ}^* - \gamma_{UCK}) \bar{\pi}_t$$

with $\gamma_{ULC} = (1 - \alpha)\sigma$ and $\gamma_{UCK} = \alpha\sigma$ from the FR-BDF 2026 CES dual (§4.2). At AU calibration $\sigma = 0.5366$, $\alpha = 0.45$: $\gamma_{ULC} \approx 0.295$ and $\gamma_{UCK} \approx 0.241$.

AU-PAC implementation note (Phase Y, 2026-05-17): AU-PAC does NOT carry the analytical π^{Q*} equation as a separately-coded structural equation. Following FR-BDF wp736 Table 4.4.4 and ECB-BASE WAPRO conventions — both of which operationalise the PAC target as a single aux regression onto the satellite VAR state, not as two parallel objects — the canonical AU-PAC target is the var_model state $\widehat{\pi Q}_t$ (`piQ_hat`) described in §4.3.3 below. The CES factor-price frontier above is the theoretical reference for the long-run anchor; in the simulation, that anchor is enforced through (i) the BGP-consistent E-SAT state vector that $\widehat{\pi Q}_t$ projects onto and (ii) the inflation-anchor coefficient $(1 - \rho_{pQ,aux})$ in the regression. Earlier versions (Phase Q–X) carried `eq_piQ_star`, `eq_pQ_star_level`, `eq_pQ_gap` as orphan diagnostic equations; these were removed in Phase Y to eliminate the dual-target ambiguity. They were verified to have zero effect on simulation dynamics (no other equation referenced them on the RHS); their removal left the Laplace LMD and MHM LMD numerically identical.

4.3.2 Short-run PAC equation (FR-BDF eq 44) The short-run dynamics use the PAC framework with Dynare’s `pac_expectation()` (eq. 14):

$$\Delta p Q_t^{level} = b_0^{pQ} (\widehat{\pi Q}_{t-1} - p Q_{t-1}^{level}) + b_1^{pQ} \Delta p Q_{t-1}^{level} + pac_exp(pac_pQ) + b_2^{pQ} \hat{y}_t + \varepsilon_t^{pQ}$$

where the ECM term pulls $p Q_t^{level}$ toward $\widehat{\pi Q}_{t-1}$ — the var_model aux regression of the VA-price target onto the augmented E-SAT state (§4.3.3) —

and `pac_exp(pac_pQ)` projects the discounted-sum future of $\widehat{\pi Q}$ via the policy-function coefficients computed by `pac.print()`. Both the ECM target and the PAC-expectation projection reference the SAME object ($\widehat{\pi Q}$), matching FR-BDF wp736 Table 4.4.4 and ECB-BASE WAPRO architecture.

Table 4.3.2: VA price short-run coefficients — Phase L2 wp1044-faithful iterative OLS vs FR-BDF

Parameter	wp1044 FR Table 3.3.3	AU L2 iterative OLS (<code>results_va_priceman</code> , N=108)	Bayesian posterior (<code>man</code> , (hybrid MCMC, 20k×2)	90% HPD
β_0 (EC speed on $p^*_Q - p_Q$)	0.05	0.258 (SE 0.057, t=4.55)	0.0098	[0.0011, 0.0221]
β_1 (πQ lag)	0.20	0.304 (SE 0.092, t=3.30)	0.1115	[0.0355, 0.1770]
β_2 (contemp $\hat{y}$)	0.09	−0.076 (SE 0.107, t=−0.71)	0.0212	[−0.0419, 0.0988]
ω (calibrated per wp1044)	0.62	0.62	0.46 (legacy)	—
χ (derived from depth-1 polynomial)	~0.03	0.174	n/a	—
Derived growth-neutrality coef (1− β_1 − ω)	0.18	0.076	—	—
R²	0.61	0.41	n/a	—
Significant dummies	—	d_08Q1=−2.15 (t=−2.82), d_20Q3=−1.80 (t=−2.02), d_21Q1=+2.17 (t=+2.79)	COVID Q2 −2.88, Q3 +1.49	—

The Phase L2 column replaces the earlier “OLS (hybrid+COVID)” column as

the structurally faithful AU estimate of wp1044 Eq 16. Three differences from the earlier AUSEPAC OLS that are worth flagging:

- $\beta_0 = 0.258$ vs the pre-L2 0.027 is a $10\times$ difference, driven by the change in the ECM regressor: the Phase L2 ECM is the structural price-level gap $\mathbf{p}^*_Q - \mathbf{p}_Q$ (constructed by cumulating $\pi^*_Q - \pi_Q$ per wp1044 Eq 17, with the wage-Phillips Eq 18 and Okun Eq 19 entering through the block-specific VAR's $\tilde{\pi}_{w_eff}$ state). The earlier pre-L2 ECM proxy was $\text{lag}(\mathbf{yhat_au})$ (the output-gap lag), which is structurally different and produced a much smaller ECM speed. The Phase L2 $\beta_0 = 0.258$ is $5\times$ the wp1044 FR value of 0.05 — see §5.1 for cross-block cataloguing of the AU-vs-FR ECM-speed gap.
- $R^2 = 0.41$ vs wp1044's 0.61: a 0.20 gap. Primary remaining gap is the wp1044 Phillips Eq 18 + Okun Eq 19 auxiliary equations are present only via the $\tilde{\pi}_{w_eff}$ state in the block VAR; a more-faithful construction would estimate π^*_Q via Eq 17 OLS and add Phillips+Okun aux equations explicitly to the VAR state (~ 1 day extra; flagged in §4.3.3 below as a residual gap).
- β_2 wrong sign on $\mathbf{yhat_t_contemp}$ (-0.08 vs wp1044's $+0.09$): t-statistic -0.71 , statistically indistinguishable from zero. Sampling noise rather than structural disagreement; the Phillips-curve channel into VA-price runs primarily through the target equation (ULC + user-cost pass-through with $\gamma_{ULC} + \gamma_{UCK} = 0.54$ — see §4.2) rather than the short-run ad-hoc output-gap term, consistent with wp1044's design.

The Bayesian posterior column is preserved from the Phase Y full-system MCMC (pre-L2); it estimates the AUSEPAC reduced-form CPI Phillips parametrization (which differs from wp1044 Eq 16 — see §4.4.0 for the CPI Phillips, distinct from the VA-price Phillips treated here). Growth neutrality is verified: $1 - \beta_1 - \omega = 1 - 0.304 - 0.62 = 0.076$, matching the derived coefficient on $\text{lag}(\mathbf{pi_Q_bar})$ in the Phase L2 regression.

4.3.3 E-SAT auxiliary regression (FR-BDF Table 4.4.4) The VA-price auxiliary equation in the `var_model` — the canonical PAC target after Phase Y — projects $\widehat{\pi Q}$ onto the augmented E-SAT state vector with Bayesian posterior coefficients from `simulation/identities/calibration.inc`:

$$\widehat{\pi Q}_t = 0.334 \widehat{\pi Q}_{t-1} + 0.043 \hat{y}_{t-1} - 0.021 \tilde{i}_{t-1} + 0.007 \tilde{\pi}_{t-1} - 0.021 \tilde{u}_{t-1} + 0.437 \tilde{\pi}_{t-1}^w + \varepsilon_t^{aux,pQ}$$

(Phase W posteriors; Phase U wage-channel re-estimate for $a_{pQ,w} = 0.437$.) Compare FR-BDF wp736 Table 4.4.4 which reports $\rho_{pQ,aux} = 0.70$ (more persistent than the AU posterior 0.334) and a contemporaneous wage coefficient of 0.59 on $\tilde{\pi}^w = \pi^w - \Delta \bar{e}$ in the analytical eq (45). AU-PAC applies the coefficient to LAGGED $\tilde{\pi}_{gap}^w = \pi^w - \bar{\pi}^{au}$ (`var_model` requires pure-VAR form; the gap subtracts the inflation anchor rather than the efficiency trend). The Phase

U AU re-estimation found 0.437 on the AU sample — the production runtime value.

Residual L2 gap: the Phase L2 iterative-OLS VA-price block (Table 4.3.2 above) gives $R^2 = 0.41$ vs wp1044's 0.61. The dominant remaining gap is that wp1044's auxiliary VAR (Table 3.3.4) includes the wage Phillips Eq 18 and Okun Eq 19 as explicit auxiliary equations that feed the construction of π^*_Q (the PV target), whereas the AU L2 implementation enters the wage/unemployment channel through the single $\tilde{\pi}_{w_eff}$ state in the 8-variable block VAR. A more-faithful construction (estimate π^*_Q separately via Eq 17 OLS; add Phillips + Okun as their own AR(1) aux equations with the wp1044 coefficients) is estimated at ~1 day of additional work and would close most of the R^2 gap.

4.3.4 Dynamic contributions *Figure 4.3.1: Historical dynamic-contribution decomposition of quarterly $\Delta \log pQ$ over 1994Q3–2024Q4 (top: full sample; bottom: pre-COVID zoom). The decomposition is a post-estimation visualisation: given the OLS parameters from Table 5.6, Dynare's Kalman smoother reconstructs the unobserved auxiliary states (pv_piQ_aux etc.), and each bar stacks the contribution from a specific channel of the structural equation. This is not a smoothing of the input data and does not affect the parameter estimates. Error correction ($b0_pQ_gap$, blue) and the residual eps_pQ (gray) are the most active channels; the PAC expectation contribution is small because Australian VA price inflation has low forward-looking content under the estimated companion matrix.*

4.4 Wages and employment (FR-BDF Section 4.5)

4.4.0 E-SAT inflation equation and expectation formation The model's CPI inflation gap $\tilde{\pi}_t^{au}$ evolves according to a reduced-form Phillips equation augmented with the FR-BDF structural deflator channels (wp736 §3.1.1) plus two AU-specific Phase V additions:

$$\tilde{\pi}_t^{au} = \lambda_\pi \tilde{\pi}_{t-1}^{au} + \kappa_\pi \hat{y}_{t-1} + \alpha_{pc}(\pi_t^Q - \bar{\pi}_t^{au}) + \alpha_{pc,lag}(\pi_{t-1}^Q - \bar{\pi}_{t-1}^{au}) + \beta_{pc,m}(\pi_t^m - \bar{\pi}_t^{au}) + \gamma_{oil} \Delta \ln pcom_t + b_{ECM,pc}(p_{t-1}^{C,*} - p_t^Q)$$

with $\alpha_{pc} = 0.20$, $\alpha_{pc,lag} = 0.11$, $\beta_{pc,m} = 0.10$, $\gamma_{oil} = 0.03$, $b_{ECM,pc} = 0.06$ (Bayesian posterior modes for α_{pc} , $\alpha_{pc,lag}$, $b_{ECM,pc}$; calibrated for the others). The reduced-form persistence λ_π and Phillips slope κ_π are jointly Bayesian-estimated with the rest of the model. The cumulative CPI target $p_t^{C,*}$ is constructed as $p_t^{C,*} = (1 - \omega_{pc})p_t^Q + \omega_{pc}p_t^M$ with the IAD-weighted import share $\omega_{pc} = 0.23$ (matching the long-run CPI target eq 79 of wp736).

Deviations from FR-BDF wp736: two terms in the equation above are AU-PAC-specific and have no counterpart in the FR-BDF Phillips curve (wp736 eq 28). - *Lagged VA-price passthrough* $\alpha_{pc,lag}(\pi_{t-1}^Q - \bar{\pi}_{t-1}^{au})$: wp736 has only the

contemporaneous α_{pc} term. We add the one-quarter lag to capture the empirical AU pattern that import-content-heavy components of the consumer basket reprice with a lag of about a quarter; the joint $\alpha_{pc} + \alpha_{pc,lag}$ total passthrough (~ 0.31) better matches the AU long-run elasticity of CPI to VA prices than α_{pc} alone. - *ECM correction* $b_{ECM,pc}(p_{t-1}^{C,*} - p_{t-1}^C)$: in FR-BDF, the long-run anchoring of CPI runs through the consumption deflator's own two-equation target-plus-ECM structure (wp736 eqs 79-80). Because AUSEP collapses that two-equation structure into the single `eq_pi_c` reduced form (see §4.9), the long-run anchor would otherwise be lost; we re-introduce it via this ECM term directly inside the Phillips curve. This is a pragmatic AU innovation, not a theoretical refinement — its function is to keep the cumulative CPI trajectory pinned to the weighted VA+import target rather than drifting freely under reduced-form cost-push shocks.

Growth neutrality holds at the steady state: $\tilde{\pi}_t^{au} = 0$, $\pi^Q = \bar{\pi}^{au}$, $\pi^m = \bar{\pi}^{au}$, $\Delta \ln pcom = 0$, and at the BGP $p^{C,*} = p^C$ so the ECM term vanishes. All RHS terms vanish at the SS.

PAC expectation formation: policy-function approach. Following FR-BDF wp1044 §2.2 + §3.2.3 (“expectations are computed as policy functions obtained from inversion of the estimated core E-SAT, supplemented by the auxiliary equations”), the discounted present values appearing as `pac_expectation(pac_X)` in the five PAC short-run equations (§§4.3.2, 4.4.2, 4.5.2, 4.6.2, 4.7.2) are pre-computed as explicit closed-form linear combinations of lagged structural state variables:

$$PV(pac_X)_t = h_0^X + \sum_{Y \in \mathcal{S}} h_Y^X \cdot Y_{t-1}$$

where $\mathcal{S}$ is the augmented E-SAT state vector $\{\hat{y}_{au}, \tilde{i}_{gap}, \tilde{\pi}_{au,gap}, \tilde{u}_{gap}, \hat{y}_{us}, \tilde{\pi}_{us,gap}, \bar{i}, \bar{\pi}_{au}, \bar{\pi}_{us}, \pi^Q, \pi^m, \Delta \ln pcom\}$ plus the PAC-specific auxiliary regression target (π^Q for the VA-price PAC, $\hat{n}$ for employment, etc.). The h-coefficients are the discounted-sum entries of the H^k powers of the reduced-form companion $H = A^{-1}B$ of the structural E-SAT VAR $AZ_t = BZ_{t-1} + \varepsilon_t$, computed once at estimation time and held fixed in simulation.

The implementation uses Dynare 7.0's `var_model + pac_model + pac.print()` + `cherrypick()` + `aggregate()` workflow (Adjemian, Brayton & Zimic, 2024). For each PAC block, an auxiliary `.mod` file declares the structural E-SAT state in pure-VAR(1) form (lagged-only right-hand sides), defines the auxiliary regression that maps the PAC target onto the E-SAT core, declares the PAC equation, and emits the closed-form expectation formula via `pac.print()`. `cherrypick()` then extracts the simulation-ready equations into a per-block `.inc` bundle. The final simulation model is composed by `aggregate()` from the five cherrypicked bundles plus a normalized identity layer that contains the non-PAC structural equations (deflators, financial block, fiscal rule, GDP identity, sectoral wealth,

etc.). Because structural shocks like ε_Q^π enter the structural variable π^Q directly and lagged π^Q appears in every PAC expectation formula, the cost-push channel propagates into every PAC equation's forward-looking term — closing the transmission loop.

4.4.1 Wage Phillips curve (FR-BDF eq 52) In the long run the labour supply curve is vertical. In the short run, wage inflation is determined by a hybrid Phillips curve indexed to **consumer-price** inflation (FR-BDF eq. 52). The unemployment-PV term enters with negative sign — a positive κ_w produces wage *deceleration* when unemployment is above its steady state. Eq. (16):

$$\pi_t^w = \lambda_w \pi_{t-1}^w + \gamma_w \pi_t^c - \kappa_w pv_{u,gap} + (1 - \lambda_w - \gamma_w) \bar{\pi}_t + (1 - \lambda_w) \Delta \ln Prod_t + \varepsilon_t^w$$

Table 4.4.1: Wage Phillips curve coefficients

Parameter	Prior	Bayesian posterior mean	90% HPD
lambda_w (persistence)	Beta(0.25, 0.10)	0.209	[0.079, 0.332]
gamma_w (CPI passthrough)	Beta(0.45, 0.15)	0.354	[0.268, 0.428]
kappa_w (unemployment PV slope)	Normal(−0.08, 0.05)	−0.103	[−0.178, −0.019]
beta_w (PV discount, calibrated)	—	0.98	—
Inflation anchor (1−λ_w−γ_w)	derived	0.437	—
stderr eps_w	InvGamma(0.30)	0.141	[0.095, 0.191]

Estimated on the ABS Wage Price Index (Cat. 6345, Total hourly rates of pay excluding bonuses, Private and Public, All industries, seasonally adjusted), the AU wage Phillips curve takes the approximate form

$$\pi^w \approx 0.21 \pi_{t-1}^w + 0.35 \pi_t^c + 0.10 pv_{u,gap} + 0.44 \bar{\pi} + 0.79 \Delta \ln Prod$$

with three notable features. First, the CPI-passthrough coefficient $\gamma_w = 0.35$ is large and sharply identified — Australia's Fair Work Commission annual wage

review anchors a substantial fraction of award-rate and enterprise-agreement floor wages to the CPI, with flow-on effects to non-award private bargaining that tracks the FWC decision. Second, the unemployment-gap slope $\kappa_w = -0.103$ is statistically significant with 90% HPD entirely negative; under the sign convention $\dots - \kappa_w \cdot \text{pv_u_gap}$ this corresponds to a *positive* structural Phillips slope (higher unemployment $\rightarrow$ lower wage growth), consistent with the FR-BDF wp736 specification and with the post-2015 RBA literature on Australian wage dynamics (Cassidy and Doyle 2018; Bishop and Cassidy 2017). Third, own-lag persistence $\lambda_w = 0.21$ is moderate; the inflation-anchor term carries the remaining 0.44 weight.

Auxiliary: Okun’s law (FR-BDF eq 53, eq. 17): $\tilde{u}_t = \rho_u \tilde{u}_{t-1} + okun \times \hat{y}_t$ with $\rho_u = 0.94$, $okun = -0.33$.

PV of unemployment gap (FR-BDF eq 137, eq. 18): $pv_{u,gap} = (1 - \beta_w) \tilde{u}_t + \beta_w \cdot pv_{u,gap}(+1)$. Forward-looking under hybrid/MCE.

4.4.2 Employment (FR-BDF Section 4.5.2) Target equation (FR-BDF eq 36 / eq 55, eq. 19): the long-run employment-growth target combines a labour-augmenting TFP term, a real-wage-gap substitution term, and a demand-side Δq channel that tracks contemporaneous demand growth:

$$\Delta \ln N_t^{*,bar} = (\hat{y}_t - \hat{y}_{t-1}) - \frac{\Delta \ln TFP_t}{1 - \alpha_K} - \sigma_{CES} \cdot rw_gap_t$$

where $rw_gap = \pi_w - \pi^Q - \Delta \ln Prod$ is the real wage growth gap and the corrected sign on the TFP term $-\Delta \ln TFP / (1 - \alpha_K)$ reflects the labour-saving nature of TFP improvements (higher TFP reduces labour demand for a given output target).

Short-run PAC equation (wp1044 Eq 30, **depth = 3** per Phase L2 audit — the earlier AUSEPAC depth-4 specification was an over-parameterisation; reducing to wp1044’s depth-3 restored convergence and lifted R^2 ; see PAC_EQUATIONS_AUDIT.md §2.3 gap #1):

$$\Delta n_{S,t} = \beta_0(n_{S,t-1}^* - n_{S,t-1}) + PV(\Delta \bar{n}_S^*)_{t|t-1} + PV(\Delta \hat{n}_S^*)_{t|t-1} + \sum_{k=1}^3 \beta_k \Delta n_{S,t-k} + (1 - \sum_k \beta_k - \omega) \Delta \bar{n}_{S,t-1}^* + \beta_4 \Delta \hat{q}_t + \delta \cdot$$

Table 4.4.4: Employment short-run PAC coefficients — Phase L2 wp1044-faithful iterative OLS (depth 3) vs FR-BDF

Parameter	wp1044 FR Table 3.4.9	AU L2		
		iterative OLS (<code>results_employment</code> , depth 3, N=124)	Bayesian posterior (<code>team_maths</code> , hybrid MCMC, 20k×2)	90% HPD
β_0 (EC speed on $\mathbf{n^*_S - n_S}$)	0.07	0.314 (SE 0.043, t=7.41)	0.0537	[0.0142, 0.0883]
β_1 ($\Delta \mathbf{n}$ lag 1)	0.44	0.295 (SE 0.065, t=4.57)	0.2890	[0.1181, 0.4411]
β_5 (contemp $\Delta \hat{q}$ / output gap)	0.13	−0.026 (SE 0.026, t=−0.97)	−0.0050	[−0.0890, 0.0775]
ω (cali- brated per wp1044)	0.34	0.34	—	—
χ (derived from depth-3 polyno- mial)	~0.43	0.205	—	—
$\mathbf{R^2}$	0.95	0.81	—	—
stderr	—	—	0.3262	[0.1186, 0.5248]
eps_n	—	—	—	—
Significant COVID dummies	—	d_20Q2=−5.45 (t=−17.0), d_20Q3=+2.47 (t=+4.5)	—	—

The Phase L2 employment block delivers the **second-best fit of the five blocks** ($R^2 = 0.81$), after consumption. Three findings stand out:

- **Depth correction:** dropping the spurious 4th lag (the earlier AUPAC `dln_n` builder used depth-4 because of a misreading of wp1044’s lag count) was the single largest fix. The pre-L2 depth-4 specification failed to converge in the earlier `estimate_all_pac_iterative.m` shortcut; depth-3 converges in 6 iterations.
- $\beta_0 = \mathbf{0.314}$ vs wp1044’s **0.07** is **4.5× faster ECM** on AU than France’s. The cross-block pattern of AU faster ECM is summarised in §5.1.
- $\beta_2, \beta_3, \beta_4$ **new all wrong-signed** but statistically insignificant (t ∈ [−0.97, +0.62]) — AU employment has less of a deep-lag and contemp-

$\Delta\hat{q}$ structure than France's. The $PV(\Delta\hat{n}^*_S)$ term in the wp1044-faithful spec absorbs most of the forward-looking signal, leaving residual noise on the free lags. The structural restriction (coef = +1 on PV) is satisfied by the L2 spec; the remaining noise is finite-sample.

The bulk of the R^2 gap vs wp1044's 0.95 is attributable to (i) the wp1044 EA-foreign block being replaced with the US-foreign block in the AU VAR state (smaller variance contribution), and (ii) wp1044's wage Phillips and Okun aux equations entering the AU VAR only via the $\hat{n}^*_S$ AR(1) state (Eq 31 implementation), not separately as wp1044 implements them.

E-SAT auxiliary (wp736 Table 4.5.7, eq. 21): $\hat{n}_t = 0.67\hat{n}_{t-1} + 0.12\hat{y}_{t-1} - 0.03\tilde{i}_{t-1} + 0.05\tilde{\pi}_{t-1} - 0.04\tilde{u}_{t-1}$

Figure 4.4.1: Historical dynamic-contribution decomposition of $\Delta\log n$ (employment growth) over 1994Q3–2024Q4. Error correction ($b0_n \cdot (n_hat - \ln n)$) is positive throughout most of the sample, reflecting steady employment convergence to potential. The four AR lags partially cancel — net AR contribution is small. The $b5_n \cdot yhat_au$ accelerator is essentially zero (posterior 0.002, weakly identified) so output-gap variation does not feed directly into $\Delta\log n$ in the AU data; transmission to employment runs almost entirely through the PAC expectation term and the eps_n shock.

4.5 Household consumption (FR-BDF Section 4.6.1)

4.5.1 Target equation The consumption target is based on permanent income (FR-BDF eq 59). Permanent income is the discounted PV of expected future output gaps (eq. 22):

$$pv_{yh} = (1 - \beta_c)\hat{y}_t + \beta_c \cdot pv_{yh}(+1)$$

with $\beta_c = 0.95$ (~25% annual discount). The consumption target (eq. 23):

$$\Delta \ln C_t^{*,bar} = \kappa_{inc}(pv_{yh,t} - pv_{yh,t-1}) + \alpha_{c,r} \cdot \Delta(\text{real lending rate gap})$$

Table 4.5.1: Consumption target coefficients

Parameter	Symbol	Value	Description
Target persistence	rho_c_star	0.95	—
Permanent income sensitivity	kappa_inc	0.050	Calibrated
PV discount	beta_c	0.95	~25% annual discount
Real rate sensitivity	alpha_c_r	-0.95	wp736 Table 4.6.14

4.5.2 Short-run PAC equation (1st-order, FR-BDF eq 61) Eq. (24), following FR-BDF eq. 61 with a forward-NPV real-lending-rate-gap channel:

$$\Delta \ln C_t^{level} = b_0^c (C_{t-1}^{*,level} - \ln C_{t-1}^{level}) + b_1^c \Delta \ln C_{t-1}^{level} + pac_exp(pac_c) + b_2^c \tilde{i}_{t-1} + b_3^c \hat{y}_t + \alpha_{c,r} \cdot pv_{r,lh,gap,t} + \varepsilon_t^c$$

The new endogenous variable `pv_r_lh_gap` is the discounted PV of the real-lending-rate gap, defined recursively with the consumption discount factor $\beta_c = 0.95$:

$$pv_{r,lh,gap,t} = (1 - \beta_c)(i_t^{LH} - \pi_t^c - r_{SS}^{LH}) + \beta_c \cdot E_t[pv_{r,lh,gap,t+1}]$$

The coefficient $\alpha_{c,r} = -0.95$ (calibrated, wp736 Table 4.6.14) carries the forward-looking real-rate channel: an expected higher real lending rate today and over the future path lowers $\Delta \ln C$ directly through the PAC short-run equation, complementing the consumption target's $\alpha_{c,r} \cdot \Delta(\text{real lending rate gap})$ channel. Under MCE the recursion runs fully forward; under VAR it operates as a backward AR(1) of past real-lending-rate gaps. Including this channel is essential for the no-forward-guidance-puzzle property under MCE — without it, the consumption block under MCE lacks a stable forward-looking real-rate anchor.

Table 4.5.2: Consumption short-run PAC coefficients — Phase L2 wp1044-faithful iterative OLS vs FR-BDF

Headline finding: AU's consumption-block ECM speed $\beta_0 = 0.23$ is broadly in line with FR-BDF's $\beta_0 = 0.29$ — the closest of any block to its French counterpart, and the only block where AU's ECM speed is **not** 4–8× faster than France's (cf. §5.1). ($\beta_0 = 0.23$ is the 2026-05-31 re-estimate on the SA-corrected, deterministically-reproducible data layer; the pre-SA layer gave 0.266. The Δc -lag coefficient is clamped at its lower bound in both, a pre-existing specification feature, so the shift reflects the SA-data correction, not the clamp.)

		AU L2 iterative OLS (results_consumptBayesian, posterior SA layer, mean (legacy pre-L2 N=103, 22 iters) MCMC) ‡			90% HPD
Parameter	wp1044 FR Table 3.5.2				
β_0 (EC speed on c* - c)	0.29	0.2307 (SE 0.078, t=2.98)	0.0519		[0.0215, 0.0805]

Parameter	wp1044 FR Table 3.5.2	AU L2 iterative OLS (results_consumption, posterior Bayesian, mean (legacy pre-L2 MCMC) ‡)			90% HPD
		SA layer, N=103, 22 iters)			
β_1 (Δc lag 1)	0.17	-0.182 (SE 0.085, t=-2.14); χ -recursion clamps β_1 to its 0.01 floor	0.0397		[0.0051, 0.0731]
α_1 (PV(r_LH gap), reduced- form coef)	-1.15 (from LR Eq 33)	-126.7 (SE 35.5, t=-3.57); reduced form $\alpha_1 \cdot \chi_{AU} =$ -1.27	— (calibrated α_{c_r} = -0.95)		—
β_{PAC} (free coef on $\Delta \bar{y}$ lag)	(free, not tabled)	+1.56 (SE 0.67, t=2.33)	0.8263		[0.3315, 1.3595]
β_2 (HtM level-diff)	0.32	+0.027 (SE 0.075, t=0.35)	0.0285 ‡		[-0.0543, 0.1134]
β_3 (impact Δr_{LH} gap)	-1.07	-0.007 (SE 0.036, t=-0.20)	-0.5796 ‡		[-0.8848, -0.2866]
ω (cali- brated per wp1044)	0.00 (gap terms)	0.00	—		—
χ (derived from depth-1 polyno- mial)	~0.17	0.010	—		—
R²	0.95	0.82	—		—
stderr	—	—	2.0145		[1.8112, 2.2523]
eps_c	—	—	—		—
Significant COVID dummies	—	d_20Q2=-13.04 (t=-13.07), d_20Q4=+3.02 (t=+2.40)	—		—

Three quantitative findings:

1. **β_0 is comparable across AU and France.** AU consumption mean-reverts to its permanent-income target at a broadly similar speed to France ($\beta_0 = 0.23$ vs 0.29 on the SA-corrected layer; was 0.266 pre-SA — both within ~ 0.06 of the French value). Of the five PAC blocks this is the closest cross-country agreement and the only block where AU’s ECM speed is **not** several times faster than France’s, validating the wp1044 PAC framework for the consumption block specifically. The structural mechanism — agents minimising a quadratic cost function around a permanent-income target plus polynomial adjustment costs — fits AU data at a comparable speed to French data. *(All AU L2 sub-coefficients in the table above are now the 2026-05-31 SA-layer re-estimate ($N=103$, 22 iters), reproduced bit-for-bit by `data/pac_blocks/estimate_pac_consumption.m`; the earlier mixing of pre-SA $\beta_1/\alpha_1/\beta_{\text{PAC}}/\text{COVID}$ snapshots with the refreshed β_0 has been resolved. The Bayesian column is the deleted pre-L2 MCMC, retained for reference.)*
2. **α_1 rescaled match (now essentially matching France).** The reduced-form forward-real-rate effect is $\alpha_1 \cdot \chi$. On AU L2 (SA layer) $\alpha_1 = -126.7$ (inflated because $\chi_{\text{AU}} \approx 0.010$ is near zero from the low β_1), so the reduced form is $\alpha_1 \cdot \chi_{\text{AU}} = -1.27$, vs wp1044’s $\alpha_1 = -1.15$ — AU’s real-rate sensitivity in consumption is now $\sim 110\%$ of France’s (it was $\sim 70\%$, i.e. -0.85 , on the pre-SA layer). The SA correction brings AU’s forward-real-rate channel into close agreement with France, and it is structurally signed correctly.
3. **β_1 near zero and HtM not identified.** Both expected: β_1 hit the lower clamp (without the clamp it goes negative); β_2 HtM is not identified. The first matches the “fast-ECM-weak-PAC” pattern from the earlier AUSEPAC L1.3a iteration. The second matches the AUSEPAC Round 1.2 finding that **b_HtM** was hard to pin down in joint Bayesian estimation (the AU sample lacks the wages+transfers level-differential variation that identifies β_2 in wp1044).

‡ The Bayesian posterior column is preserved from the pre-L2 Phase Y full-system MCMC (since deleted in the OLS pivot) and uses the *production-model* parameterization (eq. 24 above: a lagged rate-gap term $\mathbf{b}_2 \cdot \mathbf{i}_{\{\mathbf{t}-1\}}$ and an output-gap/HtM term), so its β_2/β_3 rows are aligned to the nearest economic channel rather than an identical regressor to the wp1044 Eq-35 L2 estimator. The L2 column is the structurally faithful AU estimate. The $\alpha_{\{\mathbf{c}, \mathbf{r}\}} \cdot \mathbf{pv}_{\{\mathbf{r}, \mathbf{lh}, \mathbf{gap}\}}$ real-rate NPV channel in `au_pac.mod` (§4.5.2 above this table) carries the calibrated wp1044 value $\alpha_{\{\mathbf{c}, \mathbf{r}\}} = -0.95$, which is what enters the production IRFs. The Phase L2 free- α_1 estimate confirms the calibration is in the right neighbourhood: the rescaled $\alpha_1 \cdot \chi = -1.27$ reduced form is structurally identical to the `pv_r_lh_gap` channel that already exists in the model.

E-SAT auxiliary (wp736 Tables 4.6.3-4, eq. 25): $\hat{c}_t = 0.60\hat{c}_{t-1} + 0.06\hat{y}_{t-1} - 0.04\hat{i}_{t-1} + 0.005\hat{\pi}_{t-1} - 0.03\hat{u}_{t-1} + 0.39\hat{y}_{t-1}$

The consumption auxiliary is the most complex of the seven, featuring a nested PV structure in which the income-output ratio ($\widehat{yh}$) feeds into consumption ($\widehat{c}$) and creates PV-squared dynamics (wp736 Table 4.6.4).

Figure 4.5.1: Historical dynamic-contribution decomposition of $\Delta \log c$ (real per-capita consumption growth) over 1994Q3–2024Q4. The error correction term (blue) dominates in absolute magnitude — Australian consumption converges visibly to permanent income each quarter. The rate channel $b2_c_i_gap$ (red) flips sign at major policy turns (rate cuts in 2001, 2008, 2020 supportive of consumption; tightening in 2007 and 2022 dragging). The $b_di_c_di_gap$ MP-surprise channel (purple) is small but procyclical. The 2020 COVID quarter shows a -16% $\Delta \log c$ absorbed almost entirely by the eps_c residual (gray), confirming that the smoothed structural channels alone do not explain pandemic-era consumption collapse.

4.6 Business investment (FR-BDF Section 4.6.2, wp1044 §3.5.3 Eq 46) — Phase L2 P1c structural rejection + Option 1 hybrid calibration

Reader’s overview. This section documents the **single material change** introduced by Phase L2 P1c: the wp1044 PAC restriction (coef = +1 on $PV(\Delta \hat{q}) + PV(\Delta \bar{q})$, coef = $-\sigma$ on $PV(\Delta \log \hat{r}_{KB}) + PV(\Delta \log \bar{r}_{KB})$) is statistically rejected on AU business-investment data. After 11 specification variants (Appendix G), the production model imports the BI deep parameters from wp1044 Table 3.5.13 (Option 1 of PAC_BI_AU_EXPLORATION.md §7). The §4.6.1 target equation is unchanged. The §4.6.2 short-run PAC equation is unchanged in functional form (wp1044 Eq 46 verbatim) but its six deep parameters are now wp1044-imported rather than AU-estimated.

4.6.1 Target equation (FR-BDF eq 63 / wp1044 Eq 45) Desired investment derives from the CES capital demand first-order condition (eq. 26):

$$\Delta \ln I_t^{B*,bar} = \kappa_{ib,y} \hat{y}_t - \sigma_{CES} \Delta \ln UC_t^K$$

The user cost of capital (eq. 27):

$$UC_t^K = WACC_t + \delta_K - (\pi_t^{IB} - \pi_t^Q)$$

Table 4.6.1: Business investment target coefficients (unchanged in Phase L2)

Parameter	Symbol	Value	Description
Target persistence	rho_ib_star	0.95	Calibrated

Parameter	Symbol	Value	Description
Output proportionality	kappa_ib_y	0.06	CES eq 63
CES user cost elasticity	sigma_ces	0.5366	FR-BDF 2026 calibration (Table 4.2.1) $\approx$ wp1044's 0.50
Depreciation	delta_k	0.025	$\sim 10\%$ annual

4.6.2 Short-run PAC equation (wp1044 §3.5.3 Eq 46) — structural rejection on AU data The wp1044 spec is:

$$\Delta \ln I_{B,t} = \beta_0 \log(I_{B,t-1}^*/I_{B,t-1}) + \beta_1 \Delta \ln I_{B,t-1} + \beta_2 \Delta \ln I_{B,t-2} + \text{PV}(\Delta \hat{q})_{t|t-1} + \text{PV}(\Delta \bar{q})_{t|t-1} - \sigma \text{PV}(\Delta \log \hat{r}_{KB})_{t|t-1}$$

with FOUR PV terms (gap+trend decomposition for BOTH market VA AND user cost), TWO growth-neutrality terms at derived coefficient $(1 - \beta_1 - \beta_2 - \omega)$, and a synthetic final-demand-gap term $\beta_3(\Delta df - \Delta \bar{df})$ where $df = c + i_b + i_h + g + x$ (chain-linked).

Table 4.6.2: BI coefficients — wp1044 calibration (imported via Option 1) vs Phase L2 AU iterative-OLS attempts (rejected)

Parameter	wp1044 FR Table 3.5.13 (USED IN PRODUCTION MODEL)	AU L2 baseline iterative OLS (results_business_in_AU_L2, v3-A: PV free, N=120)	Results v3-A: PV free, R ² =0.53 — diagnostic only
β_0 (EC speed on I^*_B/I_B)	0.096	0.863 (SE 0.42, t=2.04) hit clamp	0.18 (free PV)
β_1 ($\Delta \log I_B$ lag 1)	0.33	1.042 (SE 0.50, t=2.07) hit clamp	0.21
β_2 ($\Delta \log I_B$ lag 2)	0.11	0.920 (SE 0.50, t=1.84) hit clamp	0.14
β_3 (Δdf gap)	0.69	0.859 (SE 1.40, t=0.61)	0.92
PV($\Delta \hat{q}$) coef (structural restriction)	+1.0 (imposed)	+1.0 imposed (structural rejection: wedge can't close)	-5.03 (data-driven: WRONG SIGN, WRONG MAGNITUDE)

	wp1044 FR Table 3.5.13 (USED IN PRODUC- TION MODEL)		
Parameter		AU L2 baseline iterative OLS (results_business_investment, v3-A: PV free, N=120)	AU L2, v3-A: PV free, R ² =0.53 — diagnostic only
PV($\Delta \bar{q}$) coef	+1.0 (imposed)	+1.0 imposed	+2.77
PV($\Delta \log$ $\hat{r}_{KB}$) coef	$-\sigma =$ -0.50	$-\sigma$ imposed	-1.25
PV($\Delta \log$ $\bar{r}_{KB}$) coef	$-\sigma =$ -0.50	$-\sigma$ imposed	+0.001
ω (calibrated per wp1044)	0.35	0.35	—
σ_{CES} (AU labour- FOC)	(wp1044 uses 0.50)	0.5366	0.5366
R²	0.83 (wp1044 reported)	0.09	0.53




Figure 1: AU BI: 11 spec variants tested in Phase L2 P1c

Figure 4.6.2: Eleven specification variants of wp1044 Eq 46 tested in Phase

L2 P1c on AU data. Every strict-PAC variant delivers $R^2 \leq 0.11$ (red bars); free-PV variant v3-A delivers $R^2 = 0.53$ but with $PV(\Delta\hat{q}) = -5.03$ (wrong sign vs structural +1). Dashed blue line: wp1044's reported $R^2 = 0.83$ on French data. Full table in Appendix G.

The headline diagnostic: when PV regressors are free-estimated on AU data (v3-A column), the lead $PV(\Delta\hat{q})$ coefficient comes out at **-5.03** — wrong sign and 5× wrong magnitude relative to the structural +1. This is the canonical signal that the wp1044 PAC FOC is rejected: the agent's forward-looking weight on market-VA-growth expectations on AU data is *negative*, not positive at unity. No amount of dummy expansion, terms-of-trade re-targeting, target-variable substitution, or trend-regime addition closes the wedge (see Appendix G for the 11-variant search). The deeper-structural diagnostic (why $PV(\Delta\hat{q})$ is negative on AU BI data) is discussed in §5.3, where three non-exclusive hypotheses are catalogued: (A) different agent objective (heterogeneous mining vs non-mining capex); (B) mining-sector dominance; (C) sample-period contamination (GST, GFC, mining boom, COVID, multiple regime breaks in 1993–2024 — many more than France's smoother 1985–2024 sample).

Option 1 implementation. The production simulation model `dynare/au_pac.mod` lines 1746–1771 set the BI block deep parameters to their wp1044 Table 3.5.13 values:

```
% Phase L2 P1c Option 1 (2026-05-26): wp1044 BI calibration import
b0_ib = 0.096;    // wp1044 Table 3.5.13
b1_ib = 0.33;
b2_ib = 0.11;
b3_ib = 0.69;
omega_ib = 0.35;
% sigma_ces = 0.5366 (AU labour-FOC) ≈ wp1044's 0.50
```

The Bayesian estimation `au_pac_bayesian.mod` removes BI parameters from `estimated_params` (treats them as calibrated). The diagnostic free-OLS reduced form (`data/pac_blocks/estimate_pac_business_inv_au_v3.m` Variant A, $R^2=0.53$) is preserved as a historical-decomposition tool and is **not** used in the production model.

Cross-references: - §3.6 (hybrid calibration: when local data doesn't identify) — methodological rationale. - §5.3 (cross-block findings: business investment structurally rejects wp1044; three hypotheses + Option 1) — economic interpretation. - Appendix G (BI exploration) — full 11-variant table and per-variant diagnostics. - `PAC_BI_AU_EXPLORATION.md` — complete saga, 10 sections, including raw `.mat` / `.txt` artefacts per variant.

E-SAT auxiliaries (wp736 Tables 4.6.11–12; unchanged in Phase L2): two separate equations — output-gap auxiliary and user-cost auxiliary:

$$\text{Eq. (29): } \hat{ib}_t = 0.59\hat{ib}_{t-1} + 0.15\hat{y}_{t-1} + 0.04\tilde{\pi}_{t-1} - 0.02\tilde{u}_{t-1}$$

Eq. (30): $\widehat{rKB}_t = 0.55\widehat{rKB}_{t-1} + 0.24\tilde{i}_{t-1}$

These continue to drive the BI block's PAC expectation through the E-SAT VAR; the only change is in the BI equation's own deep parameters ($\beta_0 \dots \beta_3$, ω) which are now wp1044-imported.

Figure 4.6.1: Historical dynamic-contribution decomposition of $\Delta \log i_b$ (business investment growth) over 1994Q3–2024Q4, generated under the PRE-hybrid AU-estimated calibration. The accelerator $b3_ib.yhat_au$ (red) and the user-cost channel $-\sigma.pv_rKB_aux$ (brown) are the two most variable structural channels. Under the post-hybrid calibration ($b0_ib$ raised $5\times$ from 0.018 to 0.096; $b1_ib$ raised $4\times$ from 0.082 to 0.33; $b2_ib$ activated at 0.11; $b3_ib$ raised $2\times$ from 0.31 to 0.69), the ECM and AR contributions are larger and the residual ϵps_ib contribution is smaller. The re-rendered figure is pending the next MATLAB GUI session per `WORKING_PAPER_BLOCKERS.md`; the qualitative pattern (PAC expectation small but timing-relevant at 2008-09 and 2020 turning points) is preserved by the wp1044 calibration since the BI structural form is unchanged.

4.7 Housing investment (FR-BDF Section 4.6.3, wp1044 §3.5.2 Eq 37) — Phase L2 wp1044-faithful refresh with price-spread term active

4.7.1 Target equation (FR-BDF eq 66) The household investment target depends on permanent income, the mortgage rate gap, and housing price Tobin's Q (eq. 31):

$$\Delta \ln I_t^{H*,bar} = \kappa_{ih,inc}(pv_{yh,t} - pv_{yh,t-1}) - \kappa_{mort}(i_t^{LH} - i_{SS}^{LH}) + \kappa_{ph} \cdot ph_gap_{t-1}$$

Table 4.7.1: Household investment target coefficients

Parameter	Symbol	Value	Description
Target persistence	rho_ih_star	0.95	Calibrated
Permanent income	kappa_ih_inc	0.03	wp736 eq 66
Mortgage rate gap	kappa_mort	0.048	Calibrated (posterior)
Housing Tobin's Q	kappa_ph	0.03	Calibrated

4.7.2 Short-run PAC equation (2nd-order, FR-BDF eq 67) Eq. (32):

$$\Delta \ln I_t^{H,level} = b_0^{ih}(\cdot) + b_1^{ih}\Delta_{t-1} + b_2^{ih}\Delta_{t-2} + pac_exp(pac_ih) + b_3^{ih}\hat{y}_t + \varepsilon_t^{ih}$$

The direct interest rate term $b_4^{ih}\tilde{i}_{t-1}$ is dropped, as F-test diagnostics showed it statistically insignificant (F=0.001, $\Delta SSR=0.005$ on T=118). The interest-rate channel enters household investment through two structural paths: (i) the

target equation via $\kappa_{mort}(i^{LH} - SS)$ through `pac_expectation`, and (ii) the auxiliary equation $pv_{ih,aux}$ with $a_{ih,i} = -0.15$. The direct ad hoc term would have triple-counted this channel.

Table 4.7.2: Housing investment short-run PAC coefficients — Phase L2 wp1044-faithful iterative OLS vs FR-BDF

Parameter	wp1044 FR Table 3.5.7	AU L2		
		iterative OLS (<code>results_housing_investment</code> , N=70, 14 iters)	Bayesian posterior (<code>invariant</code> hybrid MCMC, 20k×2)	90% HPD
β_0 (EC speed on <code>I*_H/I_H</code>)	0.12	0.496 (SE 0.085, t=5.86)	0.0290	[0.0080, 0.0493]
β_1 ($\Delta \log I_H$ lag 1)	0.18	0.293 (SE 0.129, t=2.28)	0.1074	[0.0337, 0.1815]
β_3 (output gap accelerator)	—	—	0.2104	[0.0409, 0.3632]
β_2 (contemp $\Delta y - \hat{y}$)	0.50	−0.073 (SE 0.247, t=−0.30)	—	—
Price spread (<code>p_SH - p_IH</code>) _{{-1} - {-5}}	0.05	−3.22 (SE 4.51, t=−0.71)	—	—
ω (calibrated per wp1044)	0.05	0.05	—	—
χ (derived from depth-1 polynomial)	~0.18	0.275	—	—
R²	0.89	0.43	—	—
stderr <code>eps_ih</code>	—	—	1.8938	[0.5232, 3.3645]
Significant COVID dummies	—	d_20Q2=−6.80 (t=−3.18)	—	—

The Phase L2 housing-inv block partially replicates wp1044 Eq 37 on AU data. Three findings:

- $\beta_0 = 0.50$ is $4.6\times$ wp1044's 0.12 , consistent with the cross-block AU-faster-ECM pattern (§5.1). $\beta_1 \approx$ wp1044's β_1 (0.29 vs 0.18, within $2\times$ and structurally signed).
- β_2 **wrong sign on contemp** ($\Delta y - \hat{y}$) with $t = -0.30$ (statistically insignificant). wp1044's $\beta_2 = 0.50$ is one of their largest coefficients; the AU result is consistent with the hypothesis that AU housing-inv responds more to interest rates / Tobin's Q than to contemporaneous output growth gap.
- β_3 **price-spread term active for the first time in Phase L2**. The earlier-AUSPAC OLS lacked pSH and pIH as separate observables; Phase L2 P1b constructed them from ABS 6416 Residential Property Price Index (pSH proxy) and ABS 5206 Implicit Price Deflator for dwelling investment (pIH proxy), shortening the sample to N=70 quarters. The estimated $\beta_3 = -3.22$ is wrong-signed and insignificant ($t=-0.71$); plausibly the AU sample is too short and the housing-price spread too noisy on AU data to identify the wp1044 coefficient. R^2 gap vs wp1044 (0.43 vs 0.89) is largest of the four “fitting” blocks; primary remaining cause is the short price-spread sample plus the missing wp1044 EA-foreign block.

The **b_ph_ih** Tobin's Q channel (IV F = 432, T = 115; spliced 1959Q3-onward AU house-price series) is preserved as a separate identification — it sits inside the housing-inv long-run target (Eq 31) rather than the short-run PAC equation, so it is not directly comparable to wp1044's β_3 price-spread term, which is a short-run contemporaneous regressor inside the SR PAC equation. The two channels coexist in the production model. | SSR | — | 957.2 | — | — |

E-SAT auxiliary (wp736 Table 4.6.16, eq. 33): $\hat{ih}_t = 0.71\hat{ih}_{t-1} + 0.08\hat{y}_{t-1} - 0.08\tilde{i}_{t-1} + 0.05\tilde{\pi}_{t-1} - 0.03\tilde{u}_{t-1}$

The interest-rate channel through the auxiliary ($a_{ih,i} = -0.08$) provides the structural rate transmission. h-vector amplification is $1.90\times$, the second largest of the five PAC equations.

Housing-price gap identification. The ABS 6416 Residential Property Price Index begins in 2003Q3 (T=72 quarters), which proved too short to identify the housing-price-gap coefficient: both OLS and lag-2-ph_gap IV produced wrong-signed estimates on the 2003+ sample, reflecting reverse-causality bias from supply-side shocks. We extend the housing-price series back to 1959Q3 by chaining growth rates from a long-history Australian house-price series (`data/house_price_history_long.csv`) onto ABS 6416 levels at the 2003Q3 overlap (overlap-growth correlation 0.84, RMS difference 1.16 pp/qtr). On the resulting T=115 sample (1993Q1–2021Q4, binding constraint set by the start of housing-investment data), the lag-2 IV produces a strong first stage (F=432) and a positive sign for both OLS (+0.066) and IV (+0.0099, s.e. 0.075). The IV magnitude is statistically indistinguishable from zero —

the housing-price-gap channel in AU is small. The structural rate transmission to housing investment continues to operate through `pv_ih_aux` and the `pac_expectation` mortgage-rate term.

*Figure 4.7.1: Historical dynamic-contribution decomposition of $\Delta \log i_h$ (housing investment growth) over 1994Q3–2024Q4. The error correction term $b0_ih \cdot (ih_hat - \ln_ih)$ (blue) and the house-price gap $b_ph_ih \cdot ph_gap$ (brown) are the two largest channels, and they almost exactly offset: the smoothed **ih_hat** target is persistently above actual $\ln_ih$ (positive EC pushing investment up), while the AU house-price gap has been positive on average (positive ph_gap pulling target investment up via Tobin’s Q channel, which is already absorbed into ih_hat — so the residual ph_gap term here is consistently negative). Net, the structural channels deliver a small positive impulse balanced by AR persistence and the shock eps_ih .*

4.8 External trade (FR-BDF Section 4.7)

External trade is specified as a proper error-correction model: a long-run cointegration relationship between log volumes, log demand, and the real exchange rate, together with short-run dynamics that adjust toward that equilibrium. The earlier version of AU-PAC used a degenerate accumulator $m_gap = m_gap(-1) - \Delta \ln M$ (and symmetrically for exports) without an economic long-run target, which meant the trade block could not capture the secular rise in AU import/export openness ($\log(M/C)$ trended at +2.5% per year and $\log(X/C)$ at +2.0% over 1960–2024). The proper specification below restores the FR-BDF Section 4.7 structure with explicit long-run income elasticities (β_m, β_x) and real-exchange-rate elasticities (γ_m, γ_x).

4.8.1 Exports (FR-BDF eqs 70–73) Long-run equilibrium (eq. 71):

$$\ln X_t^{eq} = \beta_x \cdot \hat{y}_t^{US} + \gamma_x \cdot s_gap_t$$

Error-correction gap (target minus actual):

$$x_gap_t = \ln X_t^{eq} - \ln X_t^{level}, \quad \ln X_t^{level} = \ln X_{t-1}^{level} + \Delta \ln X_t$$

Short-run dynamics (eq. 73):

$$\Delta \ln X_t = b_0^x \cdot x_gap_{t-1} + b_1^x \Delta \ln X_{t-1} + b_2^x \hat{y}_t^{US} + b_3^x \cdot s_gap_t + b_4^x \Delta \ln P_t^{com} + \varepsilon_t^x$$

Table 4.8.1: Export equation coefficients

Parameter	Value	Description
beta_x (LR foreign income elasticity)	1.20	Long-run pull from world demand
gamma_x (LR RER elasticity)	+0.40	Marshall-Lerner: depreciation raises target
b0_x (EC speed)	0.05	Adjustment toward equilibrium
b1_x (persistence)	0.30	AR(1)
b2_x (SR world demand)	0.25	Impact response to US output gap
b3_x (SR exchange rate)	0.10	Impact competitiveness
b4_x (commodities)	0.15	AU-specific commodity-price channel

4.8.2 Imports (FR-BDF eqs 74-77) Imports adjust toward a long-run equilibrium that depends on the level of import-weighted demand (cumulated IAD) and the real exchange rate. With $\beta_m > 1$, the income elasticity exceeds unity, which is the mechanism by which the model captures the structural rise in AU import openness without an ad-hoc deterministic trend (cf. the MARTIN approach of adding $+\alpha \cdot \tau$ to the short-run equation, which is BGP-violating).

Long-run equilibrium (eq. 76):

$$\ln M_t^{eq} = \beta_m \cdot \ln D_t^{iad} + \gamma_m \cdot s_gap_t$$

Error-correction gap:

$$m_gap_t = \ln M_t^{eq} - \ln M_t^{level}, \quad \ln M_t^{level} = \ln M_{t-1}^{level} + \Delta \ln M_t$$

The import-weighted demand level cumulates the contemporaneous IAD growth:

$$\ln D_t^{iad} = \ln D_{t-1}^{iad} + iad_t, \quad iad_t = \sum_j w_j^{iad} \cdot \Delta \ln j_t$$

Short-run dynamics (eq. 77):

$$\Delta \ln M_t = b_0^m \cdot m_gap_{t-1} + b_1^m \Delta \ln M_{t-1} + b_2^m \cdot iad_t + b_3^m \cdot s_gap_t + \varepsilon_t^m$$

Table 4.8.2: Import equation coefficients

Parameter	Value	Description
beta_m (LR income elasticity)	1.50	>1 captures rising AU import openness
gamma_m (LR RER elasticity)	-0.40	Depreciation lowers import target
b0_m (EC speed)	0.06	Adjustment toward equilibrium
b1_m (persistence)	0.23	AR(1)
b2_m (SR IAD elasticity)	0.36	Impact response to import-weighted demand growth
b3_m (SR exchange rate)	-0.08	SR Marshall-Lerner

Table 4.8.3: IAD weights (import content of demand)

Component	Weight	Source
Consumption	0.12	ABS input-output tables
Business investment	0.25	ABS input-output tables
Housing investment	0.15	ABS input-output tables
Government	0.08	ABS input-output tables
Exports (re-export)	0.30	ABS input-output tables

Trade volumes use the ABS 5206 Seasonally-Adjusted chain volume measures (series A2304114F / A2304115J), aligned to a T=126 quarter sample (1993Q2–2024Q4) with COVID pulse dummies for 2020Q2 and 2020Q3 absorbing the trade-collapse-and-rebound outliers.

The long-run elasticities (β_m , β_x , γ_m , γ_x) are calibrated at sensible AU values and given Normal priors centred there (§6.4, Table 5.6). A bivariate Engle-Granger style OLS of log volumes on log consumption (1959Q3–2025Q4) gives $\beta_m \approx +1.73$ ($t = +113$) and $\beta_x \approx +1.56$ ($t = +119$); the model priors at 1.50 and 1.20 sit within one standard deviation. The short-run coefficient $b1_m$ is set to its OLS estimate on the SA sample (+0.232, s.e. 0.086, $t=2.71$); $b1_x$, $b2_x$, and $b2_m$ are kept at calibrated values pending Asian-PMI / commodity-demand proxies that would help identify the SR foreign-demand channel for AU exports. The COVID-period dummies are sizeable for both exports (crash -10.14, bounce -7.11; tourism and education exports did not recover in 2020Q3) and imports (crash -14.39, bounce +6.27). Adding $\Delta \ln M$ and $\Delta \ln X$ to the set of Bayesian observables (Table 5.1) gives the long-run elasticities their own data identification — previously the trade block was econometrically inert because no observable mapped to it.

4.9 Demand deflators (FR-BDF Section 4.7) — AU reduced-form simplification

AUSPAC’s six demand deflators (consumption, business investment, household investment, exports, imports, and government consumption) are modelled as **single-equation reduced forms** combining own-lag persistence with contemporaneous cost-push regressors and a growth-neutrality anchor on $\bar{\pi}$. General form (eq. 36):

$$\pi_t^j = \rho_j \pi_{t-1}^j + \alpha_j \pi_t^Q + \beta_{j,m} \pi_t^M + \dots + (1 - \rho_j - \alpha_j - \beta_{j,m} - \dots) \bar{\pi}_t + \varepsilon_t^j$$

Deviation from FR-BDF wp736 / wp1044: FR-BDF specifies each deflator as a **two-equation error-correction model** — an explicit long-run target $p^{j,*}$ (an IAD-weighted combination of p^Q , p^{MO} (non-energy import), p^{MNRJ} (energy import), plus block-specific trend or competitor terms) and a short-run equation that includes an error-correction term $\beta_{j,EC}[p_{t-1}^j - p_{t-1}^{j,*}]$ pulling each deflator back toward its target. AUSPAC carries no $p^{j,*}$ target variables and no ECM term in any of its six deflator equations; the long-run anchoring runs instead through the growth-neutrality constraint on $\bar{\pi}$. The two main consequences:

- **No energy / non-energy import split** in the deflator block — AUSPAC’s π^M is a single aggregate import-price series and the commodity channel enters separately as an exogenous AR(1) `dl_n_pcom` rather than as π^{MNRJ} inside the deflator targets. The energy-import passthrough that FR-BDF wp1044 §3.6 emphasises (especially the post-2022 surge) is therefore captured imperfectly through `dl_n_pcom + gamma_oil` rather than through a dedicated p^{MNRJ} regressor.
- **CPI long-run anchoring is recovered inside the Phillips curve** (§4.4.0 above) via the $b_{ECM,pc}(p_{t-1}^{C,*} - p_{t-1}^C)$ ECM term rather than inside `eq_pi_c`. This keeps the cumulative CPI trajectory pinned to a weighted VA+import target without re-introducing the full FR-BDF target/ECM structure across all six deflators.

The reduced-form simplification was a deliberate choice for AU estimation tractability — six additional latent target processes would have increased the estimation surface substantially with limited identification gain on the AU sample (1993Q1–2024Q4, 128 quarters). The FR-BDF wp1044 §3.6 architecture is preserved as a reference for future restoration if AU data quality or sample length permits joint identification of all 12 deflator equations + targets.

The full set of AU-estimated reduced-form coefficients:

Table 4.9.1: Consumption deflator (FR-BDF eqs 79-80)

Parameter	Value	Description
<code>rho_pc</code>	0.40	Persistence

Parameter	Value	Description
alpha_pc	0.30	VA price pass-through
beta_pc_m	0.10	Import price
gamma_oil	0.03	Commodity/energy
Anchor (1-sum)	0.17	Growth neutrality

Table 4.9.2: Business investment deflator

Parameter	Value	Description
rho_pib	0.35	Persistence
alpha_pib	0.25	VA price pass-through
beta_pib_m	0.12	Import price (high import content)
Anchor (1-sum)	0.28	Growth neutrality

Table 4.9.3: Housing investment deflator

Parameter	Value	Description
rho_pih	0.45	Persistence (sticky construction)
alpha_pih	0.25	VA price pass-through
beta_pih_m	0.08	Import price (limited)
Anchor (1-sum)	0.22	Growth neutrality

Table 4.9.4: Export deflator

Parameter	Value	Description
rho_px	0.30	Persistence
alpha_px	0.20	VA price pass-through
beta_px	-0.05	Exchange rate (world price taker)
alpha_pcom	0.10	Commodity price (AU-specific)
Anchor (1-sum)	0.45	Growth neutrality

Table 4.9.5: Import deflator

Parameter	Value	Description
rho_pm	0.30	Persistence
alpha_pm	0.15	VA price (weak domestic)
beta_pm	0.08	Exchange rate (strong FX pass-through)
beta_pm_com	0.05	Commodity price
Anchor (1-sum)	0.42	Growth neutrality

Parameter	Value	Description
-----------	-------	-------------

Table 4.9.6: Government deflator

Parameter	Value	Description
rho_pg	0.50	Persistence
alpha_pg	0.30	Public sector wages (pi_w - dln_prod, not piQ)
Anchor (1-sum)	0.20	Growth neutrality

Growth neutrality verification: At SS, $\pi^j = \pi^Q = \bar{\pi} = \pi_{SS}$ for all deflators. The sum of all coefficients on inflation-type terms equals 1 for each. Verified for all 6 deflators.

4.9.7 HICP-style headline-decomposition reporting block (Round 1.1, 2026-05-18)

Following ECB-BASE §3.2.3 and FR-BDF wp1044 §3.6.4, AUSPAC v3.1.1 adds a six-variable headline-CPI decomposition reporting layer. The decomposition is purely *one-way*: each new variable is a deterministic function of existing model objects, with **zero feedback** into the structural dynamics, the wage Phillips curve, or the Taylor rule. Adding the block leaves both the cached Bayesian posterior and the IRFs of all existing variables bit-identical (Laplace LMD = -779.30 before and after; MHM LMD = -780.47 before and after).

Components and equations. Three new core/food/energy split variables and three new tradeables/non-tradeables/trimmed-mean variables:

$$\begin{aligned}
\pi_t^{food} &= \delta_{food,Q} \pi_t^Q + (1 - \delta_{food,Q}) \bar{\pi}_t + \delta_{food,com} \Delta \ln p_t^{com} \\
\pi_t^{energy} &= \delta_{energy,M} \pi_t^M + (1 - \delta_{energy,M}) \bar{\pi}_t + \delta_{energy,com} \Delta \ln p_t^{com} \\
\pi_t^{core} &= \frac{\pi_t^{au} - w^{food} \pi_t^{food} - w^{energy} \pi_t^{energy}}{1 - w^{food} - w^{energy}} \\
\pi_t^{trad} &= \delta_{trad,M} \pi_t^M + (1 - \delta_{trad,M}) \bar{\pi}_t + \delta_{trad,com} \Delta \ln p_t^{com} + \delta_{trad,s} s_t^{gap} \\
\pi_t^{nontrad} &= \frac{\pi_t^{au} - w^{trad} \pi_t^{trad}}{1 - w^{trad}} \\
\pi_t^{trim} &= \rho_{trim} \pi_{t-1}^{trim} + (1 - \rho_{trim}) [\delta_{trim,Q} \pi_t^Q + (1 - \delta_{trim,Q}) \bar{\pi}_t]
\end{aligned}$$

The core and non-tradeables components are residual identities that close the decompositions $\pi^{au} \equiv w^{food} \pi^{food} + w^{energy} \pi^{energy} + (1 - w^{food} - w^{energy}) \pi^{core}$ and $\pi^{au} \equiv w^{trad} \pi^{trad} + (1 - w^{trad}) \pi^{nontrad}$. Adding-up is exact to machine precision (verified at -1.7×10^{-16}).

Calibration. Weights from ABS Cat. 6401.0 (CPI 2025 reweight) and RBA Bulletin tradeables share. Projection coefficients (δ) calibrated to typical AU passthrough magnitudes.

Parameter	Value	Source
w^{food}	0.17	ABS 6401 food + non-alcoholic beverages
w^{energy}	0.07	ABS 6401 automotive fuel + electricity + gas
w^{trad}	0.35	RBA Bulletin tradeables share
$\delta_{food,Q}$	0.65	Domestic VA dominance in food (meat, dairy, fresh produce)
$\delta_{food,com}$	0.20	Global agricultural commodity passthrough
$\delta_{energy,M}$	0.50	Imported-fuel share
$\delta_{energy,com}$	0.60	Oil/gas global passthrough
$\delta_{trad,M}$	0.85	Tradeables dominated by imports
$\delta_{trad,com}$	0.15	Commodity passthrough
$\delta_{trad,s}$	-0.10	FX passthrough (AUD $\uparrow \rightarrow$ tradeables $\downarrow$)
ρ_{trim}	0.85	Trimmed-mean persistence
$\delta_{trim,Q}$	0.70	Trimmed mean tracks underlying VA-price trend

IRF validation. Shock responses on impact (post-Round-1.1 IRFs, $t = 1$):

- $\pi^{food} / \varepsilon^{pcom}$: +0.600 — matches $\delta_{food,com} \sigma_{pcom} = 0.20 \times 3.0$.
- $\pi^{energy} / \varepsilon^{pcom}$: +2.431 — matches $\delta_{energy,com} \sigma_{pcom} + \delta_{energy,M} (\beta_{pm,com} \sigma_{pcom}) = 1.80 + 0.50 \times 1.26$.
- $\pi^{core} / \varepsilon^{pi}$: +0.641 — matches the residual identity $\pi_{\varepsilon_{pi}}^{au} / (1 - w^{food} - w^{energy}) = 0.487/0.76$.

Use cases. The decomposition enables direct comparison with RBA inflation

reporting (trimmed-mean, weighted-median, tradeables/non-tradeables), richer post-2022 inflation-surge attribution, and a foundation for the energy-block routing in Round 2.1 (oil + gas split, FR-BDF wp1044 Appx E), which will replace the calibrated $\delta_{energy,com}$ passthrough with a structural synthetic energy price index.

4.10 Financial variables (FR-BDF Section 4.8)

4.10.1 Term structure (FR-BDF eqs 95-97) The 10-year yield is determined by the present value of expected future short rates plus a term premium (eq. 37):

$$pv_{i,t} = (1 - \kappa_{10})i_t + \kappa_{10} \cdot pv_{i,t+1}$$

$$i_t^{10Y} = pv_{i,t} + tp_t + \varepsilon_t^{10Y}$$

with $\kappa_{10} = 0.97$ (effective duration ~10 years) and $tp_t = \rho_{tp}tp_{t-1} + (1 - \rho_{tp})tp_{SS} + \varepsilon_t^{tp}$.

SS: $i_{SS}^{10Y} = 1.049 + 0.30 = 1.349\%$ quarterly (~5.4% annual).

4.10.2 WACC decomposition (FR-BDF eqs 98-100) The weighted average cost of capital decomposes into three funding sources (eq. 38):

$$WACC = 0.50 \times i^{COE} + 0.30 \times i^{LB} + 0.20 \times i^{BBB}$$

Table 4.10.1: WACC composition

Component	Weight	Spread SS	Spread rho	Annual rate at SS
Cost of equity (i_COE)	0.50	0.80%	0.92	~8.6%
Bank lending (i_LB_firms)	0.30	0.25%	0.77	~6.4%
BBB bonds (i_BBB)	0.20	0.05%	0.94	~5.6%
WACC	1.00	—	—	~7.3%

Each rate = $i^{10Y} + \text{spread}$, where spreads follow AR(1) with shock.

4.10.3 Exchange rate (FR-BDF eq 105) Modified UIP with persistent deviations from PPP and a forward-looking NPV of the policy-rate gap:

$$pv_t^{uip} = (i_{au,t} - \bar{i}_t) + \beta_{uip} \cdot E_t[pv_{t+1}^{uip}] \quad (\text{Hybrid / MCE})$$

$$pv_t^{uip} = (i_{au,t} - \bar{i}_t) + \beta_{uip} \cdot pv_{t-1}^{uip} \quad (\text{VAR regime: backward AR(1)})$$

$$s_gap_t = \rho_s \cdot s_gap_{t-1} - \alpha_s \cdot pv_t^{uip} + \alpha_s (\tilde{\pi}_t - \tilde{\pi}_t^{US}) + \varepsilon_t^s$$

with $\rho_s = 0.775$, $\alpha_s = 0.15$, and $\beta_{uip} = 0.92$. The standard NPV form (coefficient 1 on the rate gap, β on the continuation) means `pv_i_uip` is *not* dampened on impact: under the forward recursion and a Taylor-rule persistence $\lambda_i \approx 0.85$, the impact value is $pv^{uip}(0) \approx i_{gap}(0)/(1 - \beta_{uip}\lambda_i) \approx 4.55 \cdot i_{gap}(0)$, so the spot AUD internalises about 4.5× the contemporaneous policy-rate gap. Under the VAR regime the recursion is backward AR(1) — `pv_i_uip` accumulates only as past `i_gap` realisations roll in, and the peak AUD appreciation is delayed by 10–15 quarters relative to Hybrid/MCE. This forward-looking specification is what generates the FR-BDF-style Hybrid > VAR amplification documented in §7.2.

The sign convention is $s_gap > 0$ = AUD depreciation; higher AU rates appreciate the AUD. Steady state: $pv_{SS}^{uip} = 0$, $s_gap_{SS} = 0$ (PPP holds in the long run).

This replaces the pre-Phase-Q specification $s_gap = \rho_s \cdot s_gap(-1) - \alpha_s \cdot i + \alpha_s \cdot (\tilde{\pi} - \tilde{\pi}^{US}) + \varepsilon^s$, which used the contemporaneous policy-rate gap directly and therefore did not differentiate impact responses across expectation regimes. The forward-NPV refactor is the only structural change required to deliver the FR-BDF Hybrid-amplification pattern; all other PAC blocks remain backward-looking under Hybrid by design.

4.10.4 Household bank lending rate (FR-BDF eq 68) The mortgage rate adjusts sluggishly to changes in the 10Y bond rate (eq. 40):

$$i_t^{LH} = \rho_{LH} i_{t-1}^{LH} + (1 - \rho_{LH})(i_t^{10Y} + spread_{LH}) + \varepsilon_t^{LH}$$

with $\rho_{LH} = 0.97$ (AU OLS — standard variable mortgage rates follow the cash rate with substantial lag because of fixed-rate and discount carve-outs in the mortgage book) and $spread_{LH} = 0.40\%$ quarterly (~1.6% annual).

4.10.5 Housing prices (FR-BDF eq 69) Real housing price growth follows an AR(1) with demand and credit channels (eq. 41):

$$\Delta \ln P_t^H = \rho_{ph} \Delta \ln P_{t-1}^H + \alpha_{ph,y} \hat{y}_t + \alpha_{ph,r} \tilde{i}_{t-1} + \varepsilon_t^{ph}$$

with $\rho_{ph} = 0.90$, $\alpha_{ph,y} = 0.15$, $\alpha_{ph,r} = -0.10$.

4.11 Government and GDP identity (FR-BDF Sections 4.9-4.10)

Fiscal rule (eq. 42)

$$\Delta \ln G_t = \rho_g \Delta \ln G_{t-1} + \phi_g \hat{y}_t + \varepsilon_t^g$$

Countercyclical: $\phi_g = -0.10$ means a positive output gap reduces government spending growth.

GDP expenditure identity (eq. 43)

$$\hat{y}_t^{dom} = 0.55 \Delta \ln C + 0.13 \Delta \ln I^B + 0.06 \Delta \ln I^H + 0.24 \Delta \ln G + 0.25 \Delta \ln X - 0.23 \Delta \ln M$$

Table 4.11.1: GDP expenditure weights (ABS 2023)

Component	Weight	Description
Consumption	0.55	Household final consumption
Business investment	0.13	Non-dwelling GFCF
Housing investment	0.06	Dwelling GFCF
Government	0.24	Government final consumption
Exports	0.25	Goods and services exports
Imports	-0.23	Goods and services imports

Bridge equation (eq. 44) The demand-side aggregate feeds back into the IS curve: $\hat{y}_t = \dots + \lambda_{dom} \hat{y}_t^{dom} + \varepsilon_t^q$ with $\lambda_{dom} = 0.399$ (Bayesian posterior), closing the Keynesian multiplier loop.

4.11.1 Rounds 4–8 model extensions (2026-05-20)

Six concurrent extensions added to the production model, calibrated to AU-relevant central values, none of them entering the 9-observable likelihood directly (so the Bayesian estimation surface is unchanged):

- **Round 4 — Foreign monetary policy.** Adds an explicit US policy-rate process to the previously-only-output/inflation US block. $\bar{i}_t^{us}$ follows an AR(1) toward a steady-state neutral rate $i_{ss}^{us} = 0.625$ qpp; i_t^{us} closes a simple Taylor rule on US output and inflation gaps: $i_t^{us} = \bar{i}_t^{us} + \alpha_i^{us} \pi_{t-1}^{us,gap} + \beta_i^{us} \hat{y}_{t-1}^{us} + \varepsilon_t^{i,us}$ with $\alpha_i^{us} = \beta_i^{us} = 0.5$. The variable can be used as a UIP-side conditioning observable in future identification work.
- **Round 5 — Demographic trend.** Adds a slowly-varying gap variable $\Delta \ln \overline{POP}_t$ (AR(1) persistence 0.95) that shifts long-run employment growth one-for-one through $\Delta \ln N_t^* = \dots + \Delta \ln \overline{POP}_t$. Steady-state population growth is implicit in the gap-formulation demeaning (≈ 0.375 qpp from ABS Cat 6202.0); deviations are absorbed by the new shock $\varepsilon_t^{pop,bar}$.

- **Round 6 — Tax structure decomposition.** Three AR(1) effective-rate gap variables — $\tau_t^{GST,gap}$, $\tau_t^{PAYG,gap}$, $\tau_t^{CIT,gap}$ — each feeding one structural channel:
 - GST pass-through to the consumer-price deflator: $\pi_t^c = \dots + \alpha_{GST}\tau_t^{GST,gap}$ with $\alpha_{GST} = 0.05$.
 - PAYG drag on consumption growth via the income-tax change: $\frac{\Delta \ln C_t^*}{C_t^*} = \dots - \alpha_{PAYG}\Delta\tau_t^{PAYG,gap}$ with $\alpha_{PAYG} = 0.10$.
 - CIT bump to the user cost of capital: $UC_t^K = WACC_t + \delta_K - (\pi_t^{IB} - \pi_t^Q) + \alpha_{CIT}\tau_t^{CIT,gap}$ with $\alpha_{CIT} = 0.02$.
- **Round 7 — Market vs non-market VA decomposition.** Identity-preserving split of the output gap into market and non-market sectors (ABS Cat 5206 weight $w_{market} = 0.85$ for market: manufacturing, finance, retail, mining, construction; $1 - w_{market} = 0.15$ for non-market: public administration, education, health). Non-market output is smoother and lagged: $\hat{y}_t^{nm} = \rho_{nm}\hat{y}_{t-1}^{nm} + (1 - \rho_{nm})\gamma_{nm}\hat{y}_t$ with $\rho_{nm} = 0.90$, $\gamma_{nm} = 0.30$. Market output is the residual.
- **Round 8 — RBA-style auxiliary forecasters.** Three one-way nowcast projections that can be matched against monthly / high-frequency data outside the quarterly cycle:
 - BLR_t (Bank Lending Rate nowcast): AR(1) projection of the realised lending-rate gap $i_t^{LH} - i_{ss} - \tau p_{ss} - spread_{LH}$.
 - $MAPI_t$ (Mortgage Asset Price Indicator): AR(1) projection of ph_t^{gap} .
 - $MAPU_t$ (Mortgage Asset Price Underwriting): AR(1) projection of $\Delta \ln I_t^H$. Each carries an exogenous shock for direct nowcast injection ($\varepsilon^{BLR}, \varepsilon^{MAPI}, \varepsilon^{MAPU}$).

Variable / shock counts. The extensions add 11 endogenous variables, 18 calibrated parameters, and 9 exogenous shocks; production model goes from 164 to 175 endogenous variables and from 33 to 49 shocks. Smoke-tested 2026-05-20 — model preprocesses, solves under Blanchard-Kahn, and IRFs propagate through all the new channels. None of the new variables is in `varobs`, so the Bayesian likelihood and posterior are unchanged from the v3.1.1 baseline. See §4.11.2 for the rational-expectations consolidation that wired the Round 5–6 shocks into PAC expectation formulas on 2026-05-21.

4.11.2 Rounds 4–8 expectation consolidation (2026-05-21)

The Rounds 4–8 extensions of §4.11.1 added new structural equations and shocks to `model.inc`, but the per-block `var_model` companion matrices in `dynare/aux/` were not extended at the same time. As a result, between 2026-05-20 and 2026-05-21 the production model carried an internal inconsistency: a $\tau^{CIT,gap}$ shock would (correctly) raise the user cost of capital and depress business investment in the *structural* simulation block, yet the PAC expectation formula `pac_expectation(pac_ib)` — a closed-form linear combination

over the lagged augmented E-SAT state — projected forward as if the shock did not exist. Forward-looking agents in the PAC block did not anticipate the long-run effect of tax shocks on their target, so the expectation operator returned a downward-biased response to anything driven by the new shocks.

This subsection documents the targeted fix landed in PR #6 (commit 6811940). For each new shock with a *direct* structural channel into a PAC target, the corresponding aux file was extended with (a) the new variable as a `var_model` state in its orthogonal AR(1) reduced form and (b) a loading from that state into the PAC target’s auxiliary regression. `phaseW_recherrypick.m` was then re-run to regenerate the closed-form policy-function coefficients $h_{pac,X}$, which were patched into `au_pac.mod` and `au_pac_bayesian.mod` alongside the new aux-regression loadings. No `aggregate()` re-run was needed; all Phase U/V/W manual overrides are preserved.

Table 4.11.2.1: Channels wired in PR #6

Aux file	New <code>var_model</code> state	PAC target	Aux-regression loading	New $h_{pac,X}$ coefficient
<code>aux_employment.mod</code>	$\Delta \ln \overline{POP}_t$	$\hat{n}_t$	$a_{n,pop} = 1.0$ (one-for-one)	$h_{pac,n,pop} = +0.1240$
<code>aux_consumption.mod</code>	$\tau_t^{PAYG,gap}$	$\hat{c}_t$	$a_{c,PAYG} =$ $-\alpha_{PAYG} =$ -0.10	$h_{pac,c,PAYG} =$ -0.00956
<code>aux_business_inv.mod</code>	$\tau_t^{CIT,gap}$	$\widehat{ib}_t, \widehat{rKB}_t$	$a_{ib,CIT} =$ $-\sigma_{ces}\alpha_{CIT} =$ $-0.011,$ $a_{rKB,CIT} =$ $+\alpha_{CIT} =$ $+0.02$	$h_{pac,ib,CIT} =$ -2.41×10^{-4}
<code>aux_pQ.mod</code>	$\tau_t^{GST,gap}$	$\widehat{\pi Q}_t$	$a_{pQ,GST} =$ 0.05 (indirect via CPI $\rightarrow$ wages $\rightarrow$ VA)	$h_{pac,pQ,GST} =$ $+8.21 \times 10^{-4}$

The aux-regression loadings $a_{X,Y}$ are chosen to mirror the *structural* coefficient that the same shock carries in `model.inc` (e.g. α_{PAYG} in `eq_dln_c_star_bar` maps one-for-one to $-a_{c,PAYG}$); they are explicit, calibrated, and documented inline in each aux file. The cherrypicked $h_{pac,X,Y}$ coefficients in the right-most column are the closed-form policy-function projections obtained from `pac.print()` over the extended companion matrix. Sizes vary across blocks for two reasons: (i) the structural loading itself differs (1.0 vs 0.02), and (ii) the cherrypicked discounted-sum weights the new state by its forward trajectory

under the extended VAR companion, which depends on the new state’s AR(1) persistence ($\rho_{pop} = 0.95$ vs $\rho_\tau \in [0.90, 0.94]$) and the cumulative impact through the PAC ECM speed $b_{0,X}$.

Skipped channels. Three of the six Round 4–8 extensions are intentionally *not* wired into PAC expectations:

- **Round 4** ($i^{us}, \bar{i}^{us}$). The US policy-rate process is structurally isolated under the current model: no AU PAC equation references it (UIP runs through $\bar{i}$, not $\bar{i}^{us}$), so adding it to a `var_model` would produce zero loadings on every PAC target. The deeper gap — connecting US monetary policy to AU expectations via UIP — is a separate piece of architectural work flagged in Section 7.
- **Round 7** ($\hat{y}_t^{market}, \hat{y}_t^{nm}$). Both variables are identities in $\hat{y}_{au}$, which is already the leading state in every `var_model`. PAC expectations therefore project them automatically — no separate wiring is required.
- **Round 8** ($BLR_t, MAPI_t, MAPU_t$). One-way nowcast projections with no feedback into the rest of the model. Their shocks are observation-equation only; they cannot enter PAC expectations because the variables themselves cannot affect any PAC target by construction.

Likelihood preservation. A natural concern is whether the cached MCMC posterior (Laplace -779.30 / MHM -780.36 ; see §5) remains valid for the extended model. It does, by construction. The four new shocks ($\varepsilon^{pop,bar}, \varepsilon^{\tau,GST}, \varepsilon^{\tau,PAYG}, \varepsilon^{\tau,CIT}$) are not in `varobs`, and the corresponding state variables have zero realised values throughout the 1994Q1–2024Q4 sample. The new $h_{pac,X,Y}$ terms in each `pac_expectation_pac_X` identity are therefore exactly zero in every observed period, and the Kalman filter’s recursion is bit-identical to the pre-consolidation model. Reloading the cached chain under the consolidated model reproduces MHM $= -780.362$ to numerical precision. A fresh 20k MCMC under the consolidated model was nonetheless run as an empirical check; it returned **Laplace** $= -779.285$ (vs cached -779.30 ; $+0.015$ nats, mode bit-identical within numerical precision) and **MHM** $= -780.099$ (vs cached -780.36 ; $+0.26$ nats, within typical MHM stochastic noise on a $20k \times 2$ -chain estimate). All 28 estimated posteriors are within standard error of the cached values. The fresh chain is now the authoritative cache and the new mode file lives in `dynare/au_pac_bayesian/Output/au_pac_bayesian_mode.mat`.

Companion-matrix dimensionality. After the consolidation each aux file’s `var_model` state vector has been extended by one or two states beyond the pre-consolidation count of 13–14 per block. The *core* E-SAT companion underlying §3.2’s Table 3.3 remains 12×12 — the four new states are all *block-specific* additions that live only in the aux file where they have a structural channel. The references to “ 12×12 ” in §3.2, §4.12 (E-SAT stability eigenvalue check) and Appendix A.8 / B.4 continue to refer to the E-SAT core; the per-block PAC companions are listed in Table 4.11.2.2.

Table 4.11.2.2: Per-block companion-matrix sizes post-consolidation

Block	Pre-consolidation size	Post-consolidation size	Added state(s)
aux_pQ	14	15	$\tau^{GST,gap}$
aux_consumption	14	15	$\tau^{PAYG,gap}$
aux_business_inv	14	15	$\tau^{CIT,gap}$
aux_housing_inv	13	13	(none — no Round 4–8 channel)
aux_employment	13	14	$\Delta \ln \overline{POP}$

The aux-file state vector sizes are not the same as the underlying E-SAT core because each block additionally carries its own auxiliary regression target(s) — $\widehat{\pi Q}$ for the VA price block, $\widehat{y\hat{h}}, \hat{c}$ for consumption, $\widehat{ib}, \widehat{rKB}$ for business investment, $\widehat{ih}$ for housing investment, $\hat{n}$ for employment — plus, in the case of **aux_pQ**, the wage-gap state $\tilde{\pi}_{gap}^w$ added in Phase U.

4.11.3 Round 1.2 — hand-to-mouth wage+transfer income channel (2026-05-22)

Following FR-BDF wp1044 §3.5.1 eq 35, the consumption short-run PAC equation was extended with a contemporaneous response to wage-plus-transfer income relative to per-capita output:

$$\Delta \ln C_t^{PAC} = \dots + b_{HtM} \cdot (\widetilde{wt}_t^H - \hat{g}_t^{au})$$

where $\widetilde{wt}_t^H$ is the HP-filter gap of log of real household labour-plus-transfer income, constructed from ABS 5206 Table 20 (compensation of employees A2302915V, social assistance benefits A2302919C, household HFCE IPD A2303940R) by `data/prepare_household_income.m`. The series captures the GFC stimulus peak (+4.0% in 2008Q4) and JobKeeper (+3.2% in 2020Q3) as the most extreme cyclical episodes.

$\widetilde{wt}_t^H$ enters the per-block `var_model` of `aux_consumption.mod` as a new state with a reduced-form AR(1) (procyclical wages, countercyclical transfer support reflecting AU’s JobKeeper-style fiscal response, PAYG drag). The `phaseW_recherrypick` step regenerates the consumption PAC’s policy-function coefficient $h_{pac,c,wtH}$ as exactly zero, since $\hat{c}_t$ doesn’t load directly on $\widetilde{wt}_t^H$ in the target projection — the channel is purely contemporaneous, consistent with the rule-of-thumb interpretation.

Empirical result: the calibrated channel worsens fit by 5.7 nats. A fresh 20k×2 MCMC under the Round 1.2 model (`b_HtM` = 0.32 from the FR-

BDF wp1044 posterior) yields $\text{MHM} = -785.80$ vs the pre-Round-1.2 baseline of -780.10 . The other estimated parameters barely shift relative to the pre-Round-1.2 posterior (consumption-PAC coefficients move by less than 0.05), confirming that the channel as calibrated does not match AU consumption dynamics. The negative result is documented honestly: AU consumption is well-fit by the existing rate-channel and `yh_ratio` mechanism, and the additional FR-BDF-calibrated HtM channel adds noise rather than signal at this calibration point.

The structural plumbing (new `var_model` state, `eps_wtH` shock, data column, production.mod patches) is retained because two natural extensions can plausibly improve the fit:

1. **Promote b_{HtM} to `estimated_params`** with a prior $N(0.30, 0.10)$ centred on the FR-BDF posterior. The posterior may pull b_{HtM} towards zero, eliminating the -5.7 nat penalty while preserving the structural plumbing for future use.
2. **Add `au_wt_H_real_gap` to `varobs`** so the model uses the empirical series directly rather than treating it as a latent state predicted from `yhat_au`, `u_gap`, and `tau_PAYG_gap`. With the real series in the likelihood, the channel may rebalance.

Either follow-up would require a fresh MCMC under the expanded estimation surface. The Round 1.2 work is currently best understood as **structural infrastructure for future estimation**, not as a fit-improving change.

4.12 AU-PAC modelling choices and simplifications

This subsection documents the modelling choices that shape AU-PAC. The headline monetary-transmission channels — cost of capital, exchange rate, mortgage rate, expectations, and the wage-price spiral — are all implemented and estimated on AU data. Several wp736 features are simplified or omitted; each is flagged below.

Trends specification. Gap variables (`yhat_au`, `pi_au_gap`, etc.) are constructed *once* from the raw data by HP-filter detrending and fixed-anchor calibration for $\bar{i}_{ss}$, $\bar{\pi}_{ss,au}$, $\bar{\pi}_{ss,us}$ (matching the RBA and Fed inflation targets). This is data *transformation* — extracting cyclical components from observable levels — and is the AU-PAC analogue of wp736 §4.9 eqs 127–130, which performs the same trend/cycle decomposition using exponential-smoother trends. It is not data *smoothing* in the sense of the Kalman-smoothed-targets approach we evaluated in §5.3: the gap variables are observed inputs to the model, not unobserved targets reconstructed from a fitted model. Convergence to the balanced growth path is asserted via the linearised gap formulation; §7.1.1 validates BGP stability under off-SS impulses.

Sectoral net-financial-wealth dynamics. AU-PAC includes the four sectoral wealth gap variables (`w_F`, `w_G`, `w_H`, `w_N`) with calibrated steady-state

targets (Appendix A.6) and accumulation equations following wp736 eqs 116–126. Stabilisation transfer rules ($\tau_{TF}, \tau_{TN}, \tau_{TG}$ in wp736 §4.8.5 and §4.10) are not implemented. Long-run convergence of the four wealth ratios is validated empirically under off-SS perturbations — all four converge with 2–3 quarter half-lives — but no analytical proof of BGP convergence under arbitrary shocks is provided.

Conditional projection framework. AU-PAC’s conditional forecasting (§7.5) and pseudo-real-time recursive forecast (§6.5) take the US bloc, foreign demand, and commodity-price paths as external assumptions while letting domestic policy and demand respond endogenously. The exogenisable variables are `yhat_us`, `pi_us`, `dln_pcom`, and optionally `i_au` for RBA path-conditioning experiments. No analogue to the Eurosystem BMPE harmonisation (wp736 §4.11) is required.

Single aggregate import deflator. AU-PAC §4.9 uses a single aggregate import deflator with a commodity-price loading ($\beta_{pm,com} = 0.42$) rather than the separate energy / non-energy split of wp736 §4.7.5 (eqs 88–91). Australia is a net energy *exporter*, so cost-push from energy imports is less material than for France; the aggregate-deflator approach absorbs the channel implicitly via the commodity-price loading. A formal split is listed as remaining work in Section 7.

Institutional CPI-indexation channel. Wage indexation in AU is set institutionally through the Fair Work Commission’s annual wage review, which sets award rates and award-floor enterprise agreements covering a substantial fraction of the workforce (no direct analogue to wp736 §4.5.1 eq 51’s SMIC equation). AU-PAC’s wage Phillips curve §4.4.1 captures this explicitly through the consumer-price indexation term $\gamma_w \cdot \pi^c$ with $\gamma_w = 0.35$ (posterior mean; HPD [0.27, 0.43]) — a strong passthrough that absorbs most of the cyclical wage signal once the AU institutional setting is properly specified.

E-SAT stability check. AU-PAC reports the eigenvalue check empirically: the 12×12 AU companion has all eigenvalues with modulus ≤ 1 ($\max|\text{eig}| = 0.946$ from the smoother run), confirming stationarity. A formal characterisation along the lines of wp736 Appendix A.1 is left as future work.

4.13 AU adaptations vs FR-BDF design

This subsection collects, in one place, the principal departures of AU-PAC from the FR-BDF wp736 design. The departures are organised by the *kind* of adaptation: empirical findings revealed by AU data, intentional structural simplifications, adaptations to local market structure, calibration imports from wp736, fiscal-block differences, and methodological choices. Research extensions deferred for future work are summarised in `next_session.md` under “Remaining work”.

4.13.1 AU empirical findings These are substantive AU econometric results that emerge from the Bayesian estimation rather than design decisions.

AU flat Phillips curve in price and quantity blocks. The posteriors deliver small cyclical output-gap slopes in three separate blocks: VA price ($b2_pQ \approx 0$, Table 5.6), employment ($b5_n \approx 0$, §4.4.2), and consumption ($b3_c \approx 0$, §4.5.2). In the wage Phillips curve the unemployment-PV slope is identified with the correct (FR-BDF) sign but at a small magnitude ($\kappa_w = -0.10$, §4.4.1). The CPI-indexation channel $\gamma_w = 0.35$ in the wage equation, the strong CPI-passthrough loadings in the deflators, and the persistent forward-looking real-rate channel in consumption together absorb most of the cyclical wage and demand-block signal. This is consistent with the post-2015 RBA literature on AU wage non-responsiveness (Cassidy and Doyle 2018; Bishop and Cassidy 2017) and with the broader OECD finding that the Phillips curve has flattened.

AU trade-block long-run elasticities calibrated rather than estimated. The trade-block long-run elasticities $\beta_x = 1.56$, $\beta_m = 1.73$ are calibrated from bivariate cointegration regressions on 1959Q3–2025Q4 data; the short-run trade elasticities $b1_x$, $b2_x$, $b1_m$, $b2_m$ are jointly Bayesian-estimated with the rest of the model. Direct full-MCMC identification of the long-run elasticities with a US-output-gap proxy for foreign demand is fragile because the AU mining-boom cycle (2003–2014) is collinear with the US output-gap series at the long-run frequencies that identify β_x . The bivariate cointegration approach uses AU-specific commodity-price information and is the same identification strategy adopted by FR-BDF wp736 §4.7.1.

AU consumption rate-channel weakly identified from real-time data. The direct rate-change consumption coefficient $b_{di,c}$ identified via local-projection IV using lag-instrumented Δi_{au} produces wrong-signed estimates on AU data alone. The interpretation, supported by the RBA discussions of monetary-policy surprises (Beechey and Wright 2009; Bishop and Tulip 2017), is that endogenous-policy responses to demand contaminate Δi_{au} as an instrument and that high-frequency RBA OIS surprises (a clean exogenous policy-shock series) would be required to identify the channel cleanly. Pending availability of those data, $b_{di,c} = -0.70$ is Bayesian-regularised from the FR-BDF posterior; in the IRF decomposition it contributes only a small fraction of the consumption response, which is dominated by the b_2^c interest-rate-level channel and the forward-NPV real-lending-rate channel (§4.5.2).

4.13.2 AU structural simplifications These are intentional model omissions, each motivated by AU data limitations or a substantive simplification.

Permanent income proxy. The consumption target’s permanent-income variable pv_yh (eq. 22) discounts the output gap $\hat{y}_{au}$ rather than wp736’s PV of disposable income y^H (wp736 §4.6.1, eq 58). AU national accounts (ABS 5206) do not publish a quarterly real-disposable-income series spliced consistently back to 1994Q1 at the frequency and break-adjustment needed for the wp736 specifi-

cation; the output-gap proxy is the closest stationary AU analogue and inherits the same balanced-growth-path neutrality property. Quantitatively this matters chiefly in fiscal-multiplier experiments.

No new-housing-investment deflator p_IH . The household-investment target (eq. 31) omits the γ_1 new-housing-deflator channel from wp736 eq 66. ABS 5206 publishes only an aggregate dwelling-investment deflator (Cat. 8731 series), so the AU specification absorbs the new-housing-deflator channel into the aggregate π_{ih} term rather than separately identifying a new-vs-existing housing split.

Real housing user cost simplified to nominal rate gap. The mortgage-rate term in eq. 31 is written $-\kappa_{mort}(i^{LH} - i_{SS}^{LH})$ (nominal rate gap) rather than the $PV(\pi^Q)$ -deflated wp736 form. The forward-NPV real-lending-rate channel $pv_{r,lh,gap}$ (eq. 24 for consumption) provides a separate real-rate signal that the household-investment block can pick up through cross-equation correlations, so the simplification does not preclude real-rate transmission to housing.

Housing-price equation restructured. The AU housing-price equation `eq_ph_gap` augments a simple AR(1) with explicit demand-gap and credit-conditions channels, replacing the AR(2) inflation-anchor form of wp736 §4.6.4 (eq 70). The restructuring reflects the AU empirical finding that housing-price inflation is dominated by demand-side and credit-cycle dynamics rather than a backward inflation expectations process. The spliced 1959+ AU housing series (`data/house_price_history_long.csv`) supports the demand-and-credit channels with first-stage F-statistic 432.

Wage Phillips minimum-wage channel omitted. AU-PAC’s wage Phillips (eq. 16) omits the minimum-wage equation `eq_smic` from wp736 §4.5.1 (eq 51), which models the SMIC minimum wage as a function of CPI and median earnings. The AU Fair Work Commission’s award-rate decisions are well-summarised by the CPI passthrough term $\gamma_w \cdot \pi^c$ in the standard NK Phillips form, and a separate minimum-wage equation would add no degrees of freedom relative to that channel.

4.13.3 AU adaptations to local market structure Foreign block is US, not euro area. The wp736 foreign bloc is the euro area (relevant for France as a euro-area member with fixed nominal exchange rate). AU-PAC uses the US as the foreign bloc (relevant for AU through the AUD/USD float, US Treasury yield linkages, and US import-share weighting). Concretely, $\hat{y}_{us}$, π_{us} , and the US-derived term-premium decomposition replace the wp736 euro-area analogues.

Floating AUD with endogenous long-run real rate. Because AU operates an independent inflation-targeting RBA with a floating AUD, the long-run real rate is determined endogenously by the Taylor rule plus UIP rather than imported from a euro-area-wide steady state. The closed-form long-run real-rate $\bar{r} = \bar{i} - \bar{\pi}$ is recovered from the Taylor rule’s neutral-rate calibration.

Forward-looking UIP. The UIP equation (§4.10.3) uses the forward NPV $pv_{i,uip} = (i_{au} - \bar{i}) + \beta_{uip} \cdot pv_{i,uip}(+1)$ with $\beta_{uip} = 0.92$, replacing the wp736 contemporaneous-rate-gap specification. Under the AU calibration this delivers an impact AUD appreciation amplification factor of $1/(1 - \beta_{uip} \cdot \lambda_i) \approx 4.55$ at $\lambda_i = 0.85$, broadly consistent with the empirical UIP literature (Engel 2014; Bishop and Tulip 2017).

Commodity-price channel (audit §4.6.4, Stage 11b). AU’s mining-driven export base — iron ore, coal, LNG — makes commodity prices a first-order driver of trade volumes and the AUD. AU-PAC adds an explicit commodity-price-loading $b4_x \cdot pcom_gap$ in the export equation and a smaller loading in the import deflator ($\beta_{\{pm,com\}} = 0.42$). No analogue is needed in wp736 because French exports are dominated by manufacturing and services. The commodity-price shock eps_pcom is treated as exogenous in AU-PAC and tracked as a first-order AU-specific propagation channel in §7.3.4.

4.13.4 Calibration imports from FR-BDF The following parameters are imported from wp736 rather than estimated on AU data, either because AU data are insufficient (sample length, frequency, breaks) or because the parameter governs a sub-block whose own joint identification with the rest of the AU model proved fragile. These imports are documented here for transparency.

Block	Parameter	FR-BDF value	AU-PAC adoption rationale
WACC weights (§4.10.2)	w_COE = 0.50, w_LB = 0.30, w_BBB = 0.20	wp736 eq 100 weights	AU corporate-finance-mix data sparse pre-2000; FR-BDF weights match Australian Bureau of Statistics ABS 5232 corporate-funding share ranges

Block	Parameter	FR-BDF value	AU-PAC adoption rationale
Spread persistences (§4.10)	$\rho_{\text{COE}} = 0.92$, $\rho_{\text{LB}} = 0.77$, $\rho_{\text{BBB}} = 0.94$	wp736 §4.8.3	Spread-AR(1) coefficients estimated jointly with the rest of the financial block proved weakly identified on AU monthly data; calibration imports avoid this identification fragility
Sector wealth targets (§4.8)	$w_{\text{F}_{ss}} = -2.80$, $w_{\text{G}_{ss}} = -1.60$	wp736 §4.8.5 Table 4.8.1	ABS-derived AU sectoral net financial wealth ratios bracket these targets within ± 0.15 ; FR-BDF values used for cross-country consistency
Term-structure decay (§4.10.1)	$\kappa_{10} = 0.97$	wp736 eq 95	Direct estimation on AU 10Y/cash spread (1994–) gave 0.95–0.98 with wide HPD; FR-BDF 0.97 is well inside the AU posterior
Firms revaluation (§4.8)	$\gamma_{\text{reval}} = -0.018$	wp736 §4.8.5	AU corporate equity revaluation rates not separately published; FR-BDF value used

Block	Parameter	FR-BDF value	AU-PAC adoption rationale
Consumption surprise channel (§4.5.2)	$b_di_c = -0.70$	wp736 posterior	AU IV identification failed; pending RBA OIS surprises (Bishop–Tulip RDP 2017-08)
Real-rate consumption channel (§4.5.2)	$\alpha_c_r = -0.95$	wp736 Table 4.6.14	AU joint identification with $b2_c$ is fragile

In each case the imported parameter is held fixed during estimation; the equation-by-equation OLS identifies only block-specific parameters. The audit (§6) flags these as “imported, not AU-estimated” for transparency in interpreting cross-country comparisons.

4.13.5 AU government / fiscal differences No tax-rate $\times$ tax-base decomposition. The AU fiscal block treats government revenue as a single endogenous variable `tau_G_gap` rather than wp736’s decomposition into separate effective tax rates and tax bases for GST, PAYG, and company tax. ABS publishes aggregate Commonwealth revenue (Cat. 5512) but quarterly tax-rate-by-instrument series with consistent methodology back to 1994 are not available. Consequently AU-PAC cannot simulate tax-policy changes endogenously; users wishing to study tax-policy experiments can shock `tau_G_gap` directly with appropriate magnitudes drawn from ABS revenue data. A formal tax-instrument decomposition is on the research extension list (`next_session.md`).

Countercyclical fiscal rule. The AU fiscal rule `eq_tau_G_gap` responds to the output gap ($+0.05 \cdot \hat{y}_{au}$) rather than to the sectoral asset ratio that wp736 eq 125 uses for France. This change is driven by the institutional reality that Commonwealth-budget cycles in Australia are described as countercyclical with respect to output (Treasury PEFOs / MYEFOs explicitly target a deficit-tax-revenue elasticity), not net-asset-ratio-targeting in the FR-BDF style.

Stronger Taylor-rule output-gap response. The Taylor rule coefficient on the output gap is $\rho_{stab,2} = 0.25$ (vs wp736’s 0.10). The higher coefficient reflects a structural feature of the AU monetary-policy reaction function as estimated by Cusbert and Kendall (2018) and is well within the empirical range for inflation-targeting central banks; it also ensures rational-expectations determinacy under the floating AUD and forward-NPV UIP specification.

4.13.6 Methodological choices Endogenous structural trends. AU-PAC’s structural trends `dln_n_star_bar`, `dln_q_star`, `dln_y_star`, etc. are derived endogenously from the CES production function and the trend efficiency $\bar{E}$ process (§4.2.3 two-break specification). wp736 §4.9 advocates anchoring trends to external benchmarks (e.g. the European Commission’s potential output series). The AU choice reflects three considerations: (i) ABS does not publish a long-run potential-output series consistent with the gap-form CES specification used here; (ii) the AU mining boom and post-boom adjustment imposed large structural shifts (the 2002Q2 and 2008Q3 trend breaks) that no external benchmark currently captures; and (iii) the gap formulation makes the endogenous trend a smooth function of the data, ensuring numerical stability of the MCMC posterior.

Trend-cycle decomposition via HP-filter at data prep (audit §4.9, §4.12 above). Gap variables (`yhat_au`, `pi_au_gap`, etc.) are computed once from observable levels at the data-prep stage via HP-filter trends and fixed-anchor steady states (`i_ss`, `pi_ss`, `au`, `pi_ss`, `us`). This is the AU-PAC analogue of wp736 §4.9 eqs 127–130 (exponential-smoother trends). It is not data smoothing within the model: the gap variables are inputs to Bayesian estimation, not Kalman-smoothed unobserved states.

Standard (un-normalised) NPV form in `pv_i_uip`. The UIP NPV is written $pv_{i,uip} = (i_{au} - \bar{i}) + \beta_{uip} \cdot pv_{i,uip}(+1)$ rather than the normalised form $pv_{i,uip} = (1 - \beta_{uip})(i_{au} - \bar{i}) + \beta_{uip} \cdot pv_{i,uip}(+1)$. The standard form delivers the $\sim 4.55\times$ impact amplification at $\lambda_i = 0.85$, matching the empirical UIP literature; the normalised form would deliver only a unit-scaled NPV. The two forms differ only by a multiplicative constant absorbed into α_s in `eq_s_gap`.

$\alpha_{c,r} \leftrightarrow \sigma_2$ **parameterisation correspondence.** The consumption-block coefficient $\alpha_{c,r}$ (multiplier on `pv_r_1h_gap`) corresponds to the wp736 σ_2 elasticity (eq 61); the two differ in normalisation ($\alpha_{c,r}$ is the coefficient on the level NPV; σ_2 is the elasticity on the rate gap). Under the AU calibration $\alpha_{c,r} = -0.95$ maps to a wp736-equivalent σ_2 of approximately -0.30 once the normalisation difference is unwound.

Growth-correction omission in `pv_yh`. We omit the $/\exp(\Delta\bar{y})$ growth correction to the `pv_yh` recursion that wp736 §4.6.1 (footnote) discusses. At steady state $\Delta\bar{y} \approx 0$ in the AU gap formulation, so the correction vanishes; it would matter only in transient growth-rate experiments. The correction is omitted to preserve the linear gap structure required for Dynare’s order-1 IRF protocol.

5. Cross-block findings

This section consolidates the cross-block patterns that emerge from the Phase L2 wp1044-faithful partial-replication estimates of §4.3–§4.6. Each sub-section summarises a single finding visible only when the five blocks’ coefficients are

read against each other and against wp1044's French reference; the per-block tables in §4 carry the underlying numbers.




Figure 2: Cross-block β_0 : AU L2 vs wp1044 FR

Figure 5.1: Cross-block error-correction coefficient β_0 from the Phase L2 iterative-OLS estimates (§4.3–§4.7) against the wp1044 French reference. AU is 4–8 $\times$ faster than France in four of five blocks; consumption is the exception (AU $\approx$ FR, the headline finding of §5.2). Business inv carries the wp1044 import (Option 1, §3.6).




Figure 3: Consumption headline: AU/FR β_0 ratio across blocks

Figure 5.2: AU/FR β_0 ratio across the four AU-estimated blocks. Consumption is the only block where the ratio is ≈ 1 (the others are 4–8 $\times$). Re-rendering of the same finding as Figure 5.1, oriented to highlight the consumption exception.

Figure 5.3: Cross-block R^2 of the Phase L2 fits. Four blocks deliver R^2 in [0.41,




Figure 4: Cross-block R²: 4 of 5 blocks fit wp1044 form on AU data

0.81]; *business inv* with strict wp1044 PAC delivers $R^2 = 0.09$ (the structural rejection of §5.3 and §4.6.2). The wp1044 French R^2 range is [0.61, 0.95].

5.1 AU ECM speeds are 4–8× faster than France in four of five blocks

The PAC error-correction coefficient β_0 governs how fast each variable mean-reverts to its long-run target. The Phase L2 estimates show a striking cross-block pattern on AU data: β_0 is 4–8× **wp1044’s French value in four of five blocks**, with the consumption block (β_0 AU = 0.23 vs FR = 0.29) as the sole exception.

Block	β_0 AU L2	β_0 FR (wp1044)	Ratio AU/FR	Notes
VA-price (§4.3.2)	0.258	0.05	5.2×	Eq 16
Employment (§4.4.4)	0.314	0.07	4.5×	Eq 30, depth 3
Consumption (§4.5.2)	0.23	0.29	0.79×	Exception —
Housing inv (§4.7)	0.496	0.12	4.1×	comparable to wp1044 Eq 37
Business inv (§4.6.2)	(wp1044 imported)	0.096	n/a	Option 1 hybrid

Three non-exclusive interpretations of the AU-faster-ECM pattern:

1. **More flexible price and wage setting in AU.** AU has higher import content of CPI (~23% IAD weight vs France's ~17%), variable-rate mortgages that pass policy-rate changes through quickly to household budgets (mortgage spread $\rho_{LH} = 0.97$), and a commodity-cycle pass-through to producer prices via the export sector. Each delivers a faster ECM than the French institutional analogue.
2. **Shorter business cycles in AU.** The Australian business cycle has shorter duration than France's over 1993–2024 (AU sample includes GST 2000, GFC 2008–9, mining boom 2003–14, COVID 2020–1; France's wp1044 sample has fewer regime breaks). Higher-frequency cycles deliver faster ECM mechanically.
3. **HP-trend target proxies overstate ECM speed.** The AU L2 LR target proxies are mostly HP-filter trends, which are by construction smoother than wp1044's structural targets (built from optimisation FOCs). A smoother target $\rightarrow$ larger gaps $\rightarrow$ faster apparent ECM. This is the most quantitatively important methodological caveat, and is the reason the AU β_0 values should be read against each other (the cross-block ranking) rather than against wp1044's absolute values.

The consumption-block exception (AU $\beta_0 \approx$ FR β_0) is structurally meaningful: consumption is the only block where wp1044's target **c*** (built from PV of permanent income via Eq 33) and the AU L2 target proxy converge to functionally similar series on AU data (AU L2 builds **c*** from PV of **yH_gap** per Eq 35 spec). When the targets are structurally aligned, the ECM speeds are too.

5.2 The wp1044 PAC framework validates for four of five blocks on Australian data

A second cross-block reading of the Phase L2 results: the wp1044 PAC structural form (Eqs 16, 30, 35, 37, 46) survives the AU iterative-OLS pipeline on **four of five blocks**:

Block	Coef = +1 on PV satisfied?	R ² AU L2	R ² wp1044 FR	Verdict
VA-price (§4.3.2)	Yes (imposed; growth- neutrality verified)	0.41	0.61	wp1044 form fits AU; R ² gap due to missing Phillips+Okun aux
Employment (§4.4.4)	Yes (imposed; growth- neutrality verified)	0.81	0.95	wp1044 form fits AU; depth-3 correction recovered

Block	Coef = +1 on PV satisfied?	R ² AU L2	R ² wp1044 FR	Verdict
Consumption (§4.5.2)	Yes (imposed; β_{PAC} freely estimated)	0.81	0.95	wp1044 form fits AU; $\beta_0 \approx$ wp1044
Housing inv (§4.7)	Yes (imposed; price-spread active for first time)	0.43	0.89	wp1044 form fits AU; short sample on price spread
Business inv (§4.6.2)	NO — free PV($\Delta\hat{q}$) = −5.03	(0.09 strict; 0.53 free PV)	0.83	Structural rejection — see §5.3

Three economic interpretations:

1. **Four of five validated** is the right way to read the result. Four PAC blocks accept the wp1044 structural restriction (coef = +1 on PV terms; growth-neutrality satisfied at derived coefficient) and deliver R² in [0.41, 0.81] with structurally signed dynamics. This is *consistent* with wp1044's underlying assumption that AU and France share the same agent-objective architecture for those four blocks.
2. **The consumption β_0 match is the single strongest empirical signal of structural cross-country equivalence** — and it is the block where wp1044's permanent-income/PAC framework would most naturally transfer (homogeneous-agent forward-looking consumers facing similar financial structure).
3. **One block fails** (business inv) — and the failure is structural, not finite-sample. This finding stands by itself and motivates the Option 1 hybrid calibration of §3.6 implemented in §4.6.2. See §5.3 for the three explanatory hypotheses.

5.3 Business investment structurally rejects wp1044 PAC — three hypotheses and the Option 1 implementation

Eleven specification variants of the wp1044 BI equation (Eq 46) were tested in Phase L2 P1c. Every strict-PAC variant (PV terms at coef = +1) delivers R² ≤ 0.11 on raw `dln_ib`; the best non-strict variant (PV regressors free) gives R² = 0.53 with PV($\Delta\hat{q}$) coefficient = **−5.03** (wrong sign and 5× too large in magnitude relative to the structural +1). Variants v2 (pre-residualisation), v3-B/C (strict + dummies single-shot or with loose clamps), v4 (combined PV coef = +1 + dummies), v5 (terms-of-trade + piecewise trends), v6 (replace q with `q_AU` = ToT-augmented target), wp736 form (2019 simpler 2-PV form), and the simplified spec (drops PV machinery entirely, R² = 0.33 but no PAC

structure) all yield $R^2 < 0$ or $R^2 < 0.11$. Full table in Appendix G.

Three non-exclusive hypotheses for why AU BI rejects the wp1044 PAC restriction:

- **Hypothesis A: Different agent objective.** AU business-investment agents may not minimise wp1044’s specific quadratic cost function. Heterogeneous agents (mining sector vs non-mining services), financial frictions (relationship lending in AU mining; project-finance lumpiness in LNG/iron ore), or different time-preference rates would change the agent-level FOC and the implied PV coefficient. The aggregate AU `dln_ib` then mixes forward-looking mining capex with backward-looking non-mining capex; the wp1044 single-agent FOC is mis-specified at the aggregate level.
- **Hypothesis B: Mining-sector dominance.** AU’s resource sector represents 10–15% of GDP but ~25% of business investment in mining-boom periods. Mining-sector capex is driven by: long-term commodity demand expectations (China, World steel demand); specific project lumpiness (LNG plants 2010–14; Pilbara iron-ore expansions 2003–13); and foreign direct investment decisions (Asian buyers locking in long-run supply). These dynamics don’t reduce cleanly to wp1044’s “expected market VA growth + user-cost gap” framework — the marginal AU investor is reacting to a different state vector than the wp1044 PAC FOC implies.
- **Hypothesis C: Sample-period contamination.** The AU sample (1993Q2–2023Q3) spans the GST regime change (2000Q3), the mining boom (2003–14), GFC (2008–9), COVID (2020–1), and multiple RBA inflation-targeting regime evolution episodes. Each is a structural break that disrupts long-run PAC identification. French data (longer, smoother, fewer regime breaks over the same window) provides cleaner identification of the underlying PAC restriction. Even if Hypothesis A and B were false, Hypothesis C alone could explain why the strict-PAC restriction looks rejected on the AU sample.

Option 1 implementation (decision-locked at commit 43ed22c). The production simulation model imports the BI deep parameters from wp1044 Table 3.5.13 ($\beta_0=0.096$, $\beta_1=0.33$, $\beta_2=0.11$, $\beta_3=0.69$, $\omega=0.35$, $\sigma_{CES}\approx 0.50$). The Bayesian estimation file removes BI parameters from `estimated_params` (treats them as calibrated). AU-specific BI dynamics enter the model through (i) the AU-driven E-SAT VAR state propagating into BI’s PAC expectation; (ii) the AU-estimated other 4 PAC blocks feeding BI via the GDP identity / demand block; (iii) AU-sized shocks (`stderr_eps_ib` = 2.7211 AU-estimated); and (iv) AU-specific commodity/mining-cycle shocks flowing through trade and the deflator block. The structural restriction on PV terms is preserved (coef = +1 on $PV(\Delta\hat{q}) + PV(\Delta\bar{q})$, coef = $-\sigma$ on $PV(\Delta\log \hat{r}_{KB}) + PV(\Delta\log \bar{r}_{KB})$); the deep parameters are simply borrowed from a better-identified economy. See §3.6 for the methodological rationale and Appendix G for the variant table.

Open question for future work. None of the 11 variants explored Hypothesis A or B directly. A future Phase L3 could (i) split the AU BI series into mining

vs non-mining components using ABS Cat. 5625 (Private New Capital Expenditure by industry), estimate the wp1044 PAC equation on each component separately, and test whether either component satisfies the $\text{coef} = +1$ restriction in isolation; (ii) extend the wp1044 BI equation with a financial-friction shock or a leverage-cycle variable per Bernanke-Gertler-Gilchrist; (iii) lengthen the AU sample by splicing pre-1993 ANA data (with a known regime-break adjustment); (iv) test wp1044 on a longer French sample to verify Hypothesis C's mechanical sample-length effect. None is on the AUSPAC roadmap as of this writing.

5.4 Implications for IRF analysis

The hybrid calibration has two practical implications for the §7 IRFs.

First, IRFs through the BI block are now structurally consistent ($\text{coef} = +1$ on $\text{PV}(\Delta\hat{q})$ is satisfied by construction). The §7.2 monetary-tightening IRF therefore propagates correctly through the BI accelerator ($\beta_3 \cdot \Delta\text{df gap} = 0.69$, vs the pre-hybrid 0.31), through the user-cost PV channel ($\sigma \cdot \text{PV}(\Delta \log r_{\text{KB}})$ at the wp1044 imposed $-\sigma$), and through the cost-of-capital channel via WACC. The pre-hybrid AU-estimated BI parameters ($\beta_0_{\text{ib}}=0.018$, $\beta_1_{\text{ib}}=0.082$) were small enough that the BI channel made only a modest contribution to the IRFs; the post-hybrid values ($\beta_0_{\text{ib}}=0.096$, $\beta_1_{\text{ib}}=0.33$) make the BI channel substantively larger.

Second, AU-specific BI dynamics that were captured (poorly) by free-AU estimation are now captured by the surrounding model rather than by BI's own coefficients. Commodity-price shocks (eps_{pcom}) propagate to AU BI via the trade ECM, the export deflator, and the GDP identity — not through any direct AU coefficient on a commodity regressor in the BI equation. This routing is structurally clean and Lucas-critique-immune in counterfactual analysis: under any monetary-policy regime change, the BI block's FOC still holds at $\text{coef} = +1$ on PV, and the AU-specific dynamics propagate through the other (AU-estimated and AU-shocked) blocks.

The §7 IRF figures and Table 6.3 magnitudes have been regenerated under the production model with equation-by-equation OLS calibration (wp1044 methodology), which includes ULC/UCK structural channels, deflator ECMs, endogenous spread widening, and energy/non-energy import split. The magnitudes quoted in §7.2's channel-by-channel walkthrough reflect this OLS calibration.

6. Estimation

6.1 Data

Eleven observables are used for estimation, covering the period 1993Q2-2023Q3 (122 quarters for PAC estimation) and 1993Q1-2024Q4 (128 quarters for E-SAT).

Import and export volume growth were added to the observation set together with the trade-block proper-ECM rewrite (Section 4.8) so that the long-run elasticities β_m , β_x , γ_m , γ_x are data-identified.

Table 5.1: Observable variables

#	Variable	Source	Transformation	Mean	Std
1	yhat_au	ABS GDP, HP-filtered	Log deviation from trend	0.01	1.04
2	pi_au	ABS GDP deflator	Quarterly log difference	0.66	0.60
3	i_au	RBA cash rate	Already quarterly %	1.04	0.52
4	yhat_us	US GDP, HP-filtered	Log deviation from trend	-0.55	1.88
5	pi_us	US GDP deflator	Quarterly log difference	0.54	0.36
6	pi_w	ABS 6345 Wage Price Index (Private+Public, All industries, SA)	Quarterly log difference $\times 100$	0.81	0.39
7	dln_c	ABS consumption	Quarterly log difference (demeaned)	0.00	1.78
8	dln_ib	ABS non-dwelling GFCF	Quarterly log difference (demeaned)	0.00	2.79
9	i_10y	AU 10Y govt bond rate	Annualized / 4 to quarterly	1.21	0.53
10	dln_m	ABS 5206 imports SA (A2304115J)	Quarterly log difference (demeaned)	0.00	~2.7
11	dln_x	ABS 5206 exports SA (A2304114F)	Quarterly log difference (demeaned)	0.00	~3.3

Rates and inflation are in natural units (matching model SS). Growth rates are demeaned (model SS = 0). Note: `au_irate` in `dataset.csv` is already in

quarterly percentage points. The openness drift in M/C and X/C is absorbed by demeaning each growth-rate observable independently; the identifying variation for β_m and β_x comes from cyclical co-movement of the import/export levels with cumulated domestic demand and the real exchange rate.

Figure 5.1: The nine pre-existing observable variables used in AU-PAC estimation, 1993Q1-2024Q4. The two new trade observables (dln_m , dln_x) follow the same series construction as dln_c / dln_ib .

6.2 E-SAT Bayesian estimation

Historical (retired MCMC pipeline). The LMD figure below was produced by `au_pac_bayesian.mod`, removed in cleanup 7995ce7, and is **not reproducible** from the current tree. The E-SAT auxiliary-regression posteriors it produced are nonetheless still active in `au_pac.mod` (Phase B, written back as point values); only the marginal-likelihood number is historical.

The E-SAT core was estimated using Bayesian MCMC (Metropolis-Hastings, 50,000 draws, 2 chains). Results are reported in Table 3.1. Key posterior findings relative to calibration:

- Wage persistence (λ_w): 0.225 vs calibrated 0.55. The full-system Bayesian estimate (§6.4) refines this to 0.329, consistent with moderate own-lag persistence in a standard NK wage Phillips curve.
- CPI passthrough (γ_w): 0.770 in E-SAT, refined to 0.138 in the full-system Bayesian (§6.4) — moderate contemporaneous CPI passthrough.
- Demand bridge (λ_dom): 0.399 vs calibrated 0.10 — demand feed-back four times stronger than assumed.
- Consumption persistence ($b1_c$): 0.047 vs calibrated 0.35 — less serial correlation than the prior implied.
- Exchange rate (ρ_s): 0.775 vs calibrated 0.95 — AUD half-life is approximately 3 quarters.

Log marginal density (Modified Harmonic Mean): -1095.38.

Figure 5.2: Prior vs posterior densities for E-SAT core parameters.

6.3 PAC structural estimation — intermediate methodology audit

The estimation strategy reported in this subsection — equation-by-equation iterative OLS following the wp1044 methodology — is the **production calibration** used for all IRFs, conditional forecasts, and experiments reported in §7 and §8. This is the standard FR-BDF approach: each PAC block is estimated by OLS on its own structural equation with block-specific auxiliary VARs. Earlier versions of this paper used full-system Bayesian MCMC as the authoritative calibration; the current version adopts the wp1044 OLS methodology, which lets AU data speak without Bayesian prior regularisation. The key empirical finding from the switch is that the consumption block’s direct interest-rate coefficients (β_2

$= -0.003$, $\beta_3 = -0.014$) are near zero under OLS, whereas the earlier Bayesian priors had pulled these toward economically significant values.

5.3.1 Methodology: Iterative OLS PAC structural parameters are estimated using Dynare’s `pac.estimate.iterative_ols` routine, following the ECB-Base methodology (Zimic, SemiStructDynareBasics). The algorithm:

1. Initialize PAC parameters at calibrated values
2. Compute h-vectors from the `var_model` companion matrix using `get_companion_matrix('esat_enriched', 'var')`
3. Construct the PAC expectation term using the h-vectors
4. Run OLS on the PAC equation, updating parameter estimates
5. Recompute h-vectors with updated parameters
6. Iterate until convergence (change in parameters $<$ tolerance)

5.3.2 Three dseries approaches Three approaches to constructing the auxiliary variables for estimation were compared:

Table 5.4: Three-approach comparison (SSR)

Equation	(A) Recursive	(B) Hybrid smoother	(C) Pure smoother
VA Price	40.4	40.3	1.1*
Consumption	435.4	436.8	400.8
Business Inv	973.9	972.7	968.4
Household Inv	965.8	963.1	960.0
Employment	76.1	74.3	73.4

*Pure smoother VA price SSR is artificially low — `pv_aux` absorbs all signal, leaving structural parameters unidentified.

Recommended approach: (B) Hybrid — Kalman-smoothed targets (from `calib_smoother` with `diffuse_filter`) for the EC term, with recursive `pv_aux` corrections. This avoids the over-identification problem of approach (C) while giving better target estimates than approach (A).

5.3.3 COVID dummy treatment Two pulse dummies absorb the extreme 2020 outliers: `d_covid_crash` ($= 1$ in 2020Q2) and `d_covid_bounce` ($= 1$ in 2020Q3). Each PAC equation has its own pair of COVID coefficients (10 new parameters total). The dummies fixed two sign problems: consumption AR1 flipped from -0.25 to $+0.05$, and employment AR1 from -0.26 to $+0.32$.

Table 5.5: Final PAC estimates — hybrid smoother with COVID dummies

Estimated via iterative OLS on the Kalman-smoothed hybrid dseries. The 12×12 companion matrix reflects the final AU calibration of auxiliary dynamics, deflators ($\rho_{pc}=0.67$, $\rho_{pib}=0.70$), trade ($b1_x = 0.30$, $b1_m = 0.232$), housing ($\rho_{ph} = 0.60$), and the mortgage rate ($\rho_{lh} = 0.97$).

Parameter	VA				
	Price	Consumption	Bus. Inv	Housing Inv	Employment
b0 (EC)	0.028	0.067	0.017	0.017	0.063
b1 (AR1)	0.288	0.033	0.093	0.101	0.314
b2	-0.014	-0.541	-0.045	-0.046	-0.187
b3	—	-0.024	0.344	0.289	-0.076
b4	—	—	—	—	-0.085
b5	—	—	—	—	-0.017
COVID crash	-2.877	-15.096	-4.382	-4.815	-6.601
COVID bounce	+1.488	+6.033	+2.978	+3.194	+3.851
SSR	40.6	410.4	929.8	960.0	76.3
T	118	118	118	118	118

The direct interest-rate consumption channel $b_{di,c}$ and the housing-price gap $b_{ph,ih}$ are not identified by single-equation OLS on the AU sample (wrong-signed under reverse causality). $b_{ph,ih}$ is identified via a lag-2-ph_gap IV on the 1959Q3-spliced housing-price series (+0.0099, s.e. 0.075, first-stage F = 432), at which the structural Tobin’s Q channel is statistically indistinguishable from zero. $b_{di,c}$ remains Bayesian-regularised at -0.70 pending RBA OIS-surprise data for the Romer–Romer IV.

6.4 Bayesian full-system estimation

The production AUSEPAC calibration (Table 5.6) is produced by **equation-by-equation OLS** following the wp1044 methodology — the standard FR-BDF approach. The Bayesian full-system MCMC estimation described below was used in earlier versions of this paper; **its driver au_pac_bayesian.mod was removed in cleanup 7995ce7, so the LMD figures in this subsection are historical and not reproducible from the current tree.** It is documented here only for provenance. Joint Bayesian estimation used Dynare’s `estimation()` command with `diffuse_filter` (for unit-root level accumulators), starting from calibrated values to ensure Blanchard-Kahn satisfaction. The input data were *raw* quarterly observables — no Kalman smoothing, HP-filtering, or other pre-processing is applied to the observation series.

Table 5.6: Equation-by-equation OLS posterior (hybrid calibration, BI block calibrated from wp1044)

Calibration methodology: Equation-by-equation OLS following wp1044 methodology. PAC block parameters are estimated by iterative OLS on block-specific structural equations; shock standard deviations are calibrated from OLS residuals. Three BI parameters (b0_ib, b1_ib, b3_ib) are imported from wp1044 Table 3.5.13 per the Option 1 hybrid calibration of §3.6 + §4.6.2.

Note: Earlier versions of this paper used full-system Bayesian MCMC estimation. The current production calibration uses equation-by-equation OLS following the wp1044 methodology, which is the standard FR-BDF approach. The BI block remains calibrated from wp1044 (Option 1 hybrid).

The full-system Bayesian estimation (retained as a robustness check, not the production calibration) jointly identifies 19 PAC and wage-block parameters together with 9 shock standard deviations, conditional on the calibrated supply, deflator, financial, trade, and UIP-NPV parameters quoted elsewhere in §4. The estimation is two-stage: csminwel mode search followed by Metropolis-Hastings MCMC (20,000 draws across two chains, ~51 min wall time). Posterior modes are within +/-5% of posterior means for all coefficients, indicating well-shaped log-likelihood surfaces.

Parameter	Prior	Posterior mean	90% HPD
b0_pQ	Beta(0.03, 0.015)	0.0330	[0.0064, 0.0603]
b1_pQ	Beta(0.29, 0.10)	0.2795	[0.1092, 0.4438]
b2_pQ	Normal(0.00, 0.05)	0.0077	[-0.0723, 0.0892]
b0_c	Beta(0.07, 0.03)	0.0559	[0.0248, 0.0880]
b1_c	Beta(0.05, 0.03)	0.0375	[0.0060, 0.0676]
b2_c	Normal(-0.55, 0.20)	-0.3284	[-0.6364, -0.0485]
b3_c	Normal(0.02, 0.05)	0.0147	[-0.0555, 0.0927]
b0_ib	Beta(0.02, 0.01)	0.0188	[0.0058, 0.0311]
b1_ib	Beta(0.09, 0.05)	0.0884	[0.0203, 0.1545]
b3_ib	Normal(0.34, 0.10)	0.3225	[0.1866, 0.4708]
b0_ih	Beta(0.03, 0.015)	0.0306	[0.0077, 0.0536]
b1_ih	Beta(0.11, 0.05)	0.1123	[0.0381, 0.1907]
b3_ih	Normal(0.23, 0.10)	0.2237	[0.0725, 0.3826]
b0_n	Beta(0.06, 0.03)	0.0564	[0.0143, 0.0961]
b1_n	Beta(0.31, 0.10)	0.3116	[0.1495, 0.4654]
b5_n	Normal(0.00, 0.05)	-0.0023	[-0.0767, 0.0831]
lambda_w	Beta(0.25, 0.10)	0.2091	[0.0786, 0.3315]
gamma_w	Beta(0.45, 0.15)	0.3536	[0.2676, 0.4283]
kappa_w	Normal(-0.08, 0.05)	-0.1025	[-0.1776, -0.0185]
stderr eps_q	InvGamma(0.80)	0.5372	[0.4760, 0.6013]
stderr eps_i	InvGamma(0.10)	0.1112	[0.0991, 0.1223]

Parameter	Prior	Posterior mean	90% HPD
stderr eps_pi	InvGamma(0.60)	0.5346	[0.4667, 0.6036]
stderr eps_c	InvGamma(0.50)	1.8364	[1.6397, 2.0303]
stderr eps_ib	InvGamma(1.50)	2.7469	[2.4438, 3.0454]
stderr eps_ih	InvGamma(2.00)	1.3377	[0.4890, 2.2349]
stderr eps_n	InvGamma(0.50)	0.4797	[0.1175, 0.9249]
stderr eps_w	InvGamma(0.30)	0.1407	[0.0950, 0.1906]
stderr eps_10y	InvGamma(0.10)	0.0656	[0.0540, 0.0791]

The log marginal density was **Laplace** = **−781.05** and **Modified Harmonic Mean** = **−781.39** (Brooks–Gelman PSRF convergence diagnostics passed for all 28 parameters). (*Historical: from the retired `au_pac_bayesian.mod`; not reproducible from the current tree.*)

Three findings stand out from the joint posterior. **First**, the wage Phillips slope $\kappa_w = -0.103$ identifies cleanly with 90% HPD entirely negative; under the FR-BDF eq. (52) sign convention $\dots - \kappa_w \cdot \text{pv_u_gap}$ this corresponds to a *positive* structural Phillips slope (higher unemployment $\rightarrow$ lower wage growth), consistent with both the wp736 specification and the post-2015 RBA literature on AU wage non-responsiveness to slack (Cassidy and Doyle 2018; Bishop and Cassidy 2017). **Second**, CPI indexation $\gamma_w = 0.354$ is large and sharply identified, consistent with the substantial fraction of the AU wage bill anchored to Fair Work Commission CPI-indexed award-rate decisions plus flow-on effects to enterprise bargaining. **Third**, the consumption-block rate channel $b_2^c = -0.328$ is significantly negative (90% HPD $[-0.636, -0.049]$), and the consumption error-correction speed $b_0^c = 0.056$ is moderate; together with the forward-NPV real-lending-rate channel ($\alpha_{c,r} \cdot \text{pv}_{r,lh,gap}$, eq. 24) these identify a robust interest-rate sensitivity for AU households despite their high ($\sim 70\%$) variable-rate mortgage exposure.

6.5 Pseudo-out-of-sample forecast evaluation

Reproducibility note (2026-05-30, Wave 4). The original full-system exercise used the Kalman smoother + Dynare `forecast=8` block of the MCMC pipeline, removed in cleanup 7995ce7 (the production model no longer carries a `varobs/estimation` interface). The driver `forecast_eval.m` has been **reconstructed** as a reproducible 1-step-ahead recursive evaluation in the spirit of the equation-by-equation OLS methodology; it runs on `estimation_data.mat` under MATLAB R2026a with no Dynare dependency. The multi-horizon numbers in the prior version of Table 5.8 came from the lost driver and are not reproducible; the reproducible 1-step results are below.

Over the last 24 one-step transitions ($\sim 2018\text{Q1}–2023\text{Q4}$), with coefficients held fixed, each observable’s 1-step RMSE is compared across three naive benchmarks

— random walk (RW), recursively re-fit AR(1), and the unconditional mean (= 0 in gap form) — and, for the two observables whose model equation is a clean *predetermined* autoregression (the CPI Phillips `pi_au` and the E-SAT IS curve `yhat_au`), the model’s fixed-coefficient structural 1-step forecast.

Table 5.8: Reproducible 1-step recursive RMSE (24 origins, ~2018Q1–2023Q4)

Observable	RW	AR(1)	mean=0	MODEL
<code>pi_au</code> (CPI)	0.919	0.784	0.805	0.745
<code>pi_w</code> (wages)	2.958	2.286	2.120	—
<code>yhat_au</code> (gap)	1.944	1.720	2.016	1.802
<code>dln_c</code>	5.885	5.703	3.821	—
<code>dln_ib</code>	2.625	2.162	2.179	—
<code>i_au</code>	0.027	0.027	0.190	—
<code>i_10y</code>	0.096	0.097	0.709	—
<code>yhat_us</code>	2.367	2.322	2.201	—
<code>pi_us</code>	0.459	0.472	0.672	—

The headline finding is positive: the AU-estimated **CPI Phillips beats all three naive benchmarks at h=1** (MODEL 0.745 < AR(1) 0.784 < mean 0.805 < RW 0.919) — so despite its flat *in-sample* R^2 (§6.6), the equation carries genuine, if modest, out-of-sample inflation-forecasting skill. The **IS curve** (`yhat_au`, MODEL 1.802) likewise beats the random walk (1.944) and the unconditional mean (2.016) and roughly matches a recursive AR(1). As expected, the highly persistent policy and long rates (`i_au`, `i_10y`) are best forecast by the random walk, while the noisy growth observables (`dln_c`, `dln_ib`, `pi_w`) are best forecast by their unconditional mean — the evaluation window contains the 2020 COVID quarter (`dln_c` $\approx -16\%$), which no equation can forecast. (The reproduced `pi_au` and `yhat_au` h=1 RMSEs, 0.745 and 1.802, line up with the prior table’s 0.782 and 1.819, cross-validating the reconstruction.)

Figure 5.5: Recursive forecast paths (light grey, 24 origins $\times$ 8 quarters) vs realised observable (black). The four panels show output gap, CPI inflation, real consumption growth, and policy rate. The model consistently under-shoots the 2020Q4 inflation rebound and the 2022–23 hiking cycle for `i_au`, reflecting the absence of forward guidance in the unconditional zero-shock projection.

Figure 5.6: Pseudo-real-time RMSE by horizon, all observables. Output gap and CPI inflation are the best-forecasted variables in proportion to their volatility; cash rate and 10Y yield grow monotonically with horizon as expected. The model’s strongest comparative advantage relative to a random walk is in the financial variables (`i_au`, `i_10y`) at h=1 (RMSE < 0.10).

Full recursive re-estimation (REESTIMATE=true in `forecast_eval.m`, which re-runs `csmnwel` per origin) is left as a separate computation; the qualitative

RMSE pattern is not expected to change materially because the posterior is informed by the full sample and the model is reasonably well-identified at the posterior mode (see the identification analysis in Appendix F).

6.6 Phase L2 OLS audit (2026-05-28)

A final equation-by-equation OLS pass replaced remaining calibrated coefficients with AU point estimates where data permitted. Two motivating observations:

1. **Orphan parameters.** Three calibrations ($b_{di_c} = -0.701$, $b_{ph_ih} = 0.0099$, b_{PAC_c}) were active in the .mod file but not part of any wp1044 estimation specification — they amplified responses through channels with no OLS identification. The audit zeroed them where genuinely orphan and traced wiring for `pv_r_lh_gap` into the consumption SR PAC equation (wp1044 Eq 35).
2. **CPI/trade/deflator blocks** were carrying wp1044 (France) calibrations rather than AU point estimates. Where AU data exists (ABS 5206 IPD, RBA I02 commodity index, RBA F11 TWI), single-equation OLS was run and the estimates written back.

Data: ABS 5206 chain volumes + IPD percentage-changes (cols 75–77), RBA I02 commodity prices (526 monthly obs, 1982–2026), RBA F11 trade-weighted index (672 monthly obs after splicing the pre-2009 archive). The OLS pipeline is `data/pac_blocks/estimate_{cpi_phillips,deflators,trade_exports,trade_imports,wage_phillips}` all using `ols_with_se` from `data/pac_helpers/`.

Table 6.6: Phase L2 OLS audit — coefficients replaced with AU point estimates

Block	Coefficient	wp1044/prior	AU OLS	t-stat	N	R ²	Status
CPI Phillips (Eq 51)	λ_{π}	0.402	0.259	2.76	121	0.06	identified
CPI Phillips	κ_{π}	0.098	−0.034	−0.61			insig, sign flip
CPI Phillips	α_{pc}	0.385	0.001	0.01			insig
CPI Phillips	α_{pc_lag}	0.023	0	—			dropped (mc)
CPI Phillips	β_{pc_m}	0.10	−0.004	−0.21			insig
Exports SR (Eq 70)	$b0_x$	0.05	0.299	−7.66	103	0.73	identified

Block	Coefficient	wp1044/pOLS	AU pOLS	t-stat	N	R ²	Status
Exports SR	b1_x	0.30	0.867	14.70			identified, high persistence
Exports SR	b2_x	0.25	0.022	0.19			insig
Exports SR	b3_x	0.10	-0.361	-0.22			insig
Exports SR	b4_x	0.15	0.010	1.11			insig
Imports SR (Eq 74)	b0_m_ne	0.06	0.158	-3.95	103	0.64	identified
Imports SR	b1_m_ne	0.23	0.743	10.02			identified, high persistence
Export deflator π_x	ρ_{px}	0.21	0.043	1.40	139	0.89	insig
Export deflator	α_{px}	0.20	-0.033	-0.24			insig
Export deflator	β_{px}	-0.05	-0.244	-0.11			insig
Export deflator	α_{pcom}	0.10	0.584	30.34			commodity-dominated
Import deflator π_m	ρ_{pm}	0.28	0.223	3.08	139	0.34	identified
Import deflator	α_{pm}	0.38	0.057	0.25			insig
Import deflator	β_{pm}	0.09	-8.59	-2.16			reverted (BK)
Import deflator	β_{pm_com}	0.42	0.221	7.56			identified

Block	Coefficient	wp1044/pOLS	AU OLS	t-stat	N	R ²	Status
BI de- flator π_{ib}	ρ_{pib}	0.70	0.474	7.37	139	0.39	identified
BI de- flator	α_{pib}	0.19	-0.086	-1.71			insig, sign flip
BI de- flator	β_{pib_m}	0.12	0.097	5.95			identified
HI de- flator π_{ih}	ρ_{pih}	0.49	0.366	4.49	139	0.10	identified
HI de- flator	α_{pih}	0.40	-0.055	-0.65			insig
HI de- flator	β_{pih_m}	0.08	-0.019	-0.71			insig
Govt defla- tor π_g	ρ_{pg}	0.13	-0.471	-5.95	122	0.23	sig, sign flip (oscilla- tion)
Govt defla- tor	α_{pg}	0.37	0.036	0.63			insig
Wage Phillips λ_w	0.21	-0.082 /	-1.64	121	0.77 / 0.30	†BK- constrained (§6.7)	
Wage Phillips γ_w	0.46	1.594 / 0.857†	15.27			†BK- constrained	
Wage Phillips κ_w	0.054	0.371 / 0.359†	3.18			†BK- constrained	
Orphan	b_{di_c}	-0.701	0	—			not in wp1044 Eq 35 L2
Orphan	b_{ph_ih}	0.0099	0	—			skipped (no AU data)
LR trade β_x	1.20	0.005	0.01			reverted (insig)	

Block	Coefficient	wp1044/p1044	AU OLS	t-stat	N	R ²	Status
LR trade γ_x	0.40	−9.50	−1.66				reverted (excessive)
LR im-ports β_{m_ne}	1.50	0.073	1.36				reverted (insig)
LR im-ports γ_{m_ne}	−0.40	23.38	1.09				reverted (wrong sign)
Trade-deflator s_gap	0.09	−8.59	−2.16				reverted (sign flip)
β_{pm_ne}							

Read-out (bolded rows = wp1044 calibration replaced verbatim with AU OLS):

- **25 of 35 audited coefficients** were written back as AU OLS point estimates (per the *use-OLS-unless-Dynare-fails* convention adopted for this audit).
- **10 coefficients reverted** for one of two reasons. (a) *Blanchard–Kahn failure* — the unconstrained wage Phillips OLS gave $\gamma_w = 1.59 > 1$, which produces a wage-price spiral that violates BK; LR trade and import-deflator elasticities had insignificant or wrong-signed estimates whose magnitudes pushed eigenvalues outside the unit circle. (b) *Spreads and UIP* (α_s , ρ_s , $\kappa_{spread_LB/BBB}$, β_{uip}) were not estimated — these require GMM on forward recursions (pv_i_uip), not single-equation OLS.

The headline empirical finding is the *AU flat-Phillips* result on the CPI side: with $N = 121$ quarterly observations (1993Q3–2023Q3), only the lag coefficient $\lambda_\pi = 0.259$ is identified ($t = 2.76$); output-gap, VA-price, lagged-VA-price, and import-price channels are all insignificant under single-equation OLS. The R^2 of 0.063 is extremely low — AU CPI inflation gaps are essentially noise around their lagged value at quarterly frequency. The same data also reveal a **wage-side anomaly**: unconstrained OLS gives $\gamma_w = 1.59 > 1$, implying over-indexation of nominal wages to CPI (which BK-fails the structural model and is replaced by the AU pre-MCMC OLS estimate $\gamma_w = 0.46$). On the price side the **export deflator is commodity-dominated** with $\alpha_{pcom} = 0.584$ ($t = 30.3$, $R^2 = 0.89$), confirming AU’s commodity-cycle pricing structure.

Implication for Table 6.3. The post-audit IRF peaks for real activity ($\ln_Q$, $yhat_au$) and exchange rate (s_gap , 10Y) are essentially unchanged because monetary transmission to those variables runs through channels (CES user cost, UIP forward recursion, term structure) that were not touched by the audit. The CPI IRF response weakens further because the AU CPI OLS finds the

structural channels insignificant; this is the same “flat AU Phillips” finding that wp736 calibrated around, now reproduced in the L2 OLS pass.

6.7 BK-constrained wage Phillips (2026-05-28)

The Phase L2 unrestricted wage Phillips OLS gave $\gamma_w = 1.59$ ($t = 15.3$) — strongly identified but BK-infeasible because the wage-price spiral closes (CPI passes through to wages at $1.59\times$ while wages feed back into the CPI deflator and VA price). A re-estimation under the structural functional form

$$\text{pi_w} - \text{pibar} - \text{dln_prod} = \lambda_w \cdot (\text{pi_w}(-1) - \text{pibar} - \text{dln_prod}) + \gamma_w \cdot (\text{pi_c} - \text{pibar}) - \kappa_w \cdot \text{u_gap} + \varepsilon$$

with the BK-stability constraint $\lambda_w + \gamma_w \leq \mathbf{0.95}$ and non-negativity on all coefficients, solved as a quadratic program against $N = 121$ quarterly observations, yields:

Table 6.7: BK-constrained wage Phillips OLS

Coefficient	Unconstrained OLS	Constrained OLS	Status
λ_w (own-lag persistence)	0.219 ($t=3.48$)	0.093	inactive bound (≥ 0)
γ_w (CPI passthrough)	1.471 ($t=10.9$)	0.857	binding (sum=0.95)
κ_w (Phillips slope)	0.065 ($t=0.40$)	0.359	inactive (positive)
$\lambda_w + \gamma_w$	1.689	0.950	—
R^2	0.424	0.301	loss 11.7 nats

The constraint is **binding at the sum**: AU wages are nearly fully indexed to CPI inflation in the small open-economy sense, but stability requires holding the wage-price-spiral coefficient just below unity. The Phillips slope $\kappa_w = 0.359$ (Q-bar AU unemployment-gap units) is now substantial and structurally identified — substantially larger than the wp736 calibration centre of 0.05 and the FR-BDF point estimate around 0.10. Heuristic standard errors are not reported because the active constraint makes classical inference inappropriate; the loss of 11.7 nats vs unconstrained corresponds to a likelihood ratio in the high-significance range, consistent with the wage Phillips structure being a meaningful binding constraint on AU data.

Implication for IRFs. With κ_w now identified, the wage IRF response to a 100bp MP tightening peaks at -0.011 pp y/y at Q7 (vs -0.002 pp at Q2 under the pre-MCMC unrestricted values), a $5\text{--}10\times$ amplification reflecting the active unemployment channel. Real-activity and exchange-rate IRFs are unchanged

because those run through capital/UIP/term-structure mechanics, not through wage formation.

6.8 Phase L2.A architectural fixes (2026-05-28)

Following the L2 OLS audit, a deeper architectural review revealed that three model state variables were being defined by E-SAT-style reduced-form equations rather than the structural identities used in FR-BDF. Per wp736 §2.3, §4.3.3, §4.4 (eq 47), and §4.5.1, the E-SAT IS curve, Phillips curve, and Okun’s law are tools agents use to **forecast** future paths of $\hat{y}$, π , u . They are not the **defining** equations for the simulator’s actual state variables, which come from supply-side (CES) potential output, demand-side identity for actual GDP, weighted-aggregate deflator block for CPI, and labour-market accounting for the unemployment gap. The relevant FR-BDF quote from §4.3 is unambiguous: “FR-BDF has a fully specified supply block which enables the model to endogenously converge toward the long-run or natural level of GDP, **as well as to define the output gap used in E-SAT**” — that is, the supply block *defines* the gap, and E-SAT *consumes* it.

In the production AU-PAC model up to commit 6cd4fb9, the equations at au_pac.mod:1051, 1071, and 1217 were:

$$\begin{aligned} \text{yhat_au} &= \delta \cdot \text{yhat_us} + \lambda_q \cdot \text{yhat_au}(-1) - \sigma_q \cdot (\text{i_gap}(-1) - \text{pi_au_gap}(-1)) \\ &\quad + \lambda_{\text{dom}} \cdot \text{yhat_dom} + \varepsilon_q \quad (\text{E-SAT IS curve defining yhat_au}) \\ \text{pi_au_gap} &= \lambda_\pi \cdot \text{pi_au_gap}(-1) + \kappa_\pi \cdot \text{yhat_au}(-1) \\ &\quad + \alpha_{\text{pc}} \cdot (\text{piQ} - \text{pi_bar_au}) + \dots + \varepsilon_{\text{pi}} \quad (\text{E-SAT Phillips defining pi_au_gap}) \\ \text{u_gap} &= \rho_u \cdot \text{u_gap}(-1) + \text{okun_coeff} \cdot \text{yhat_au} \quad (\text{Okun shortcut defining u_gap}) \end{aligned}$$

In the architectural fix (commits dd45c87, f276ebe, 053bf57), these were replaced with FR-BDF-fidelitous structural identities:

Table 6.8: Phase L2.A architectural fixes

Variable	Pre-fix (E-SAT shadow)	Post-fix (structural identity)	FR-BDF reference
yhat_au	E-SAT IS curve	yhat_au = yhat_au(-1) + yhat_dom + ε_q (accumulation of demand-identity growth)	wp736 §2.3, §4.3.3

Variable	Pre-fix (E-SAT shadow)	Post-fix (structural identity)	FR-BDF reference
pi_au	derived from E-SAT Phillips via $\pi_{au} = \pi_{au_gap} + \pi_{bar_au}$	$\pi_{au} = w_{food} \cdot \pi_{au_food} + w_{energy} \cdot \pi_{au_energy} + (1 - w_{food} - w_{energy}) \cdot \pi_{au_core}$ (structural aggregator)	wp736 §4.7 (demand deflators)
pi_au_gap	E-SAT Phillips	$\pi_{au_gap} = \pi_{au} - \pi_{bar_au}$ (definitional)	wp736 §4.7
pi_au_core	back-out residual	$\pi_{au_core} = \pi_c + \varepsilon_{\pi}$ (consumption deflator + CPI noise)	wp736 §4.7.1
u_gap	Okun shortcut with persistence 0.946	$u_{gap} = -\ln_n_level$ (negative of cyclical employment gap; log-linear from N/LF identity)	wp736 §4.5.1

The L2 OLS estimates of the now-demoted equations (`lambda_pi`, `kappa_pi`, `alpha_pc`, `b_ECM_pc`, `okun_coeff`, `rho_u_gap`) are preserved in `data/pac_blocks/results_cpi_phillips.tx` and the parameter declarations of `au_pac.mod`. They can be re-attached to E-SAT auxiliary forecasting equations when the `var_model` machinery is rebuilt.

Table 6.9: Impact on the 100bp annualized monetary tightening IRF
Comparison of pre-L2.A (commit 6cd4fb9, all L2 OLS coefficients written back + BK-constrained wage Phillips) vs post-L2.A (commit 053bf57, all three architectural fixes applied). Same calibration except for the three structural re-definitions.

Variable	Pre-L2.A	Post-L2.A	Δ
Real GDP ($\ln_Q$)	-0.088% Q8	-0.135% Q9	+53% larger response

Variable	Pre-L2.A	Post-L2.A	Δ
Output gap (yhat_au)	−0.069% Q8	−0.096% Q9	+40%
CPI inflation (pi_au, y/y)	−0.020 pp Q16	−0.004 pp Q3	smaller and earlier
pi_au_food (new decomp)	—	+0.019 pp Q1	—
pi_au_energy (new decomp)	—	−0.048 pp Q9	—
pi_au_core (new decomp)	—	+0.001 pp Q11	$\approx$ pi_c
VA price (piQ, y/y)	+0.030 pp Q1	+0.032 pp Q1	unchanged
Consumption growth	−0.020% Q1	−0.027% Q2	+35%
Business inv growth	−0.055% Q8	−0.093% Q9	+69%
Housing inv growth	−0.027% Q4	−0.028% Q4	unchanged
Employment growth	−0.005% Q5	−0.007% Q12	+40%, later
Wage inflation (y/y)	+0.012 pp Q47 (overshoot)	−0.005 pp Q2	reverses sign; now structurally responsive
Exchange rate (s_gap, %)	−0.980% Q8	−0.988% Q8	unchanged
10Y yield (annualised)	+0.103 pp Q1	+0.105 pp Q1	unchanged

Three findings emerge:

1. **Real-activity transmission is larger** under the structural identity. The pre-L2.A E-SAT IS curve had $\sigma_q = 0.065$ with persistence $\lambda_q = 0.70$ acting in parallel to the demand-side aggregator (with weight $\lambda_{\text{dom}} = 0.40$); the structural fix removes this dampening parallel channel, letting the PAC blocks (business-investment WACC, housing mortgage, consumption pv_r_lh_gap, trade real-exchange-rate) fully aggregate into yhat_au via the demand identity. The result: AU-PAC moves from the “conservative outlier” position in the RBA-suite comparison (Table 6.4) to a more central position.
2. **The CPI response decomposes interpretably** under the structural aggregator. The headline pi_au peak is now smaller (−0.004 pp vs −0.020 pp pre-L2.A) because the structural deflator components partially offset: pi_au_food rises (+0.019 pp on impact via VA-price and commod-

ity passthrough) while `pi_au_energy` falls (-0.048 pp at Q9, dominated by import-price and commodity passthrough as the AUD appreciates), and `pi_au_core` (consumption deflator) barely moves ($+0.001$ pp). This matches the L2 OLS empirical finding ($R^2 = 0.06$; only persistence identified): AU CPI inflation is dominated by import/commodity cycles, not by the output-gap Phillips channel.

3. **The wage Phillips slope $\kappa_w = 0.359$ now transmits** to nominal wages. Pre-L2.A, `u_gap` was driven by the Okun shortcut with persistence 0.946, taking dozens of quarters to respond; the wage Phillips κ_w channel saw essentially no signal. Post-L2.A, `u_gap` = $-\ln n_{\text{level}}$ inherits the PAC employment dynamics (faster, structurally identified), so the wage Phillips responds on realistic horizons. The pre-L2.A `pi_w` peak of $+0.012$ pp at Q47 was an overshoot in the tail; the post-L2.A peak of -0.005 pp at Q2 is the structural impact response.

The architectural changes are documented in detail in `ESAT_ARCHITECTURE_AUDIT.md`, which also lays out a checklist for future development to avoid the same pattern.

Known caveats: - The first known caveat — stale `h_pac_*` coefficients — is **resolved in §6.9 below** (2026-05-28). - ~~The L2-OLS estimates of the PAC equation parameters (`b0_c`, `b1_c`, `b0_pQ`, etc.) were obtained using the OLD `yhat_au` as a regressor in the iterative-OLS procedure. They are not formally re-optimal under the new structural definition. A re-fit is left for future work.~~ **Resolved (2026-05-31):** this concern is moot for the single-equation iterative-OLS. The estimators use the *observed* output gap (`au_ygap` from `dataset.csv`, via `prepare_12_data.m:151`) as the `yhat_au` regressor — not the model’s internal `yhat_au` equation. The §6.8 fix changed the *simulator’s* `yhat_au` definition (E-SAT shadow $\rightarrow$ structural identity), and per the §6.8 logic the E-SAT VAR remains the expectation-forming tool for the PV term; neither estimation input is affected. A full re-run of the 5-block pipeline (`data/run_all_12_ols.m`) confirms this: VA-price reproduces `b0_pQ` = 0.2581 / `b1_pQ` = 0.3039 *exactly* (committed 0.2580 / 0.3039), and employment is within rounding. (Two blocks show a drift unrelated to `yhat_au` — consumption `b0` 0.27 $\rightarrow$ 0.23 and housing `b0` 0.50 $\rightarrow$ 0.60, at identical R^2/N . This is the SA-data re-estimation documented in §6.14’s “Full re-estimation check”; it was **written back 2026-05-31** after confirming the data layer regenerates deterministically. Separately, the χ verification (`dynare/verify_pac_chi_pv.m`) found the employment block’s estimation-side χ was computed by an *approximate* depth-1 solver; fixing it to the exact depth-m solver (§6.9) raised R^2 0.82 $\rightarrow$ 0.86 and moved `b0_n` 0.31 $\rightarrow$ 0.48, also written back.)

6.9 PAC policy-function regeneration (2026-05-28)

Per the §6.8 follow-up, the five `h_pac_*` policy-function coefficient blocks in `au_pac.mod` were regenerated by re-running Dynare’s `pac.print()` against the post-L2-OLS / post-L2.A calibration. Five auxiliary `.mod` files were

reconstructed under `dynare/aux/` — `aux_pQ.mod`, `aux_consumption.mod`, `aux_business_inv.mod`, `aux_housing_inv.mod`, `aux_employment.mod` — each declaring a pure-VAR-form companion (E-SAT core + AR(1) reduced forms for the structural state variables that are no longer E-SAT-defined) plus the block-specific auxiliary regression. The aux files use the L2-OLS-overridden parameter values from production `au_pac.mod` for the E-SAT VAR ($\lambda_{\pi} = 0.2588$, $\kappa_{\pi} = -0.0336$, $\alpha_{pc} = 0.0006$, $\beta_{pc_m} = -0.0042$, $\rho_{u_gap} = 0.946$, $\text{okun_coeff} = -0.13$, $\rho_{pm} = 0.2230$). The PAC short-run coefficients `b0_*`, `b1_*` use the **L2 OLS estimates verbatim for all five blocks**. (A §6.10 attempt to dampen an IRF oscillation by reverting the housing and employment `b0/b1` to `wp1044` reference was itself reverted in §6.11, which shows the oscillation originates in the trade error-correction block, not the PAC blocks; the PAC `b0/b1` therefore carry their AU L2 OLS values.) Discount factor `beta_pac = 0.98` is unified across blocks.

Following the FR-BDF design choice, the aux files implement **pure VAR-based agents**: agents believe the E-SAT reduced form (with the L2 OLS coefficients), while the simulator runs the post-L2.A structural identities for `yhat_au`, `pi_au`, and `u_gap`. The closed-form policy function for each block is `h_constant +  $\Sigma_k$  h_var_state_k_lag_1 · state_k(-1)`.

Files added: - `dynare/aux/aux_pQ.mod`, `aux_consumption.mod`, `aux_business_inv.mod`, `aux_housing_inv.mod`, `aux_employment.mod`

Files modified: - `dynare/au_pac.mod` lines 555–684: all `h_pac*_constant` and `h_pac*_var*_lag_1` parameter assignments replaced; each block now carries a `regen` comment with the source aux file path.

6.10 PAC oscillation diagnosis and the mixed `wp1044/L2` hybrid for housing inv and employment (2026-05-29) — superseded; see §6.11

Correction (2026-05-30). The diagnosis below misattributed the IRF oscillation to “PAC-block $\leftrightarrow$ demand-identity feedback” and applied a mixed `wp1044/L2` hybrid to the housing-inv and employment `b0/b1` coefficients to shrink the policy-function loadings. **§6.11 supersedes this.** Direct eigenvalue testing shows the dominant 11Q mode is invariant to every PAC lever §6.10 touched (`b3` accelerators, `b0` ECM speeds, policy inertia) and moves *only* with the trade error-correction growth-persistence coefficients `b1_x` and `b1_m_ne`. The §6.10 hybrid changed IRF *amplitude* (it shrank the restored §6.9 PAC loadings) but could not and did not change the *eigenvalues*, so it never removed the oscillation. The hybrid was therefore **reverted on 2026-05-30**: all five PAC blocks again carry their L2 OLS `b0/b1` (§6.9), and the oscillation is cured at its source in §6.11. The section is retained below as a record of the (incorrect) intermediate diagnosis.

The first §6.9 pass — using L2 OLS values for **all** PAC short-run coefficients

including $b0_ih = 0.4956$, $b0_n = 0.3145$, $b3_ih = -0.0728$, $b5_n = -0.0257$ — produced visibly oscillating IRFs for $\Delta \log I_H$, $\Delta \log N$, and (with smaller amplitude) $\Delta \log C$, $\Delta \log I_B$, and headline $\ln_Q$. The dominant oscillation period is $\approx 11Q$, with secondary 13.5Q content.

6.10.1 Diagnosis Eigenvalue analysis of the state-transition matrix from `oo.dr.eigval` identifies two complex pairs that drive the cycle:

Pair	$ \lambda $	period	Mode interpretation
11.03Q complex pair	0.9315	11.0Q	dominant; tied to PAC-block $\leftrightarrow$ demand-identity feedback
13.54Q complex pair	0.9056	13.5Q	secondary; cross-block expectation interaction

Both eigenvalues exist **in both pre-regen and post-regen models with essentially unchanged moduli**. What the §6.9 regeneration did was substantially amplify the *loading* on these modes through the policy-function expansion `oo.dr.ghx`, especially for the housing-inv (`h_pac_ih_var_yhat_au_lag_1` jumped $50\times$) and employment (`h_pac_n_var_yhat_au_lag_1` jumped $5.5\times$) blocks. The IRF amplitude (max – min over Q1–40) under monetary shock changed accordingly:

Variable	Pre-regen	Post-regen (§6.9)	Δ
yhat_au	0.060	0.059	–1%
$\Delta \log C$	0.020	0.021	+5%
$\Delta \log I_B$	0.072	0.072	0%
$\Delta \log I_H$	0.018	0.030	+73%
$\Delta \log N$	0.006	0.009	+54%
π_au, π_Q	0.003, 0.022	0.003, 0.022	$\approx 0\%$

The lever is $b0_*$ (PAC ECM speed), not β_pac (PAC discount factor). Lowering β_pac $0.98 \rightarrow 0.92$ in all aux files changed `h_pac_ih_var_yhat_au_lag_1` only from 0.0586 to 0.0580 — a null result. Setting $b0_ih = 0.12$ (wp1044 reference) in the housing aux file, by contrast, shrinks the same h-vector to 0.0111, a $5.3\times$ reduction. The h-vector formula scales linearly in $b0_*$.

A second mechanism — true contemporaneous simultaneity — operates through the demand identity $yhat_au_t = yhat_au_{t-1} + yhat_dom_t$

where $\hat{y}_{dom,t} = w_c \cdot \Delta \log c_t + w_{ib} \cdot \Delta \log I_{B,t} + w_{ih} \cdot \Delta \log I_{H,t} + \dots$, and the PAC equations each contain a contemporaneous accelerator $b3_k \cdot \hat{y}_{au,t}$ (or $b5_n \cdot \hat{y}_{au,t}$ for employment). Each accelerator coefficient feeds back through the demand identity to $\hat{y}_{au,t}$, creating an algebraic loop with gain $\sum_k b3_k \cdot w_k$. Under wp1044's reference values ($b3_{ib} = 0.69$, $b3_{ih} = 0.50$, $b5_n = 0.13$) the loop is large; under L2 OLS for housing and employment ($b3_{ih} = -0.0728$, $b5_n = -0.0257$ — both wrong-signed and insignificant) the loop is essentially zero.

6.10.2 Solution: mixed wp1044/L2 hybrid for ih and n A clean fix uses **wp1044 reference for b0_* and b1_*** (slow PAC ECM and smooth AR persistence, so the regenerated h-vectors stay small) but **L2 OLS for the contemporaneous accelerator b3_ih and b5_n** (the AU data identify these as wrong-signed and insignificant, i.e., zero; this avoids reintroducing the contemporaneous loop). The full calibration:

Coefficient	wp1044 ref	L2 OLS	Adopted	Rationale
b0_ih	0.12	0.4956	0.12	wp1044 — fast L2 ECM speed amplifies the 11Q mode loading 5×
b1_ih	0.18	0.2934	0.18	wp1044 — for consistency with b0
b3_ih	0.50	-0.0728	-0.0728	L2 OLS — wp1044's b3 reintroduces contemp loop
b0_n	0.07	0.3145	0.07	wp1044 — same logic as housing
b1_n	0.44	0.2950	0.44	wp1044 — for consistency with b0
b2_n, b3_n	0, 0	-0.0278, 0.0261	0, 0	wp1044 depth-1; L2 lags 2-3 were small noise

Coefficient	wp1044 ref	L2 OLS	Adopted	Rationale
b5_n	0.13	−0.0257	−0.0257	L2 OLS — wp1044’s b5 reintroduces contemp loop

This mixed hybrid is philosophically consistent with the §3.6 hybrid-Option-1 pattern already used for the business-investment block: PAC structural deep parameters (ECM speed, AR persistence) come from wp1044; AU L2 OLS overrides apply where the AU data identify economically interpretable estimates and is preserved where the AU data delivers a useful “this channel is zero on AU” finding.

6.10.3 Verification Dynare BK stability preserved (`sdim = 126`, `edim = 5`). The 11Q and 13.5Q complex pairs remain (these are intrinsic to the PAC + demand-identity structure) but their *loadings* are now smaller for the affected blocks. IRF amplitude (max–min over 40Q) under the 100bp monetary shock:

Variable	Pre-regen	§6.9 regen	§6.10 mixed hybrid
yhat_au	0.060	0.059	0.059 (recovered)
$\Delta \log C$	0.020	0.021	0.022
$\Delta \log I_B$	0.072	0.072	0.072 (unchanged)
$\Delta \log I_H$	0.018	0.030	0.012 (−33% vs pre-regen)
$\Delta \log N$	0.006	0.009	0.0035 (−44% vs pre-regen)
π_au, π_Q	0.003, 0.022	0.003, 0.022	0.003, 0.022

$\Delta \log I_H$ and $\Delta \log N$ now have smaller amplitude than even the pre-regen baseline, because the wp1044 b0/b1 produce smoother adjustment than the original calibrated b’s used pre-regen too. The headline-GDP response is preserved ($\ln_Q$ peak −0.139% vs pre-regen −0.135%, within 3%).

6.11 The oscillation is a trade-ECM root, not a PAC-block artefact (2026-05-30)

§6.10 treated a symptom. Re-solving the model in Dynare and decomposing the 126×126 state-transition matrix (`A = oo_.dr.gfx(M_.nstatic+1 : M_.nstatic+M_.nspred, :)`) shows the dominant 11Q complex mode loads almost entirely on the trade-block state variables — `ln_x_level`, `ln_m_ne_level`, `x_gap`, `m_ne_gap`, `ln_ib_level` — while **yhat_au participation is only 0.149**. The PAC blocks are driven *by* the cycle; they are not its source.

6.11.1 The trade error-correction blocks are a momentum oscillator

Both trade quantity blocks (exports `dln_x`, non-energy imports `dln_m_ne`) share the FR-BDF one-step error-correction structure (`au_pac.mod` ~L1353–1410):

```
ln_x_level = ln_x_level(-1) + dln_x           // level = running sum of growth (integrates)
dln_x      = b0_x*x_gap(-1) + b1_x*dln_x(-1) + ... // growth = error-correction + AR momentum
x_gap      = ln_x_eq - ln_x_level             // gap = equilibrium - level
```

Substituting the gap and the level identity into the growth equation collapses the block to a second-order difference equation in the level,

$$L_t = (1 - b_0 + b_1) L_{t-1} - b_1 L_{t-2},$$

whose characteristic roots are complex with **modulus exactly** $\sqrt{b_1}$ when $b_1 > \frac{1}{4}(1-b_0+b_1)^2$. This is the classic momentum-ECM / Samuelson multiplier–accelerator oscillator: growth has inertia, so when the level overshoots its equilibrium the error-correction term reverses growth but the AR momentum carries it past, producing a damped cycle. High growth-persistence b_1 pushes the roots toward the unit circle and the cycle barely damps over the 40-quarter IRF horizon.

The §6.6 audit had replaced the wp1044 trade growth-persistence values ($b_{1,x} = 0.30$, $b_{1,m_ne} = 0.23$) with the AU single-equation OLS estimates ($b_{1,x} = 0.8673$, $b_{1,m_ne} = 0.7427$). Those AU estimates put the roots at $\sqrt{0.8673} = 0.931$ and $\sqrt{0.7427} = 0.862$ — exactly the moduli §6.10 had measured and (wrongly) attributed to the PAC blocks.

Mode	Analytic $\sqrt{b_1}$	Period	Solved eigenvalue	Source
Export ECM	$\sqrt{0.8673} = 0.931$	11.0Q	0.9314, 11.0Q	<code>b1_x = 0.8673</code>
Import ECM	$\sqrt{0.7427} = 0.862$	15.5Q	0.8999, 13.8Q (coupled)	<code>b1_m_ne = 0.7427</code>

6.11.2 The mode is invariant to every §6.10 PAC lever Perturbing each §6.10 candidate and re-solving confirms the dominant mode does not respond to anything in the PAC blocks, and moves only with `b1_x`:

Experiment	Dominant mode after
<code>b3_ib = 0</code> (BI accelerator off)	0.9314 — unchanged
all <code>b3 = 0</code> (no contemporaneous demand loop)	0.9314 — unchanged
<code>lambda_i = 0.50</code> (less policy inertia)	0.9329 — unchanged
all <code>b0 × 4</code> (faster PAC ECM)	0.9316 — unchanged
<code>b1_x → 0.30</code>	mode gone (collapses to $ \lambda \approx 0.55$)

6.11.3 Fix: revert the trade growth-persistence to wp1044; restore §6.9 PAC blocks Two coordinated changes (2026-05-30):

1. **Trade growth-persistence reverted to wp1044 reference** — $b1_x$: $0.8673 \rightarrow 0.30$, $b1_m_ne$: $0.7427 \rightarrow 0.23$. This is the actual fix: it moves the export-ECM root from $\sqrt{0.8673} = 0.931$ to $\sqrt{0.30} = 0.548$ (half-life \$ \$1Q instead of \$ \$10Q) and the import-ECM root to $\sqrt{0.23} = 0.480$. The error-correction speeds $b0_x = 0.299$ and $b0_m_ne = 0.158$ are left at their AU L2 OLS values — they set the cycle *frequency* but not the root *modulus*, and the mode is heavily damped once $b1$ is low. A quarterly export-growth AR(1) of 0.87 is implausibly high and most likely reflects an unremoved trend / near-unit-root in the AU growth series; the full revert ($w = 0$ in the shrinkage $b_1 = w \cdot \text{\texttt{AU}} + (1-w) \cdot \text{\texttt{wp1044}}$) is adopted.
2. **§6.10 mixed hybrid reverted** — because §6.10's premise was wrong, the housing-inv and employment blocks are returned to their §6.9 L2 OLS values ($b0_ih = 0.4956$, $b1_ih = 0.2934$, $b0_n = 0.3145$, $b1_n = 0.2950$, $b2_n = -0.0278$, $b3_n = 0.0261$; $b3_ih$, $b5_n$ were already at L2 OLS), and the $h_pac_ih_*$ / $h_pac_n_*$ policy-function vectors are regenerated to match. This restores the larger (data-identified) PAC loadings the §6.10 hybrid had suppressed.

Verification (100bp monetary shock, 40Q; Dynare BK preserved, $sdim = 126$, $edim = 5$). The trade fix removes the oscillation at its source; direction reversals over Q1–40 fall from ≈ 7 (≈ 3.5 sustained cycles) to 1–3 (a normal hump):

Variable	Turning pts before	Turning pts after
$\ln_Q$	7	2
$yhat_au$	7	2
$\Delta \log I_B$	7	3
$\Delta \log I_H$	7	2
$\Delta \log C$	7	3
$\Delta \log N$	7	4

Because the §6.10 amplitude suppression is also removed, the PAC growth-rate responses are larger than the §6.10 table reported (e.g. $\Delta \log I_H$ trough -0.049% at Q4 vs §6.10's -0.020%), but they no longer oscillate. The headline monetary transmission to real activity is materially *weaker and slower* than the §6.10 calibration: the trade-block growth-persistence collapse changes net-trade dynamics, so the $\ln_Q$ trough shrinks from -0.139% (Q9) to **-0.039% (Q17)** and real GDP drifts to $+0.076\%$ by Q40. (Note: the abstract, §6.8, and §7.2.2 still quote the pre-§6.11 headline magnitudes and have not yet been propagated.)

The $b1_x = 0.30$, $b1_m_ne = 0.23$ values here were a blunt first pass (full revert to wp1044). §6.12 supersedes them with a *stability-constrained estimate* $b1_x = b1_m_ne = 0.49$ after establishing that the high AU values are genuine. Table 6.3 reflects §6.12.

6.12 Stability-constrained estimate of the trade growth-persistence (2026-05-30) — superseded by §6.13

Superseded. §6.13 shows the high $b1$ this section worked around was itself an artefact of estimating on the ABS *Trend* (smoothed) series. On Seasonally Adjusted data $b1_x = 0.09$, so no constraint is needed. The robustness evidence below (that $b1 \approx 0.78$ is stable across specs/samples) is correct *conditional on the Trend series*; it does not survive the switch to SA. Retained as the record of how the constraint was derived before the data bug was found.

§6.11's full revert to wp1044 ($b1_x=0.30$) discards AU signal arbitrarily (it is $w=0$ in the shrinkage $b_1 = w \cdot AU + (1-w) \cdot wp1044$ with no basis for w). Before accepting it we asked: **is the high AU $b1$ real, or a specification artefact?**

6.12.1 The high AU export-growth persistence is genuine $b1_x$ was re-estimated under seven specifications and five sub-samples (data/pac_blocks/estimate_exports_*.m):

Specification	$b1_x$	$\sqrt{b1}$
HP-detrended level (original §6.6 audit)	0.867	0.931
Consistent cyclical (model's $dln_x = \Delta$ cyclical level, not raw growth)	0.802	0.895
Cointegrating target vs OECD trading-partner GDP	0.742	— (no cointegration)
Piecewise-linear trend, K=0...3 breaks (data-driven: 2003Q3, 2011)	0.69–0.80	—
+ short-run foreign-demand-growth term $\Delta ln y_world$ (sig., t=2.3)	0.749	0.865

5 sub-samples (pre/post-GFC, mining-boom core, excl-GFC)	0.77–0.79	~0.88
---	-----------	-------

Two findings of independent value emerged: (a) the original estimation mixed an HP-detrended *level* (for the EC term) with *raw* dln_x (LHS/AR), inconsistent with the model’s cyclical dln_x — the consistent spec is preferable; (b) a **data-driven piecewise trend with breaks at 2003Q3 and 2011** (mining-boom onset and the investment-peak→production-ramp transition) strengthens the error-correction speed b0 from insignificant to $\text{t}=-8.9$, and a **short-run foreign-demand-growth term** (the analog of imports’ *iad*, missing from the export equation) is significant — both are spec improvements worth folding in later. **But none lowers b1 below ~ 0.7 .** AU exports do *not* cointegrate with foreign demand (OECD residual $\text{AR}(1)=0.97$; $\beta=2$ but unit-root residual), because AU export-*volume* growth is **supply/capacity-driven** (iron-ore/coal/LNG capacity coming online in discrete ramps), not foreign-demand-driven — so no demand-side target can absorb the momentum. $\text{b1}_x \approx 0.78$ is one of the most robust facts in the model.

6.12.2 The constrained estimate: project the OLS onto the IRF-stability frontier Since b1 is genuine and the oscillation modulus is $\sqrt{\text{b1}}$ (invariant to b0 , the target, and the trend — §6.11.1), the principled resolution is a **dynamic-stability constraint** (the device already used for the wage Phillips, §6.7): keep the data’s b1 capped at the largest value giving a non-oscillatory GE IRF. Mapping turning points (100bp monetary, 40Q) over a grid of the common cap $c = \text{b1}_x = \text{b1}_{m_ne}$:

c (= b1)	$\sqrt{c}$	$\ln_Q$	yhat_au	$\Delta \log$ I_B	$\Delta \log$ I_H	$\Delta \log$ C	$\Delta \log$ N	$\ln_Q$ trough
0.30	0.55	2	2	3	2	3	2	−0.036%
(§6.11)								
0.40	0.63	2	2	3	2	3	4	−0.038%
0.49	0.70	2	2	3	2	3	4	−0.043%
0.56	0.75	2	2	3	2	5	4	−0.052%
0.65	0.81	2	2	3	6	5	4	−0.057%

The headline $\ln_Q/\text{yhat_au}$ stay clean across the whole grid, but the **PAC growth-rate IRFs bind first** — higher trade momentum leaks into $\Delta \log \text{C}$ (oscillates at 0.56) and $\Delta \log \text{I_H}$ (at 0.65) through the demand identity. The largest cap keeping *every* variable non-oscillatory is $\sqrt{\text{b1}} \leq 0.70$, i.e. $\text{b1}_x = \text{b1}_{m_ne} = 0.49$ — the projection of the joint OLS estimate (0.79, 0.74) onto the all-variable IRF-stability frontier. (A symmetric 0.65 cap fails: it

cleans `ln_Q` but leaves $\Delta \log I_H$ oscillating, and raising imports alone to 0.65 reintroduces the 14Q import mode.)

Adopted: `b1_x = b1_m_ne = 0.49`. This retains genuine AU export/import momentum (vs the over-damped 0.30/0.23), deepens the real-GDP trough to **-0.043% (Q11)** from §6.11's -0.039%, and keeps all IRFs non-oscillatory (turning points 2-3; BK preserved `sdim=126`, `edim=5`). It is defensible precisely because the unconstrained estimate is robust and the constraint binds at a clear, demonstrated threshold — no arbitrary `w`.

Queued structural follow-up. The deeper fix is to split exports into resource vs non-resource (mirroring the import energy/non-energy split): resource volumes are capacity-driven (high `b1`, belong in a trend/capacity process, not a demand-ECM), non-resource are demand-driven (lower `b1`). That would remove the oscillation structurally and let the cap be dropped — see `NEXT_PROJECT_export_resource_split.md` (needs ABS 5368 + commodity deflation).

6.13 Root cause: the trade block used the ABS *Trend* series, not Seasonally Adjusted (2026-05-30) — resolves §6.6/§6.11/§6.12

Pursuing the §6.12 mining/non-mining split (which requires goods-level volumes) surfaced the actual root cause. Pulling ABS Balance-of-Payments goods chain-volume exports (cat. 5302.0 Table 6, incl. SA series A3535046J non-rural goods) and estimating `b1` per component gave a uniform surprise: **every goods-export-volume component has $b1 \approx 0$** — mining (non-rural) -0.06, rural 0.19, total goods 0.04 — nothing like the model's 0.74. That is, the mining hypothesis was rejected (mining-volume *growth* is anti-persistent quarter-to-quarter; the capacity ramps are low-frequency trend), and the high `b1` is not in goods volumes at all.

Tracing the discrepancy: **the model's export and import volume series were the ABS *Trend* (Henderson-filter-smoothed) chain-volume series, not Seasonally Adjusted.** `data/prepare_trade_price_data.m` hardcoded `ix_x = 39`, `ix_m = 40`, `ix_g = 6` into `abs_5206_vol.xlsx`, and columns 39/40/6 are the Trend block (series A2304238J / A2304239K / A2304206R). A Trend series is a centred moving average; its $AR(1) \approx 0.74$ is mechanical smoothing, not economic momentum. Estimated on the **Seasonally Adjusted** column (121, A2304114F) the same series gives `b1 = 0.008`; on Original (col 206), -0.07. **This single data-selection error is the source of the entire IRF oscillation:** $\sqrt{0.74} \approx 0.86$ put the trade-ECM root near the unit circle. The §6.6 audit (`b1_x=0.87`), §6.11 revert (0.30) and §6.12 stability constraint (0.49) were all treating the symptom.

Fix. Switched the volume columns to Seasonally Adjusted — `ix_x = 121` (exports, A2304114F), `ix_m = 122` (imports, A2304115J), `ix_g = 86` (govt FCE, A2304080V) — rebuilt `trade_price_data.mat`, and re-estimated. The defla-

tor/IPD columns were audited and are correctly SA already; only the three volume series were affected. New SA OLS (written back verbatim per the OLS-over-calibration convention):

Coef	Trend (bug)	SA OLS (adopted)	t	$\sqrt{b1}$
b0_x (export ECM speed)	0.299	0.363	−4.4	—
b1_x (export Δ -persistence)	0.867	0.092	0.97	0.30
b0_m_ne (import ECM speed)	0.158	0.309	−4.6	—
b1_m_ne (import Δ -persistence)	0.743	0.185	1.87	0.43

Result. $\sqrt{b1} = 0.30 / 0.43 \rightarrow$ the trade-ECM roots are well inside the unit circle, so the oscillation is gone *with no constraint*. The 100bp-monetary IRF has turning points of **1–2 across every variable** (the cleanest of any calibration), BK preserved (sdim=126, edim=5), and the real-GDP trough is **−0.072% (Q11)** — deeper and more RBA-suite-plausible than the over-damped §6.11/§6.12 values, because the well-identified SA ECM speeds (b0_x=0.36, b0_m=0.31) are stronger than the Trend-data estimates. This **supersedes §6.11 and §6.12** (the wp1044 revert and the stability constraint are no longer needed). The **ln_Q** Q40 positive drift (+0.078%) is unchanged — it is the separate CES trend-labour **rw_gap** hysteresis (IRF_TRANSMISSION_DRIFT_INVESTIGATION.md), independent of the trade block. The mining/non-mining split is no longer required for the oscillation, though the audit incidentally confirmed all goods-volume components are low-persistence.

6.14 Full-model seasonal-adjustment data audit (2026-05-30)

The §6.13 bug (a Trend series read where Seasonally Adjusted was intended) prompted a full audit of every series feeding the model, checking the ABS “Series Type” metadata of each hardcoded-column pick and the SA status of each FRED series. **One further instance of the same class was found and fixed; everything else is SA or a benign/documented exception.**

- **Found:** data/prepare_supply_data.m selected employment, the unemployment rate, and the labour force via a cols(3) “3rd = SA” heuristic, but the occurrence order in abs_6202_labour_force.xlsx is Trend(1)/SA(2)/**Original(3)** — so it read the **Original (NSA)** series. Fixed to the robust pick_sa_col helper (reads the Series-Type row), as hours and WPI already did.
- **Impact: negligible**, unlike §6.13. The corrected SA unemployment rate feeds the wage Phillips and u_gap; re-estimating moves the constrained wage-Phillips slope only κ_w : 0.359 $\rightarrow$ 0.343 (and γ_w : 0.857 $\rightarrow$ 0.863, λ_w : 0.093 $\rightarrow$ 0.087) — AU unemployment-rate seasonality is mild and enters as one regressor. The CPI Phillips, deflators, and

market-VA blocks are unchanged (their inputs were already SA). The employment PAC is unaffected (it uses the SA FRED `au_employment`). The 100bp-monetary IRF is materially identical to §6.13 (real-GDP trough -0.073% Q11 at this step, turning points 1–2, BK preserved); the subsequent employment- χ / housing / consumption-SA ECM-speed writebacks (the “Full re-estimation check” below) then deepen the trough to -0.144% Q11 (§7.2).

Audit result (all series feeding the model): seasonally adjusted everywhere after the fix — ABS 5206 trade volumes (§6.13), ABS 6202 employment/unemployment/labour-force (§6.14), hours, WPI, industry GVA (market/non-market, via a self-checked +109 SA offset), all IPD deflators, household income (by SA series ID), and the FRED activity/labour series (OECD “S” suffix). **Benign exceptions, flagged not fixed:** `abs_6302_awe.xlsx` is Trend-only (but AWE is unused — the wage Phillips uses the SA WPI); `abs_5204_*` (compensation, capital stock, productivity) are Original-only low-frequency supply series with no quarterly seasonality; FRED CPI is NSA but feeds only a legacy synthetic-ULC proxy, not the WPI-based wage Phillips; the 10-year bond yield is a rate (SA N/A). **Durable lesson:** when reading ABS time-series spreadsheets by hardcoded column index, always verify against the “Series Type” metadata row — Original / Seasonally Adjusted / Trend blocks sit in different columns and silently mis-selecting Trend injects spurious persistence (§6.13) or Original injects seasonality (§6.14).

Full re-estimation check (resolved, written back 2026-05-31). With the data layer rebuilt on the corrected SA series, all PAC blocks were re-estimated. The VA-price block is identical (its inputs were already SA, reproducing `b0_pQ` = 0.2581 to 4 d.p.); consumption, employment, and housing shift moderately because the rebuilt `12_data_layer` carries the SA unemployment/employment through `u_hat/n_hat_star` (housing `b0_ih` 0.50→0.60, consumption `b0_c` 0.27→0.23, employment `b0_n` 0.31→0.48 — the latter also reflecting the exact- χ fix of §6.9). These re-estimates were **initially deferred** (the base `dataset.csv` had been recovered from git history, and a consumption Δc -lag clamp seemed to confound the SA effect). The deferral was **reversed** after (i) confirming the data layer regenerates **deterministically** from `dataset.csv` (max-diff 0 across `c_star/piQ/n_hat_star_S`), which resolves the provenance concern, and (ii) recognising that the Δc -lag clamp is constant across the pre-SA and SA runs, so the shift reflects the SA-data correction rather than the clamp. The SA-corrected estimates are now the production values, consistent with the rest of the model using SA data. The headline IRFs (Table 6.3) deepen slightly (real-GDP trough to $\approx -0.14\%$ from the faster housing/employment ECM speeds) and remain clean hump shapes reverting at Q200.

7. Model Properties

Scope note. All IRFs and figures in this section are computed under **equation-by-equation OLS calibration** following the wp1044 methodology (not Bayesian MCMC). PAC block parameters are estimated by iterative OLS on block-specific structural equations; shock standard deviations are calibrated from OLS residuals. The CES calibration of Table 4.2.1 and the trade ECM of §4.8 are unchanged. Peaks and troughs quoted in the text are read from the saved IRF binaries; the corresponding figures are reproduced from the same artefacts.

Calibration note. The production model uses the **hybrid Option 1 calibration** of §3.6 + §4.6.2 (b0_ib=0.096, b1_ib=0.33, b2_ib=0.11, b3_ib=0.69, $\omega=0.35$ — all imported from wp1044 Table 3.5.13), plus ULC/UCK structural channels, deflator ECMs, endogenous spread widening, and energy/non-energy import split. Table 6.3 and the §7.2.2 channel-by-channel walkthrough reflect IRFs from this OLS production specification.

7.1 Steady state

At the balanced growth path (gap model, all growth rates = 0):

Table 6.1: Steady-state values

Variable	SS value (quarterly)	Annual equivalent	Interpretation
yhat_au, yhat_us	0	—	Output gaps closed
pi_au, piQ, pi_w, all pi_j	0.625%	~2.5%	Inflation = RBA target
pi_us	0.500%	~2.0%	Inflation = Fed target
i_au	1.049%	~4.2%	RBA neutral rate
i_10y	1.349%	~5.4%	10Y bond yield
wacc	1.834%	~7.3%	Weighted avg cost of capital
uc_k	1.859%	~7.4%	User cost of capital
i_lh	1.749%	~7.0%	Mortgage lending rate
All gaps, growth rates	0	—	Balanced growth

7.1.1 Long-run convergence to balanced growth path To verify that the model is well-behaved under no shocks — i.e., that the gap variables converge to zero and the level variables follow their deterministic trends — we run the linearised decision rule forward from three off-steady-state initial conditions for 80 quarters with all shocks set to zero. Because the model is solved to first order, an off-SS initial state is mathematically equivalent to a sum of impulses, so the convergence paths can be visualised directly from the model’s long-horizon IRFs

(Figure 6.0). Conditions: boom-start (+1 s.d. output-gap impulse), commodity-cycle start (+1 s.d. eps_pcom impulse), and tight-money start (+1 s.d. policy-rate impulse).

Figure 6.0: AU-PAC long-run convergence under hybrid expectations. Output-gap, inflation, real-activity, and demand-component gaps all decay smoothly to zero within 20–40 quarters; the policy rate, 10Y yield, and exchange rate exhibit longer-lived but monotonic adjustment. No cyclical instability or divergence is observed under any of the three starting conditions, confirming BK stability and convergence to the balanced growth path.

7.2 Monetary policy transmission

Following the wp736 Section 6.2 protocol, we assess the impact of a 100 basis point annualised monetary policy tightening (= 0.25 quarterly percentage points). At order = 1, IRFs scale linearly: we compute them from Dynare’s 1 s.d. responses multiplied by $0.25/\sigma_{\varepsilon^i}$, where $\sigma_{\varepsilon^i} = 0.111$ is the posterior mean of the monetary-policy shock standard deviation.

The walkthrough below describes the **Phase T single-regime architecture** documented in §3.4 — the current production model in `au_pac.mod`. Figures 7.1–7.2 below visualise three IRF traces (VAR / Hybrid / MCE) from the earlier three-regime workflow; they are kept as documentation of the architectural evolution. The numerical values in Table 6.3 below and the channel-by-channel walkthrough in §7.2.2 reflect the Phase T (formerly “Hybrid”) trace — the regime that is closest to the current single-regime policy-function expectations.




Figure 5: Monetary policy IRFs

7.2.1 Impulse responses *Figure 6.1: AU-PAC responses to a 100bp annualized monetary policy tightening, all key variables. Shock = 0.25 qpp impact on the cash-rate residual ε_i . Three traces visualised — VAR (backward-looking-only expectations), Hybrid (3 forward-looking recursive objects + backward PAC), MCE (model-consistent expectations across all PAC equations) — from the earlier three-regime workflow; the Phase T production regime is closest to the Hybrid trace.*

Table 6.3: Peak IRF (100bp annualized monetary tightening) — current calibration (post-§6.13/§6.14 SA-data fixes + 2026-05-31 employment- χ / housing / consumption-SA ECM-speed writebacks)

GDP and growth-rate variables in % deviation from baseline (FR-BDF convention); inflation and yield variables in annualised pp. Peaks taken within the first 40 quarters. Calibration: post-Phase-L2 single-equation OLS for CPI Phillips, trade quantities, trade and other deflators (§6.6), BK-constrained restricted OLS for wage Phillips (§6.7), Phase L2.A architectural fixes for $yhat_au$, pi_au , and u_gap (§6.8), all five h_pac_* policy-function blocks regenerated against the post-L2.A VAR (§6.9), the §6.10 mixed hybrid **reverted** so all PAC blocks carry their L2 OLS $b0/b1$, and the **§6.13 trade-data fix**: export/import volumes switched from the ABS *Trend* (smoothed) series to *Seasonally Adjusted*, with the trade short-run coefficients re-estimated on SA data ($b0_x=0.363$, $b1_x=0.092$, $b0_m=0.309$, $b1_m=0.185$; written back verbatim). This removes the IRF oscillation at its source — $\sqrt{b1} = 0.30/0.43$, well inside the unit circle — with **no stability constraint** (§6.11/§6.12 superseded). For activity variables the table reports the (negative) transmission **trough**; $ln_Q/yhat_au$ additionally exhibit a slow positive drift over the back half of the horizon (noted below).

Variable	Peak	Quarter
Real GDP (ln_Q)	−0.144% (trough)	Q11
Output gap ($yhat_au$)	−0.086% (trough)	Q11
CPI inflation (pi_au , y/y)	−0.004 pp	Q3
pi_au_food (decomp)	+0.020 pp	Q1
pi_au_energy (decomp)	−0.054 pp	Q9
pi_au_core (decomp; $\approx$ pi_c)	+0.001 pp	Q12
VA price inflation (piQ , y/y)	+0.030 pp	Q1
Consumption growth (Δlog)	−0.036%	Q1
Business inv growth (Δlog)	−0.062%	Q8
Housing inv growth (Δlog)	−0.051%	Q4
Employment growth (Δlog)	−0.010%	Q4
Wage inflation (y/y)	−0.008 pp	Q3

Variable	Peak	Quarter
Unemployment gap (u_gap, pp)	≈ 0 (not retained in regen IRF set)	—
Exchange rate (s_gap, %)	−0.985%	Q8
10Y yield (annualised)	+0.105 pp	Q1

See §6.8 (Table 6.9) for the pre-vs-post-L2.A architectural comparison; §6.9 for the policy-function regeneration; §6.10 for the (superseded) PAC-oscillation diagnosis; §6.11/§6.12 for the trade-ECM root cause and the (now-superseded) stability constraint; and **§6.13** for the SA-data fix that produced this table. The IRF oscillation is gone with **no constraint** — direction reversals over 40Q are **1–2 for every variable** (the cleanest of any calibration; vs ≈ 7 before §6.11). The real-GDP trough is **−0.144% (Q11)** on the current calibration — deeper than the §6.13 SA-fix value (−0.072%) after the 2026-05-31 employment- χ / housing / consumption-SA ECM-speed writebacks raised the real-side adjustment speeds (b0_n 0.31→0.48 under the exact depth-3 χ , b0_ih 0.50→0.60), and comfortably within the RBA-suite range. **ln_Q** then drifts to +0.058% and **yhat_au** to +0.047% by Q40 — the separate CES trend-labour **rw_gap** hysteresis (IRF_TRANSMISSION_DRIFT_INVESTIGATION.md), monotone after the trough, not a cycle, and unaffected by the trade-data fix.

Comparison with pre-audit values: real-GDP / output-gap / s_gap / 10Y are essentially unchanged; the **CPI channel weakens further** (pi_au peak shifts to Q16 with −0.020 pp from earlier −0.038 pp at Q9; pi_c becomes effectively flat) because the AU CPI Phillips OLS identifies only persistence ($\lambda\pi = 0.26$, $t = 2.76$) and finds output-gap, VA-price, and import-price channels insignificant. **Consumption-growth response on impact is larger** (−0.036% at Q1 on the current calibration vs −0.004% at Q7 pre-audit) because the orphan **b_di_c·di_gap** calibration was zeroed (it had been amplifying consumption via a channel not in wp1044 Eq 35); the remaining response now reflects only OLS-identified channels (pv_r_1h_gap, b_HtM, PAC growth-neutrality term). **The wage Phillips response is several× larger** than pre-audit (−0.008 pp at Q3 vs −0.002 pp earlier) because the BK-constrained wage Phillips OLS (§6.7) identifies a substantial Phillips slope $\kappa_w = 0.359$ and a near-unit CPI passthrough $\gamma_w = 0.857$.

7.2.2 Channel-by-channel walkthrough The aggregate **−0.144% real GDP** trough (Q11), **−0.086% output gap** trough (Q11), and **−0.004 pp CPI inflation** decline (Q3) decompose across six transmission channels documented in Table 2.2. Under the equation-by-equation OLS calibration (wp1044 methodology), the direct consumption rate channel is near-zero (the L2 SA estimates give a HtM coefficient $\beta_2 = +0.027$ and an impact-rate coefficient $\beta_3 = -0.007$ — AU data shows no significant direct interest-rate sensitivity in consumption under the wp1044 spec), and monetary transmission operates primar-

ily through the exchange rate (-0.99%) and investment (business-investment growth -0.062% , housing-investment growth -0.051%).

(i) Term structure and WACC. The 10Y yield rises **+0.105 pp on impact** because the term-structure NPV (eq. 48) $i_{10Y,t} = (1 - \kappa_{10})i_t + \kappa_{10} \cdot E_t[i_{10Y,t+1}] + s_{10Y,t}$ with $\kappa_{10} = 0.97$ internalises the entire expected mean-reverting cash-rate path on impact. The 10Y yield enters the weighted average cost of capital (eq. 49) via the term-debt weight, lifting WACC by approximately 60–80 bp at impact under the AU calibration; the remaining cost-of-capital absorption goes through the bank-lending and equity-cost spreads.

(ii) User cost of capital and business investment. WACC enters the user-cost identity (eq. 27) $UC_t^K = WACC_t + \delta_K - (\pi_t^{IB} - \pi_t^Q)$. The higher user cost lowers the desired capital stock through the CES first-order condition with elasticity $\sigma_{CES} = 0.5366$. The business-investment PAC equation (eq. 28) then propagates this through the long-run target $-\sigma_{CES} \cdot pv_{rKB,aux}$ and the accelerator $b_3^{ib} = 0.69$ (post-hybrid, wp1044 Table 3.5.13; was 0.32 pre-hybrid) via the depressed output gap. Under the current production calibration (Option 1, §4.6.2) the combined effect is a **-0.062%** peak in $\Delta \ln I^B$ at Q8. The structural form (coef = +1 on $PV(\Delta \hat{q}) + PV(\Delta \bar{q})$, coef = $-\sigma$ on $PV(\Delta \log \hat{r}_{KB}) + PV(\Delta \log \bar{r}_{KB})$) is unchanged; only the deep parameters are now wp1044-imported. See §4.6.2 and §5.3 for the rationale.

(iii) Mortgage rate and housing investment. Australia’s variable-rate mortgage market passes cash-rate changes through quickly to i^{LH} (eq. 56) with a spread $\rho_{lh} = 0.40$ above the 10Y yield. The household-investment target (eq. 31) includes a mortgage-rate term $-\kappa_{mort}(i^{LH} - i_{SS}^{LH})$ with $\kappa_{mort} = 0.048$, plus a separate house-price (Tobin’s Q) channel. The housing-investment PAC equation (eq. 32) absorbs these via its long-run target and the cyclical $b_3^{ih} = 0.22$ accelerator. Net peak in growth: $\Delta \ln I^H = -0.051\%$ at Q4. The level ($\ln_{IH}$) continues to decline through the 40-quarter horizon — the longest-lagged response in the demand block, consistent with Australia’s out-sized housing-construction cycle.

(iv) Exchange rate (forward-NPV UIP). UIP (eq. 38) is written as $s_gap = \rho_s \cdot s_gap(-1) - \alpha_s \cdot pv_i_uip + \alpha_s \cdot (\pi_au_gap - \pi_us_gap) + \varepsilon_s$, where pv_i_uip is the forward-NPV of the policy-rate gap defined recursively as $pv_i_uip = (i_au - ibar) + \beta_uip \cdot pv_i_uip(+1)$ with $\beta_uip = 0.92$. The forward recursion produces $pv_i_uip(0) \approx 1/(1 - \beta_uip \cdot \lambda_i) \cdot i_gap(0) \approx 4.55 \cdot i_gap(0)$ on impact — the spot AUD internalises the entire expected rate path immediately. Result: s_gap peaks at **-0.99% at Q8**. Through eqs. (40)–(41) the appreciation depresses export volumes via the trade elasticity b_2^x and dampens imports — net trade contributes a small early-quarter negative to GDP.

(v) Permanent income, real-rate NPV, and consumption. Under the L2 equation-by-equation OLS calibration (SA layer), the consumption block’s direct interest-rate coefficients are near zero: the HtM coefficient $\beta_2 = +0.027$

and the impact-rate coefficient $\beta_3 = -0.007$ — AU data shows no significant direct interest-rate sensitivity in consumption under the wp1044 spec. The combined peak consumption-growth response is **−0.036% at Q1** — small, and it is a general-equilibrium income effect (via the now-amplified output gap feeding the permanent-income target) rather than a direct rate response. This remains the headline finding of the OLS calibration: Australian consumption does not respond *directly* to interest rates in the wp1044 structural equation, and monetary transmission therefore operates primarily through other channels (the exchange rate and business investment). The near-zero direct response contrasts sharply with the earlier Bayesian MCMC calibration (which gave -0.131% at Q2 under Bayesian-regularised rate coefficients); the OLS result reflects what AU data actually identifies without prior regularisation.

(vi) Wage–price spiral. Wage inflation responds negatively, with peak -0.008 pp at Q3, driven by the CPI-passthrough channel $\gamma_w \cdot \pi^c$ ($\gamma_w = 0.35$) as CPI inflation falls. The unemployment-PV channel $-\kappa_w \cdot pv_{u,gap}$ (with the Phillips slope $\kappa_w = -0.10$ corresponding to a positive structural Phillips slope under the FR-BDF sign convention) contributes a small reinforcing negative. VA-price inflation (piQ) **jumps UP +0.032 pp on impact (Q1)**, reflecting the ULC cost-push channel: higher interest rates raise unit labour costs through the user-cost-of-capital channel, and the ULC pass-through with $\gamma_{ulc} = 0.295$ transmits this into VA prices before the demand-side deflationary forces dominate at longer horizons. The CPI deflator inflation (pi_c) is effectively flat (peak $+0.001$ pp at Q12) — the flat-Phillips passthrough leaves core consumer prices barely moved.

The channels together transmit monetary policy primarily through the exchange-rate channel (-0.99%) and the investment cost-of-capital channels (housing-investment growth -0.051% , business-investment growth -0.062%). The direct consumption rate channel is near-zero under OLS calibration (the -0.036% consumption-growth trough is a general-equilibrium income effect, not a direct rate response), a striking contrast to earlier Bayesian-regularised results and to DSGE models where intertemporal substitution in consumption is typically the dominant transmission mechanism. AU’s variable-rate mortgages and the strong CPI-indexation wage Phillips shape the relative magnitudes of the non-consumption channels.

7.2.4 Comparison with the RBA’s model suite A useful external benchmark for AU-PAC’s monetary-transmission results is the RBA’s own model suite, surveyed in Mulqueeney, Ballantyne and Hambur (2025, RBA Bulletin, April). That paper reports the response to a 100 bp cash-rate increase across four Australian models — Beckers (2020) VAR, MARTIN (the RBA’s semi-structural workhorse), DINGO (the RBA’s DSGE), and Murphy (an external benchmark) — and decomposes the transmission into four channels: exchange rate, asset prices / wealth, savings and investment (i.e., intertemporal substitution), and cash flow.

Table 6.4: AU-PAC vs RBA model suite, 100bp cash-rate increase

Peak GDP / output-gap and inflation responses across models. Mulqueeney et al. (2025) report year-ended inflation effects in percentage points; AU-PAC’s quarterly CPI peak is converted to a year-ended equivalent by multiplying by 4.

Model	Peak GDP fall	Peak quarter	Peak inflation fall (y/y, ppt)
Beckers (2020) VAR	-1.5%	~Q4	-0.40 to -0.50
MARTIN (semi-structural)	-0.45 to -0.50%	Q4	-0.10 to -0.15
DINGO (DSGE)	-0.60 to -0.75%	Q5-7	-0.20
Murphy (external)	-0.30%	Q4	-0.40 to -0.50
AU-PAC (OLS calibration, this paper)	-0.144%	Q11	-0.015

There are three plausible reasons AU-PAC’s peak magnitude sits at the bottom of the RBA suite, the first reinforced by the FR-BDF 2026 calibration adopted here:

1. **CES production with $\sigma = 0.54$.** The AU calibration under the FR-BDF 2026 labour-FOC method gives $\sigma = 0.5366$ — modestly below the 0.5–0.7 range typically used in DSGE calibrations of the same vintage. Importantly, the *higher* σ (vs the previous AUSAAC calibration of 0.34) actually delivers a *smaller* peak IRF: firms reallocate factors more readily, so factor prices need to move less to clear a given demand shortfall. The mechanism is well-known but its direction is often misread — higher CES elasticity dampens, rather than amplifies, the output response to a relative-price shock when the wage-price-spiral feedback is small ($\kappa_w = 0.09$ is small for AU).
2. **Estimated PAC coefficients remain moderate in AU data.** The L2 OLS estimates (Table 5.6) give error-correction speeds in the ≈ 0.23 –0.60 range (consumption 0.23, VA-price 0.26, employment 0.48 under the exact depth-3 χ , housing 0.60), AR coefficients in the 0.09–0.33 range, and accelerator coefficients ($b3_{ib} = 0.69$ imported, $b3_{ih} = 0.22$, $b5_n \approx 0$). The 2026-05-31 employment- χ / housing writebacks raised these speeds (and correspondingly deepened the real-GDP trough from -0.098% to -0.144%), but PAC adjustment-cost smoothness still dampens responses relative to DSGE / VAR architectures. Additionally, the consumption block’s near-zero direct rate sensitivity under OLS ($\beta_2 = +0.027$, $\beta_3 = -0.007$) removes the strongest demand-side channel, further reducing the aggregate response.

3. Linearised gap formulation. AU-PAC’s output is $y_{\text{hat_au}}$, the *gap* from a slowly-moving balanced-growth path. RBA models report responses against the *baseline level*, including the slow-moving potential trajectory the model would have followed absent the shock. The peak-gap reading is therefore a lower bound on the cumulative GDP impact.

Timing. RBA models report peak effects “around one to two years” (Q4–Q8). AU-PAC’s output-gap and real-GDP troughs sit at Q11 — just beyond the upper end of this range, reflecting the slower hump under the SA-data calibration. The peak response timing is broadly comparable to MARTIN/DINGO’s horizon.

Exchange-rate channel weight. Mulqueeney et al. (2025, Graphs 6–7) report that 25–67% of the GDP response and one-third to two-thirds of the inflation response come from the exchange-rate channel in MARTIN / DINGO. AU-PAC’s UIP channel works in the same direction (AUD appreciation on impact) and the magnitude is broadly in line with the RBA suite: AUD appreciation peak -0.99% at Q8. Under the OLS calibration, the exchange-rate channel is the dominant transmission mechanism (with near-zero consumption response), making the exchange-rate channel weight in AU-PAC even larger proportionally than in the RBA models. With $\rho_s = 0.775$ (half-life ≈ 3 quarters) AU-PAC’s mean-reversion to PPP is faster than typical estimates, but the impact magnitude — achieved on Q1–Q8 through the $p_{v_i_uip}$ forward NPV — is comparable to MARTIN/DINGO.

Housing channel. The RBA paper notes that “housing is a sensitive part of economic activity”, and the dwelling-investment contribution to MARTIN’s peak GDP fall is sizeable. AU-PAC’s housing-investment growth IRF is one of the larger demand-side components and the level ($\ln_IH$) continues to decline through the 40-quarter horizon — a long-tailed response shape rather than a sharp peak. Two factors explain the lagged level dynamics: the housing PAC has higher-order AR dynamics with strong adjustment cost; and the Tobin’s Q feedback coefficient $b_{ph,ih} \approx 0$ in AU data, so the house-price channel that RBA models likely include is muted in AU-PAC.

Cash-flow channel. Mulqueeney et al. (2025) find the cash-flow channel is small in aggregate (MARTIN: “small but faster than the savings / investment and asset prices channels”) because borrower and saver responses partially offset. Under the OLS calibration, AU-PAC’s consumption direct rate-sensitivity coefficients are near zero ($\beta_2 = +0.027$, $\beta_3 = -0.007$), so the cash-flow / intertemporal-substitution channel through consumption is effectively absent — even more consistent with the RBA’s finding of a small aggregate cash-flow channel than the earlier Bayesian-regularised calibration.

Verdict. AU-PAC’s output-gap trough (-0.086% at Q11) is the smallest peak demand response to a 100 bp tightening among comparable Australian models (vs Murphy’s -0.30% and the RBA workhorses’ -0.45% to -1.5% in GDP terms). The real GDP trough (-0.144% at Q11) is also at the conservative end of the suite. The trough timing (Q11) is at/just beyond the MARTIN/DINGO

short-horizon range. The smaller AU-PAC peak magnitude reflects three factors: PAC adjustment-cost frictions (Tinsley 1993, FRB/US tradition); the equation-by-equation OLS calibration that lets AU data speak without Bayesian regularisation; and the resulting near-zero direct consumption rate sensitivity ($\beta_2 = +0.027$, $\beta_3 = -0.007$), which removes the consumption channel that is typically the largest demand-side transmission mechanism in DSGE and semi-structural models. AU-PAC therefore sits within the robust-policymaking spirit of “diversity supports more robust policymaking” (Mulqueeney et al. 2025, p. 11) rather than aiming to match the average response. Refinements that would narrow the gap include adopting the broader FR-BDF 2026 financial-block extensions (NFC accelerator, DSR-based mortgage block) to amplify the cost-of-capital channel, and revisiting the consumption equation with external high-frequency RBA OIS-surprise data for rate-channel identification.

7.3 Impulse responses to other shocks

All shocks are sized at policy-relevant magnitudes rather than 1 s.d. At order = 1 (linear), IRFs scale exactly by the ratio $\text{target} / \sigma_{\text{shock}}$.




Figure 6: Monetary policy shock

7.3.1 Monetary policy shock (eps_i) — 100bp annualized *Figure 6.3: 100bp annualized monetary policy tightening (scale = $0.25 / \sigma_i = 0.25/0.1105 \approx 2.26$).*

The exchange rate appreciation ($s_gap -0.99\%$ at Q8) is the dominant transmission channel under OLS calibration, followed by business-investment growth (-0.062% at Q8) and housing-investment growth (-0.051% at Q4). Consump-

tion growth troughs at -0.036% at Q1 — the direct rate channel remains near-zero ($\beta_2 = +0.027$, $\beta_3 = -0.007$; no significant direct interest-rate sensitivity in AU consumption), so this is the general-equilibrium income effect. Employment growth declines -0.010% at Q4. The $\ln_IH$ level continues to decline through the 40-quarter horizon (long-tailed housing-investment response, consistent with AU's outsized construction cycle). Real GDP ($\ln Q$) troughs at -0.144% at Q11, and the output gap troughs at -0.086% at Q11.




Figure 7: Foreign demand shock

7.3.2 Foreign demand shock (eps_q_us) — 1pp US output gap *Figure 6.4: 1pp US output-gap shock (scale = $1.0 / \sigma_{q_us} = 1.0/1.1118 = 0.899$). Peak AU output gap $+0.11\%$ at Q4 (real GDP $+0.18\%$ at Q4), peak export growth $+0.46\%$ at Q2, peak business-investment growth $+0.10\%$ at Q4, peak housing-investment growth $+0.014\%$ at Q5 (foreign-demand channel identified through the proper trade ECM). The specification routes most of the foreign-demand impulse through direct trade ($\beta_x = 1.2$ LR pull) rather than the housing channel.*

7.3.3 Government spending shock (eps_g) — 1% of GDP *Figure 6.5: 1% of GDP government spending shock (scale = $1.0 / \sigma_g = 1.0/0.3 = 3.333$). The output gap peaks at $+0.91\%$ at Q10 and real GDP at $+1.56\%$ at Q11 — a peak GDP multiplier above unity, materially larger than under the prior calibration ($+0.18\%$ output gap), reflecting the 2026-05-31 employment- χ / housing ECM-speed writebacks that amplified the real-side blocks together with the $\lambda_{dom} = 0.40$ domestic-demand feedback; the exchange-rate appreciation still provides partial crowding-out. Imports rise $+0.30\%$ at Q6 as domestic demand pulls in foreign goods through the proper-ECM income-elasticity channel ($\beta_m = 1.5$).*




Figure 8: Government spending shock

(The larger multiplier is a consequence of the recent writebacks and may warrant a sensitivity check on λ_{dom} / the government-spending share.)

7.3.4 Commodity price shock (eps_pcom) — 10% increase *Figure 6.6: 10% commodity-price increase (scale = $10.0 / \sigma_{\text{pcom}} = 10.0/3.0 = 3.333$). AU-specific propagation: peak output gap +0.22% at Q1, housing-investment growth +0.054% at Q3, export growth +1.47% at Q1 (commodity-price loading in $b4_x$), import growth +0.49% at Q1 via the income channel. The commodity-export economy benefits in real terms; with the proper trade ECM, more of the impulse is now visible directly in $dl n_x$ rather than indirectly through the GDP identity.*

7.3.5 Cost-push shock (eps_pQ) — 1pp VA-price inflation *Figure 6.7: 1pp VA-price inflation cost-push shock (scale = $1.0 / \sigma_{\text{pQ}} = 1.0/0.6975 \approx 1.43$). Cost-push transmission operates through both the structural deflator channels in the E-SAT inflation equation (§4.4.0) and the level price-block dynamics. The VA-price PAC equation (eq_piQ_pac) absorbs ϵ_{pQ} into π_Q on impact (+1.35 qpp, with general-equilibrium amplification above the unit shock). Through the structural channels in eq_au_phillips, π_Q propagates into the CPI inflation gap immediately (+0.15 qpp impact); but under the flat AU Phillips and high policy smoothing the real-activity response is negligible — the output gap stays within $\pm 0.002\%$ over the entire horizon.*

The cost-push transmission mechanism in AU-PAC is the wp736 §3.1.1 / §5.2.6 structural channel: π_Q rises on impact, propagates into both consumer-price




Figure 9: Commodity price shock




Figure 10: Cost-push shock

inflation (via the deflator coefficients in eq. 4.4.0) and the policy-rate path (via the Taylor rule responding to the CPI gap), and the output gap contracts as the higher real interest rate dampens demand. Quantitatively, the AU response has two features that distinguish it from the FR-BDF benchmark (wp736 §5.2.6 reports -0.45% real GDP at Q8 in France):

1. **Very high policy smoothing:** AU's Taylor-rule persistence $\lambda_i = 0.96$, versus wp736's 0.89 — meaning AU policy-rate adjustments take roughly three times longer to build up than France's. Even a structurally correct CPI signal cannot translate into a sharp short-horizon policy tightening.
2. **Weaker AU $\pi_Q \rightarrow$ CPI passthrough:** the AU-estimated deflator coefficient $\alpha_{pc} = 0.17$ is much smaller than FR-BDF's 0.71 (wp736 Table 4.7.1), reflecting AU's lower share of domestic value-added in the consumption basket (large mining-led commodity-export sector, large import share of consumption goods).

Together these institutional features produce a real-GDP response to a +1pp cost-push shock that is essentially nil relative to FR-BDF: the output gap stays within $\pm 0.002\%$ (peak -0.001% on impact) and the GDP level $\ln Q$ remains within $\pm 0.002\%$ throughout the IRF horizon (-0.002% on impact, $+0.001\%$ by Q20). The qualitative cost-push transmission — CPI rising on impact, the policy rate tightening with a lag — matches FR-BDF; the near-zero real-activity magnitude reflects the AU institutional features documented above (flat Phillips, weak VA $\rightarrow$ CPI passthrough, very high policy smoothing), which leave the demand block almost untouched by a pure VA-price shock.




Figure 11: TFP shock

7.3.6 TFP shock (eps_tfp) — 1% level shift *Figure 6.8: TFP shock (eps_tfp_LR, plotted at scale $1.0/\sigma$ per gen_tfp_chart.m). The figure's purpose is the **supply-neutrality** property: output ($\ln_Q$) and potential output ($\ln_{QN}$) move together, so the output gap $yhat_au = \ln_Q - \ln_{QN}$ stays at zero by construction, and gap-variable panels (yhat_au, dln_n) plus unrelated nominal variables (piQ, i_10y) sit at floating-point noise. Under the current calibration the estimated long-run-TFP innovation is small ($\sigma = 0.01$) and, after the §6.8 structural-identity refactor ($\ln_Q$ defined by CES potential + demand aggregation rather than an E-SAT reduced form), its effect on the output **level** is negligible at 1 s.d.; employment edges down only $\approx -0.016\%$ over 40 quarters and wage inflation moves $\approx +0.001$ pp. (An earlier version of this figure showed a $+13\% \ln_Q / +10\% \ln_N$ response under a pre-refactor calibration with $\sigma_tfp = 0.2$; that magnitude is no longer representative. The weak TFP→output-level transmission under the current architecture is flagged for review — see model-review notes.)*




Figure 12: Term premium shock

7.3.7 Term premium shock (eps_tp) — 100bp annualized *Figure 6.9: 100bp annualized term-premium shock (scale = $0.25 / \sigma_tp = 0.25/0.05 = 5.0$; hybrid regime). Raises long rates by 0.25 pp without short-rate movement. As with the cost-push shock (Figure 6.7), the term-premium channel in AU-PAC works through level variables rather than short-run gaps — the demand block absorbs the financing-cost effect via permanent-income / target-level adjustments. Consumption level troughs at -0.28% , business investment level at -0.50% , housing investment level at -0.03% (all around Q12); consumption permanent income ($\ln_C_star$) tracks the consumption level closely. Under the current cal-*

ibration the financing-cost effect now also feeds the output gap ($yhat_au$ troughs at -0.079% at Q8) and VA-price inflation ($\pi_Q + 0.083$ pp on impact) through the amplified demand blocks — no longer floating-point noise, but still modest relative to the direct cash-rate shock. This is broadly the architectural pattern that makes the APP-style experiment (Figure 6.14) deliver moderate level effects rather than dramatic gap responses.




Figure 13: Output gap overview

7.3.8 Output gap overview — all shocks (policy-relevant sizes) *Figure 6.10: Output gap responses to all shocks at policy-relevant sizes. Shock sizes: monetary 100bp, foreign demand 1pp, govt spending 1pp, commodity 10%, TFP 1 s.d., term premium 50bp.*

7.4 AU vs FR-BDF IRF comparison

For the four PAC blocks where AU L2 estimates fit the wp1044 structural form (VA-price, employment, consumption, housing inv — §4.3–§4.7), it is useful to compare the AU-model IRF to the FR-BDF-reported equivalent under the same shock. wp1044 §5.2.6 (cost-push), §5.2.4 (monetary policy), and §5.2.7 (foreign demand) report Banque-de-France IRFs over a 40-quarter horizon; AU-PAC IRFs over the same horizon are in §7.3 above. The qualitative cross-block agreement is strong: the consumption response on impact is dominated by the real-lending-rate channel in both models; the BI response builds gradually via the user-cost channel; the housing-inv response is large and long-tailed; the wage-Phillips response builds via the CPI-passthrough channel.

Quantitative differences trace back to four sources, in descending order of magnitude:

1. **AU faster ECM speeds in 4 of 5 blocks** (§5.1). AU PAC blocks reach peak response $4\text{--}8\times$ faster than FR-BDF in the affected blocks, partially offsetting the differences from (2)–(4) below.

2. **Different Taylor-rule persistence.** wp1044's $\lambda_i = 0.89$ vs AU's 0.96 (cf. §7.3.5 cost-push discussion). AU policy rate adjustments take $3\times$ longer to build up than France's, dampening short-horizon transmission.
3. **Different VA-price-to-CPI passthrough.** AU's $\alpha_{pc} = 0.17$ vs wp1044's 0.71 (§4.4.0). AU CPI is less responsive to VA-price shocks because of the large import share of the consumption basket and the large commodity-export sector that decouples producer prices from consumer prices.
4. **Different exchange-rate regime.** AU has a fully floating AUD with strong commodity-price exposure; France has the EUR/EA monetary regime. The UIP and trade ECM mechanisms transmit foreign shocks differently in consequence.

For the BI block specifically (calibrated from wp1044 via Option 1, $\$3.6 + \$4.6.2 + \$5.3$), the IRF through the BI channel is structurally identical to wp1044's by construction; only the surrounding model differs. Mining-cycle dynamics enter via E-SAT + commodity shocks rather than through the BI block's own coefficients.

The full side-by-side AU-vs-wp1044 IRF panels are produced by `data/make_paper_charts.m` and saved to `dynare/paper_artifacts/irf_au_vs_wp1044_<block>.png` (pending the next MATLAB GUI session per `WORKING_PAPER_BLOCKERS.md`).

7.5 Conditional forecasting

The model supports conditional forecasting via residual inversion (ECB-Base pattern). Given a desired path for conditioned variables, the algorithm solves for the shock sequence that replicates it using the decision rule matrices (g_x , g_u).

RBA Tightening Scenario: +25bp/quarter for 4 quarters (100bp total), hold 4 quarters, normalize over 4 quarters:

Table 6.5: Conditional forecast — RBA tightening scenario

Variable	Q1	Q4	Q8	Q12	Q16
i_au (conditioned)	+0.063	+0.250	+0.250	+0.050	0.000
yhat_au	-0.000	-0.034	-0.097	-0.099	-0.041
pi_au	-0.000	-0.001	-0.004	-0.006	-0.003
dln_c	-0.000	-0.023	-0.059	-0.054	-0.017
dln_ib	-0.000	-0.043	-0.103	-0.085	-0.016
dln_ih	-0.000	-0.058	-0.155	-0.123	+0.001
dln_n	-0.000	-0.018	-0.073	-0.094	-0.049
i_10y	+0.026	+0.103	+0.103	+0.021	+0.001

Figure 6.11: Conditional forecast for an RBA tightening cycle. Housing invest-

ment is the most sensitive (-0.16% at Q8) and employment the most sluggish (peak at Q12); the 10Y rate moves approximately 40% of the policy rate (yield-curve flattening). The conditional-forecast driver was retired in the 2026-05-18 Phase S cleanup; the figure remains valid as a baked artefact and the scenario is reproducible under the Phase T architecture by a short residual-inversion script invoking `au_pac.mod`.

7.6 Forward guidance experiment

The model does not suffer from the forward guidance puzzle. Testing with superposition of N-quarter rate cuts (25bp each):

Table 6.6: Forward guidance scaling test

Normalised peak output response to N consecutive quarterly 25bp rate cuts, scaled so the N=1 response equals 1.0.

Duration N	Standard NK	Discounted NK	AU-PAC	Linear ref
1	1.00	1.00	1.00	1.00
2	1.44	1.45	1.99	2.00
4	1.72	1.73	3.93	4.00
6	1.78	1.79	5.75	6.00
8	1.79	1.80	7.40	8.00
10	1.79	1.80	8.89	10.00
12	1.79	1.80	10.14	12.00

AU-PAC’s near-linear scaling comes from three features: a high permanent-income discount ($\beta_c = 0.95$); a discounted term structure ($\kappa_{10} = 0.97$); and PAC adjustment-cost frictions that prevent explosive compounding. The standard NK and discounted NK models both saturate around $N = 5$ (ratio plateau at ≈ 1.79 – 1.80) — the textbook forward-guidance puzzle. AU-PAC tracks the linear reference to within 16% even at $N = 12$, confirming that PAC-style adjustment costs eliminate the puzzle for this model class. The mild sublinearity at long horizons (10.14 vs linear 12 at $N = 12$) reflects PAC’s intrinsic decay in marginal effect.

Figure 6.12: Forward guidance experiment — output gap responses for different durations.

Figure 6.13: Forward guidance puzzle test — AU-PAC (solid) vs standard NK (dashed). AU-PAC scales approximately linearly.

7.7 APP-style experiment — 200bp persistent term-premium compression

The RBA did not implement an asset-purchase programme directly analogous to the ECB’s APP or the Fed’s QE during 2015–2018, but a hypothetical scenario is

straightforward to simulate in AU-PAC: an unscheduled persistent compression of the term premium tp by 200 basis points annualised ($= -0.50$ qpp). The shock is propagated via eq. (eq_term_premium), $tp_t = \rho_{tp} \cdot tp_{t-1} + (1 - \rho_{tp}) \cdot tp_{ss} + \varepsilon_t^{tp}$, which feeds directly into the 10Y yield via eq. (48). Scale: $-0.50/\sigma_{tp} = -0.50/0.05 = -10.0$.

Figure 6.14: APP-style 200bp annualised term-premium compression. The 10Y yield falls 0.5 qpp on impact and gradually recovers over approximately 40 quarters. Consumption level rises by +0.21% peak (Q40, via the permanent-income channel through $\ln_C_star$), business investment by +0.16% peak (Q1, via WACC), and housing investment by +0.32% peak (Q40, via the mortgage-rate channel). The short-run output gap response is small: the term-premium shock works through the level / target channels rather than through the short-run PAC error terms — a structural feature of the PAC architecture in which persistent financial-condition shocks enter via $\ln_X_star$ (the target / permanent-income aggregate) rather than the cyclical $\ln_x_level$ deviation. An RBA APP of this magnitude would deliver roughly 20–30 basis points of stimulus to household consumption and approximately 30 basis points to housing investment over a 10-year horizon.

8. Conclusion

This paper has presented AU-PAC, a semi-structural macroeconomic model for Australia adapted from the FR-BDF framework. The model combines Polynomial Adjustment Costs with explicit expectations from an enriched 12-equation satellite VAR, enabling systematic analysis of how expectation formation affects monetary policy transmission.

Three Phase-L2 findings are the most publishable contributions of this version of the paper.

1. **The wp1044 PAC framework validates on Australian data for four of five behavioural blocks.** A wp1044-faithful partial-replication of all five PAC blocks (Eqs 16, 30, 35, 37, 46) was implemented in Phase L2 via the iterative-OLS pipeline of §3.5: block-specific VARs (6–9 state variables per block), the structural target constructions wp1044 specifies ($p_Q - p_Q$ for VA-price; $n_S - n_S$ for employment; c_* from permanent-income for consumption; I_H/I_H with price-spread for housing inv; I_B/I_B for business inv), the exact χ from the depth-m characteristic polynomial, growth-neutrality at the derived coefficient, and the full dummy set. Four of five blocks pass the wp1044 structural restriction (coef = +1 on PV) with R^2 in $[0.41, 0.81]$ and economically signed dynamics (§4.3–§4.5, §4.7). This is *consistent* with wp1044’s underlying assumption that Australia and France share a common agent-objective architecture for those four blocks.

2. **Consumption ECM speed is comparable across AU and France** (β_0 AU = 0.23 vs FR = 0.29 — within ~ 0.06). Of the blocks estimated in Phase L2, this is the closest cross-country agreement and the only block where AU’s ECM speed is not several times faster than France’s. The match is structurally meaningful: consumption is the block where wp1044’s permanent-income/PAC framework most directly applies (homogeneous-agent forward-looking consumers facing similar financial structure), and the AU L2 implementation builds the long-run target c^* from the same Eq 33 specification wp1044 uses on French data. The agreement validates the FR-BDF permanent-income channel for Australian consumption specifically (§4.5.2, §5.2). For comparison, AU’s ECM speeds in the other four blocks are 4–8 $\times$ wp1044’s French values (§5.1) — the consumption block is the exception that proves the rule.
3. **AU business investment structurally rejects the wp1044 PAC restriction; the paper adopts the Option 1 hybrid calibration.** Eleven specification variants of wp1044 Eq 46 were tested in Phase L2 P1c (Appendix G; full saga in PAC_BI_AU_EXPLORATION.md). Every strict-PAC variant delivers $R^2 \leq 0.11$ on raw `dln_ib`; the best non-strict variant (PV regressors free) gives $R^2 = 0.53$ but with the lead $PV(\Delta\hat{q})$ coefficient at **−5.03** — wrong sign and 5 $\times$ wrong magnitude vs the structural +1. This is not finite-sample noise: the coefficient mismatch is order-of-magnitude, persists across spec changes (target substitution, lag structure, trends, dummies), and three non-exclusive hypotheses (different agent objective; mining-sector dominance; sample-period contamination) are catalogued in §5.3. The production model adopts **Option 1**: import the BI deep parameters from wp1044 Table 3.5.13 ($\beta_0=0.096$, $\beta_1=0.33$, $\beta_2=0.11$, $\beta_3=0.69$, $\omega=0.35$, $\sigma_{CES}\approx 0.50$), preserve the wp1044 structural form (coef = +1 on PV, coef = $-\sigma$ on user-cost PV), and let AU-specific BI dynamics enter via the E-SAT VAR, the other 4 PAC blocks, AU-sized shocks, and AU-specific commodity/mining-cycle channels. The Bayesian estimation removes BI parameters from `estimated_params` (calibrated, not estimated). This is the standard small-open-economy modelling approach when local data identifies poorly (§3.6, §4.6.2, §5.3).

Additional findings carried over from the pre-L2 paper (numerically updated, methodologically unchanged):

4. **The Australian wage Phillips curve has standard New Keynesian structure.** Estimated on the ABS Wage Price Index, $\gamma_w = 0.354$ (90% HPD [0.268, 0.428]), $\lambda_w = 0.209$ (HPD [0.079, 0.332]), and $\kappa_w = -0.103$ (HPD [−0.178, −0.019]) — moderate contemporaneous CPI passthrough, moderate own-lag persistence, and a statistically significant positive structural Phillips slope (sign convention $\dots - \kappa_w \cdot pv_u_gap$). The Fair Work Commission’s institutional CPI indexation operates on roughly 15% of the workforce (award rates plus enterprise-agreement floors); most AU wage growth is set by private bargaining

responsive to the unemployment-gap PV. AU wages therefore adjust to labour-market slack via a Phillips-curve channel and to inflation via a moderate contemporaneous CPI coefficient plus the inflation-anchor weight ($1 - \lambda_w - \gamma_w = 0.44$).

5. **COVID dummies are essential for PAC estimation.** Without explicit treatment of the 2020Q2–Q3 outliers, two of five PAC equations (consumption and employment) produce wrong-signed AR(1) coefficients. In the Phase L2 replication every block carries two-or-more COVID dummies (e.g. $d_{20Q2} = -13.04$, $t = -13.07$ in consumption); the inclusion is mechanical (COVID was unique) and the dummy magnitudes are stable across specifications.
6. **Monetary policy transmits primarily through the exchange rate and business investment, not consumption.** Under the equation-by-equation OLS calibration (wp1044 methodology), a 100 bp annualised cash-rate tightening produces a real-GDP trough of -0.144% at Q11, an output-gap trough of -0.086% at Q11, a CPI-inflation fall of -0.004 pp at Q3, an AUD appreciation of 0.99% at Q8 via the forward-NPV UIP, and a 10-year yield rise of 0.105 pp on impact via the term-structure NPV. The direct consumption rate channel is near-zero (the L2 SA OLS gives HtM sensitivity $\beta_2 = +0.027$ and impact-rate sensitivity $\beta_3 = -0.007$, both insignificantly different from zero). Monetary transmission therefore operates through the exchange-rate channel (-0.99%), business investment (-0.062%), and housing investment (-0.051%), with AU's variable-rate mortgage market and strong CPI-indexation wage Phillips shaping the relative magnitudes of these channels.
7. **Higher CES substitution elasticity dampens, not amplifies, the IRF.** The FR-BDF 2026 recalibration raised σ from 0.34 to 0.54 — well below Cobb-Douglas but more substitutable than the previous AU calibration. The peak responses fell modestly as a result, because firms with higher input substitutability require smaller factor-price movements to clear a demand shortfall. The structural intuition is well-known but its direction is often misread: higher CES elasticity dampens output responses when wage-price-spiral feedback (κ_w) is small, as it is in AU.
8. **Australia-specific channels matter.** Variable-rate mortgages ($\rho_{lh} = 0.97$) create the strongest demand channel, with banks adjusting lending rates slowly. The commodity-price channel feeds export volumes, the export deflator, and the import deflator. The endogenous Taylor rule creates a feedback loop (demand $\rightarrow$ output gap $\rightarrow$ RBA $\rightarrow$ rates $\rightarrow$ demand) that closes the IS curve.

Comparison with the RBA model suite. §7.2.4 benchmarks AU-PAC against the four Australian models surveyed in Mulqueeney, Ballantyne and Hambur (2025). AU-PAC's output-gap trough (-0.086% at Q11) and real-GDP trough (-0.144% at Q11) are at the low end of the RBA suite in

demand-response terms. The conservative response reflects three features: PAC adjustment-cost frictions that distinguish FRB/US-style models from VAR / DSGE alternatives; the FR-BDF 2026 CES calibration with higher capital–labour substitutability that dampens factor-price movements; and the equation-by-equation OLS calibration that reveals near-zero consumption rate sensitivity in AU data (removing the typically dominant demand-side transmission channel). AU-PAC therefore complements the RBA suite on the conservative end of “diversity supports more robust policymaking” (Mulqueeney et al. 2025, p. 11).

Open extensions for future work, in priority order:

1. **Full wp1044 fidelity in the four fitting blocks**: add the wp1044 Phillips Eq 18 + Okun Eq 19 auxiliary equations as separate AR(1) states (currently absorbed via the $\tilde{\pi}_{w_eff}$ and $\hat{n}^*_S$ single states); replace the AU L2 VAR(1) Phi with a Bayesian Minnesota-prior VAR(p); lift the four-block R^2 values toward wp1044’s 0.61–0.95 range. Estimated ~1 week of additional work.
2. **Test Hypotheses A/B/C from §5.3 on AU BI**: split `dln_ib` into mining vs non-mining components using ABS Cat. 5625, re-estimate wp1044 Eq 46 on each component separately, test whether either passes the $\text{coef} = +1$ restriction. Lengthen the pre-1993 sample by splicing ANA data. Run wp1044 on a longer FR sample to check Hypothesis C mechanically. Estimated ~2 weeks.
3. **Adopt FR-BDF 2026 financial-block extensions**: the Dees et al. (2022) NFC financial accelerator and the Bové et al. (2020) household debt-service-ratio block, both of which are natural amplification mechanisms for the cost-of-capital and mortgage channels in AU-PAC.
4. **High-frequency RBA OIS-surprise IV estimation** for the consumption rate-change coefficient `b_di_c` (Bishop–Tulip 2017; Beechey–Wright 2009). The current value is Bayesian-regularised because AU IV without external monetary-surprise data gives wrong-signed estimates.
5. **Channel-decomposition exercise** of the type Mulqueeney et al. (2025) report for MARTIN and DINGO, in which each transmission channel (exchange rate, asset prices/wealth, savings/investment, cash flow) is “turned off” in turn to isolate its contribution to the overall response.

9. Australia-Specific Features

Variable-rate mortgages

Australia’s predominantly variable-rate mortgage market creates the strongest housing channel among all demand components. The mortgage lending rate (i^{LH}) adjusts to the 10Y rate with persistence $\rho_{LH} = 0.97$ (estimated from RBA F5 data, $T = 127$) and a spread of 0.40% quarterly. The 0.97 persistence implies a half-life of approximately 23 quarters, reflecting the stickiness of standard

variable mortgage rates in Australia. The housing-investment target is directly sensitive to the mortgage rate gap ($\kappa_{mort} = 0.048$), and the auxiliary equation has a large interest-rate loading ($a_{ih,i} = -0.15$).

Commodity-price channel

Australia is a net commodity exporter, so commodity prices enter primarily as an *income* shock via the export deflator and terms of trade. Commodity prices follow an exogenous AR(1) process:

$$\Delta \ln P_t^{com} = \rho_{pcom} \Delta \ln P_{t-1}^{com} + \varepsilon_t^{pcom}$$

with $\rho_{pcom} = 0.42$ estimated from IMF / RBA G3 commodity-price data — reflecting the high quarterly volatility of Australia’s broad commodity basket. Commodity prices feed into export volumes ($b_4^x = 0.15$), the export deflator ($\alpha_{pcom} = 0.10$), and the import deflator ($\beta_{pm,com} = 0.42$, estimated from ABS data).

Endogenous monetary policy

AU-PAC’s short rate is set endogenously by an RBA Taylor rule responding to AU inflation and the output gap. This closes a policy feedback loop — demand shock $\rightarrow$ output gap $\rightarrow$ Taylor rule $\rightarrow$ interest rate $\rightarrow$ demand components $\rightarrow$ output gap — producing endogenous stabilisation. The estimated rule has inertia $\lambda_i = 0.828$, inflation weight $\alpha_i = 0.279$, and output weight $\beta_i = 0.135$.

US as foreign bloc

The US enters as the foreign bloc in E-SAT. The AU–US demand spillover ($\delta = 0.20$) captures the trade and financial channel linking US conditions to Australia, while the US inflation process influences Australian import prices and exchange-rate dynamics.

Architectural deviations from FR-BDF

AU-PAC diverges from FR-BDF wp736 / wp1044 along three architectural dimensions that are not market-structure-driven (the four subsections above) but model-design simplifications adopted for AU estimation tractability. They are documented at their primary locations in §3.1 (IS curve), §4.4.0 (Phillips curve), and §4.9 (demand deflators); this subsection collects them for a reader specifically interested in the AU-vs-FR-BDF map.

AU innovation	Where	FR-BDF analogue	Rationale
Domestic-demand feedback term $\lambda_{dom}\hat{y}_t^{dom}$ in the IS curve	§3.1, eq (1)	No analogue — wp736 IS curve has only $\hat{y}^{US}$ + interest gap	Closes the Keynesian multiplier loop; data strongly supports a $4\times$ larger feedback than the prior centre.
Two Phase V additions in the Phillips curve: lagged VA-price passthrough $\alpha_{pc,lag}(\pi_{t-1}^Q - \bar{\pi}_{t-1}^{au})$ and an ECM correction $b_{ECM,pc}(p_{t-1}^{C,*} - p_{t-1}^C)$	§4.4.0	$\alpha_{pc,lag}$: no analogue. $b_{ECM,pc}$: replaces the long-run-target ECM that lives inside <code>eq_pi_c</code> in wp736 eq 80.	$\alpha_{pc,lag}$: captures empirical AU lagged import-component repricing. $b_{ECM,pc}$: pragmatic device to restore CPI long-run anchoring after collapsing the deflator ECM structure (next row).
Demand-deflator reduced-form collapse: 6 single-equation AR(1)+cost-push deflators (<code>eq_pi_c</code> , <code>eq_pi_ib</code> , <code>eq_pi_ih</code> , <code>eq_pi_x</code> , <code>eq_pi_m</code> , <code>eq_pi_g</code>) instead of 12 paired target+ECM equations	§4.9	wp736 eqs 79-88 / wp1044 eqs 51-55, 112-114: each deflator carries an explicit IAD-weighted long-run target $p^{j,*}$ and a short-run equation with an ECM term	Six additional latent target processes would substantially enlarge the estimation surface with limited identification on 128 quarters of AU data. Long-run anchoring is recovered via the Phillips-curve ECM term ($b_{ECM,pc}$ above).

These deviations are intentional and were preserved through Phase Q–X estimation cycles. They are flagged explicitly here (rather than buried in the equation

specifications) so that a researcher porting AU-PAC to another country or extending it back toward the full FR-BDF target/ECM structure can identify the exact code-level entry points.

Appendix A: Complete Variable List

A.1 E-SAT Core (11 variables)

Variable	Description	SS value
yhat_au	AU output gap (%)	0
i_au	RBA cash rate (quarterly %)	1.049
pi_au	AU GDP deflator inflation (quarterly %)	0.625
yhat_us	US output gap (%)	0
pi_us	US GDP deflator inflation (quarterly %)	0.500
ibar	Long-run interest rate anchor	1.049
pibar_au	AU inflation anchor	0.625
pibar_us	US inflation anchor	0.500
i_gap	Interest rate gap ($i_{\text{au}} - \text{ibar}$)	0
pi_au_gap	AU inflation gap	0
pi_us_gap	US inflation gap	0

A.2 VA Price Block (2 variables)

Variable	Description	SS value
piQ	VA price inflation (quarterly %)	0.625
pQ_level	VA price log-level accumulator	0

(Pre-Phase-Y also carried `piQ_star`, `piQ_star_bar`, `pQ_gap`, `pQ_star_level` as diagnostic-only orphan variables. These were removed 2026-05-17 — see §4.3.1.)

A.3 Supply Block (5 variables)

Variable	Description	SS value
dln_k	Capital services growth	0
dln_y_star	Potential output growth	0
dln_tfp	TFP growth	0
dln_ulc	Unit labor cost growth	0.625
dln_prod	Labor productivity growth	0

A.4 Labor Market (10 variables)

Variable	Description	SS value
u_gap	Unemployment gap	0
pv_u_gap	PV of future unemployment gaps	0
pi_w	Wage inflation (quarterly %)	0.625
dln_n	Employment growth	0
dln_n_star	Employment target growth	0
dln_n_star_bar	Employment trend growth	0
n_gap	Employment gap	0
dln_n_1, dln_n_2, dln_n_3	Employment lag auxiliaries	0

A.5 Demand Block (14 variables)

Variable	Description	SS value
dln_c, dln_c_star, dln_c_star_bar	Consumption growth, target, trend	0
c_gap	Consumption gap	0
pv_yh	Permanent income PV	0
dln_ib, dln_ib_star, dln_ib_star_bar	Business investment growth, target, trend	0
ib_gap	Business investment gap	0
dln_ib_1	Business investment lag auxiliary	0
dln_ih, dln_ih_star, dln_ih_star_bar	Housing investment growth, target, trend	0
ih_gap	Housing investment gap	0
dln_ih_1	Housing investment lag auxiliary	0

A.6 Financial Block (15 variables)

Variable	Description	SS value
i_10y	10Y bond yield	1.349
tp	Term premium	0.300
pv_i	PV of expected future short rates	1.049
wacc	Weighted average cost of capital	1.834
i_COE, i_LB_firms, i_BBB	Component rates	2.149, 1.599, 1.399
s_COE, s_LB_firms, s_BBB	Component spreads	0.800, 0.250, 0.050
s_gap	Exchange rate gap	0
uc_k, dln_uc_k	User cost of capital and growth	1.859, 0

Variable	Description	SS value
i_lh	Household bank lending rate	1.749
dln_ph, ph_gap	Housing prices growth and gap	0, 0

A.7 Trade, Deflators, Government (20 variables)

Variable	Description	SS value
dln_x	Export growth	0
ln_x_level	Log exports level (deviation, accumulates dln_x)	0
ln_x_eq	Export LR equilibrium ($\beta_x \cdot \hat{y}_{us} + \gamma_x \cdot s_{gap}$)	0
x_gap	Export EC term ($\ln_x_{eq} - \ln_x_{level}$)	0
dln_m	Import growth	0
ln_m_level	Log imports level (deviation, accumulates dln_m)	0
ln_m_eq	Import LR equilibrium ($\beta_m \cdot \ln_{d_{iad}} + \gamma_m \cdot s_{gap}$)	0
m_gap	Import EC term ($\ln_m_{eq} - \ln_m_{level}$)	0
ln_d_iad	Log import-weighted demand level (accumulates iad)	0
dln_pcom	Commodity price growth	0
pi_c, pi_ib, pi_ih, pi_x, pi_m, pi_g	6 demand deflators	all 0.625
dln_g	Government spending growth	0
yhat_dom	Domestic demand gap	0
iad	Import-adjusted demand	0
rw_gap	Real wage growth gap	0

A.8 var_model Shadow + PAC Levels (27 variables)

12 var_model shadow variables (y_gap_var, i_gap_var, etc.) plus 7 auxiliary targets (piQ_hat, n_hat, c_hat, etc.) plus 6 pv_aux correction terms plus log-level accumulators (ln_c_level, ln_ib_level, ln_ih_level, ln_n_level, pQ_level).

Appendix B: Complete Shock List

B.1 E-SAT core shocks (8)

Shock	Std dev	Description
eps_q	0.80	AU demand shock
eps_i	0.10	Monetary policy shock
eps_pi	0.60	AU cost-push shock
eps_q_us	1.10	US demand shock
eps_pi_us	0.40	US cost-push shock
eps_ibar	0.01	Long-run rate anchor
eps_pibar_au	0.01	AU inflation anchor
eps_pibar_us	0.01	US inflation anchor

B.2 Structural shocks (12)

Shock	Std dev	Description
eps_pQ	0.50	VA price cost-push
eps_w	0.60	Wage push
eps_n	0.50	Employment
eps_c	0.50	Consumption
eps_ib	1.50	Business investment
eps_ih	2.00	Housing investment
eps_tfp	0.10	TFP
eps_x	1.00	Export
eps_m	0.80	Import
eps_g	0.30	Government spending
eps_pcom	1.00	Commodity price
eps_lh	0.05	Mortgage rate

B.3 Financial shocks (5)

Shock	Std dev	Description
eps_10y	0.10	Long rate residual
eps_tp	0.05	Term premium
eps_COE	0.10	Equity spread
eps_LB_firms	0.05	Bank lending spread
eps_BBB	0.02	BBB bond spread

B.4 Other shocks (8 deflator/trade + 12 var_model + 2 COVID dummies)

Exchange rate (eps_s, 0.50), deflators (eps_pc, eps_pib, eps_pih, eps_px, eps_pm, eps_pg), housing price (eps_ph), plus 12 var_model shadow shocks and 2 COVID pulse dummies.

B.5 Rounds 4–8 extension shocks (9, added 2026-05-20)

Shock	σ	Channel
eps_pop_bar	0.05	Round 5: demographic-trend gap ($\rightarrow$ eq_dln_n_star_bar shifter)
eps_ibar_us	0.01	Round 4: US LR rate anchor (mirror lambda_ibar)
eps_i_us	0.15	Round 4: US policy-rate Taylor residual
eps_tau_GST	0.10	Round 6: GST effective-rate gap ($\rightarrow$ eq_pi_c)
eps_tau_PAYG	0.20	Round 6: PAYG effective-rate gap ($\rightarrow$ eq_dln_c_star_bar Δ form)
eps_tau_CIT	0.30	Round 6: CIT effective-rate gap ($\rightarrow$ eq_uc_k)
eps_BLR	0.05	Round 8: Bank-Lending-Rate nowcast injection
eps_MAPI	0.50	Round 8: Mortgage Asset-Price Indicator nowcast
eps_MAPU	0.30	Round 8: Mortgage Asset-Price Underwriting nowcast

Appendix C: Growth Neutrality Proofs

C.1 PAC equations

At the balanced growth path, the sum of the AR lag coefficients plus the expectations weight (ω) plus the error-correction coefficient equals 1:

Equation	b0 (EC)	Sum AR	omega (h-vector)	GN residual	Sum
VA price	0.060	0.500	0.452	-0.012	1.00
Consumption	0.060	0.149	0.678	0.113	1.00
Business inv	0.030	0.281	0.501	0.188	1.00
Household inv	0.049	0.290	0.569	0.092	1.00
Employment	0.040	0.470	0.446	0.044	1.00

C.2 Demand deflators

All deflator coefficients on inflation-type terms sum to 1:

Deflator	rho	alpha	beta_m	gamma/other	Anchor	Sum
Consumption	0.40	0.30	0.10	0.03	0.17	1.00
Business inv	0.35	0.25	0.12	—	0.28	1.00
Housing inv	0.45	0.25	0.08	—	0.22	1.00
Export	0.30	0.20	-0.05	0.10	0.45	1.00
Import	0.30	0.15	0.08	0.05	0.42	1.00
Government	0.50	0.30	—	—	0.20	1.00

C.3 Wage Phillips curve

$\lambda_w + \gamma_w + (1 - \lambda_w - \gamma_w) = 0.329 + 0.138 + 0.533 = 1.00$. On the balanced growth path with productivity growth g : $\pi_{SS}^w = \pi_{SS} + g$. The inflation-anchor coefficient (0.533) carries genuine weight under the WPI calibration — wage inflation in the long run is anchored to expected steady-state inflation through this channel, not solely through the contemporaneous CPI term ($\gamma_w = 0.138$) and own-lag ($\lambda_w = 0.329$).

Appendix D: h-Vector Decomposition

D.1 Summary

PAC equation	h-param indices	h-vector sum	GN coeff	Amplification
VA price	[173, 174, 175, 176]	0.452	0.048	1.0x
Consumption	[153, 154, 155, 156]	0.678	0.173	1.84x
Business inv	[158, 159, 160, 161]	0.501	0.218	1.43x
Household inv	[163, 164, 165, 166]	0.569	0.141	1.90x
Employment	[168, 169, 170, 171]	0.446	0.084	1.49x

Detailed per-state h-vector element weights require extraction from `M_.params` at the h-param indices. The enriched `var_model` (12x12) produces h-vectors where only the target variable element is non-zero in the current Dynare 7.0 implementation (the companion matrix concentrates weight on the target state). Element-level decomposition is deferred to future work.

Appendix E: Complete Parameter List

See the comprehensive parameter tables throughout Section 4. A machine-readable list of all 262 parameters with their values is available alongside the model code at <https://github.com/DavidAS Stephan/AUSPAC>.

The four trade long-run elasticities introduced in the Section 4.8 ECM rewrite are summarised here for convenience:

Parameter	Value	Prior (for MCMC re-run)	Role
beta_m	1.50	Normal(1.50, 0.30)	LR income elasticity of imports (FR-BDF eq 76)
gamma_m	-0.40	Normal(-0.40, 0.20)	LR RER elasticity of imports

Parameter	Value	Prior (for MCMC re-run)	Role
beta_x	1.20	Normal(1.20, 0.30)	LR foreign- income elastic- ity of exports (FR- BDF eq 71)
gamma_x	+0.40	Normal(+0.40, 0.20)	LR RER elastic- ity of exports

Appendix F: Identification analysis

Driver: `identification_analysis.m`. Two approaches are reported below.

F.1 Dynare identification command — not run

Dynare 7.0's `identification` command requires analytic derivation of the model, which is **incompatible with `diffuse_filter`** — the filter setting used during estimation to handle the model's unit-root level accumulators (`ln_Q`, `ln_K`, `ln_C`, etc.). Running the command produces:

```
ERROR: analytic derivation is incompatible with diffuse
filter
```

For Iskrev (2010) and Komunjer–Ng (2011) tests at the posterior mode we would need either (a) a stationarised version of the model (replacing level accumulators with their deviations from steady state), or (b) a future Dynare release that supports analytic derivation under diffuse filtering. Neither is in scope for this paper; we report the posterior-based identification proxy below instead.

F.2 Posterior-width-based weak-identification flagging

For each of the 28 estimated parameters, we report the 90% HPD interval from the MCMC posterior (20,000 draws $\times$ 2 chains). Absolute HPD width is the headline measure of how much the data has constrained the parameter; the normalised width ($\text{HPD_width} / |\text{posterior_mean}|$) is informative only when the posterior is well away from zero.

Table F.1: Posterior HPD widths

Sorted by absolute width, descending. Parameters at “near-zero” identification are flagged where the posterior mean is essentially zero.

Parameter	HPD low	HPD high	width	normalised width
b2_c	-0.6029	-0.0504	0.553	1.67
b3_ih	0.0623	0.3971	0.335	1.48
b1_n	0.1527	0.4848	0.332	1.03
b1_pQ	0.1202	0.4512	0.331	1.14
lambda_w	0.1342	0.4471	0.313	1.08
b3_ib	0.1629	0.4642	0.301	0.97
b3_c	-0.0680	0.0965	0.164	8.27 (<i>mean near zero</i>)
b2_pQ	-0.0796	0.0825	0.162	<i>huge (mean $\approx$ zero)</i>
b1_ih	0.0301	0.1888	0.159	1.38
kappa_w	0.0102	0.1694	0.159	1.65
b5_n	-0.0707	0.0885	0.159	22.06 (<i>mean near zero</i>)
b1_ib	0.0186	0.1381	0.120	1.49
gamma_w	0.0796	0.1960	0.116	0.84 (<i>tightly identified at 0.138</i>)
b0_n	0.0108	0.0976	0.087	1.52
b0_c	0.0231	0.0909	0.068	1.13
b1_c	0.0046	0.0629	0.058	1.65
b0_pQ	0.0057	0.0517	0.046	1.50
b0_ih	0.0065	0.0498	0.043	1.50
b0_ib	0.0050	0.0317	0.027	1.42

Findings:

1. **gamma_w is tightly identified** (90% HPD [0.080, 0.196] — width 0.12 around a mean of 0.138, normalised width 0.84). The data sharply distinguishes the moderate contemporaneous CPI passthrough from larger values, and the relative weights of the own-lag, CPI, and unemployment-gap channels in the wage Phillips curve are robustly identified.
2. **Three “near-zero” parameters** show large *normalised* widths (b5_n, b2_pQ, b3_c) but small *absolute* widths (≤ 0.16). The data has sharply identified them at zero — the AU output gap does not have a direct cyclical accelerator effect on either employment growth or VA-price inflation, and the output-gap channel into consumption is essentially absorbed by the rate channel b2_c. This is a sharp identification result rather than weak identification.
3. **b2_c (consumption interest-rate channel)** is the structural parameter with the widest absolute HPD: [-0.61, -0.05]. The data identify the sign

(negative — rate hikes depress consumption) and statistical significance, but the magnitude is uncertain over almost an order of magnitude. External RBA OIS-surprise data (Bishop–Tulip 2017; Beechey–Wright 2009) would sharpen this.

4. **The five PAC EC speeds** ($b0_pQ$, $b0_c$, $b0_ib$, $b0_ih$, $b0_n$) are tightly identified with normalised HPDs around 1.0–1.5 — the Beta priors are informative and the data provide material updates. Posteriors are all in the 0.02–0.07 quarterly range.
5. **Shock standard deviations** (Table 5.6): eps_ih and eps_n have the widest HPDs proportionally ($[0.55, 2.23]$ and $[0.12, 0.79]$ respectively). The wide HPD for eps_n is a by-product of $b5_n$ being identified at zero — most of the employment-equation residual variance is absorbed by eps_n rather than apportioned across structural channels.

Figure F.1: Posterior densities (blue) from the MCMC run, with posterior mean (red dashed) and posterior mode (green solid) markers. Most posteriors are unimodal and roughly symmetric; the shock standard deviations eps_ih and eps_n are the only parameters showing notable asymmetry consistent with weak identification.

F.3 What a future Dynare run would add

Given a stationarised version of the model without `diffuse_filter`, the following diagnostics would be available:

- **Iskrev (2010) rank test** at the posterior mode (yes/no answer on local identification);
- **Komunjer–Ng (2011) test** for structural identification of the linear DSGE under partial-information assumptions;
- **Identification strength** (moment-condition based) for each parameter — a quantitative ranking of how informative each moment of the data is for each parameter;
- **Pairwise collinearity matrix** — which parameter pairs are near-perfectly correlated under the model’s mapping from parameters to moments, flagging redundancies that the posterior HPD-width approach cannot directly detect.

These remain pending future work; the HPD-width analysis here is sufficient to flag which parameter magnitudes are well determined and which would benefit from additional data.

Appendix G: Business Investment Exploration (the wp1044 PAC rejection on AU data)

This appendix tabulates the 11 specification variants of wp1044 Eq 46 tested in Phase L2 P1c, summarises the diagnostic for each, and lists the corresponding `data/pac_blocks/` artefact. The full prose narrative is in `PAC_BI_AU_EXPLORATION.md` (10 sections); this appendix is a deliberately compact summary intended to be self-contained for the paper.

G.1 The structural restriction in question

wp1044 §3.5.3 Eq 46 specifies the business-investment PAC equation. The FOC of the agent’s quadratic-cost minimisation gives four PV terms with structural coefficients:

- $PV(\Delta \hat{q})_{t|t-1}$ at coef **+1** (market VA gap growth)
- $PV(\Delta \bar{q})_{t|t-1}$ at coef **+1** (market VA trend growth)
- $PV(\Delta \log \hat{r}_{KB})_{t|t-1}$ at coef $-\sigma$ (user cost gap growth)
- $PV(\Delta \log \bar{r}_{KB,t-1})_{t|t-1}$ at coef $-\sigma$ (user cost trend growth)

where $\sigma = \sigma_CES$. These coefficients are structural identities of the wp1044 model, not free parameters. wp1044 reports them as exactly +1 and $-\sigma$ in the fitted equation (Table 3.5.13), with $R^2 = 0.83$ on French data.

G.2 Variant table

All variants estimated on the same AU 122-obs sample (1993Q2–2023Q3), `dln_ib = 100·Δlog(au_gfcf_nondwelling)` as LHS, $\sigma_CES = 0.5366$ (AU labour-FOC). wp1044 BI ingredients: ECM on `log(I*_B/I_B)`, depth-2 lags, the 4 PV terms above, derived growth-neutrality $(1 - \beta_1 - \beta_2 - \omega)(\Delta \bar{q}_{t-1} + \Delta \log \bar{r}_{KB,t-1})$, β_3 on synthetic df gap (`= c + h_inv + exports`), $\omega=0.35$ calibrated, 3 COVID dummies (`d_20Q1, d_20Q2, d_20Q3`).

Table G.1: Eleven specification variants of wp1044 Eq 46 on AU data

Variant	Spec	R ² (raw)	Verdict
baseline (<code>estimate_pac_business_invest.m</code>)	strict wp1044 (<code>PACs_inv.m</code>)	0.09	Coefficients hit clamps; barely any structural signal
v1 (<code>estimate_pac_business_invest.m</code>)	strict PAC + 10 AU dummies (<code>AU_dummies.m</code>) (GST, GFC, mining-boom phases, COVID)	0.11	Dummies don’t penetrate the over-fit PAC structure

Variant	Spec	R ² (raw)	Verdict
v2 (estimate_pac_business_inv_v2.m)	pre-residualise in business_inv_v2.m PAC on residuals	-23.7	Two-stage approach inconsistent; PV depends on un-cleaned VAR Best fit; PV($\Delta\hat{q}$) coef = -5.03 → PAC structurally rejected Single-shot doesn't fix the wedge Diverges; β_1 , β_2 hit clamps
v3-A (estimate_pac_business_inv_v3.m, Variant A)	PV regressors with FREE coefficients + dummies	0.53	
v3-B	strict PAC + dummies, single shot	-2.20	
v3-C	strict PAC + dummies, iterative, looser clamps	-67.0	
v4 (estimate_pac_business_inv_v4.m)	combined PV at in business_inv_v4.m + dummies	-33.0	Weaker structural restriction still fails R ² _adj = 0.73 but raw R ² = -10.7; wedge persists Even with commodity-augmented target, strict PAC fails
v5 (estimate_pac_business_inv_v5.m)	strict PAC + in business_inv_v5.m + 3-segment piecewise trends	-10.7	
v6 (estimate_pac_business_inv_v6_tot.m)	Option 2: replacing q_au_v6_tot.m (market VA) with q_AU (ToT-augmented target) , strict PAC + dummies	-39.1	
wp736 (estimate_pac_business_inv_wp736.m)	wp736 Eq 64 in business_inv_wp736.m 2-PV form) + dummies	-0.75	
			The pre-COVID original spec also fails on AU

Variant	Spec	R ² (raw)	Verdict
simplified (estimate_pac_business_investment_simple.m)	drops PV lags + ECM + Δ df gap + dummies	0.33	$\beta_1=0.23 \approx$ wp1044's 0.33; $\beta_3=0.82 \approx$ wp1044's 0.69. No PAC structure but useful as a reduced-form benchmark for the post-Phase-L2 question of what AU data wants.

G.3 The killer diagnostic from variant v3-A

When the four PV regressors are estimated as free OLS regressors on AU data (v3-A, the highest-R² variant), the structural restrictions are flatly rejected:

PV term	wp1044 imposed	v3-A AU free-OLS	Diagnostic
PV($\Delta\hat{q}$)	+1.00	−5.03	WRONG SIGN, 5× wrong magnitude
PV($\Delta\bar{q}$)	+1.00	+2.77	Right sign, 2.8× wrong magnitude
PV($\Delta\log \hat{r}_{KB}$)	$-\sigma = -0.54$	−1.25	Right sign, 2.3× wrong magnitude
PV($\Delta\log \bar{r}_{KB}$)	$-\sigma = -0.54$	+0.001	Wrong sign, magnitude near zero

The PV($\Delta\hat{q}$) lead-coefficient mismatch (−5.03 vs +1.00) is the structural killer: AU business-investment agents respond to forward-looking market-VA-growth expectations in the *opposite* direction from wp1044's FOC prediction, and with 5× the magnitude. No subset of the FOC restrictions is satisfied by the data.

G.4 The variant v6 (Option 2) diagnostic

The most plausible alternative to “AU agents are different” (Hypothesis A of §5.3) is that AU business-investment agents respond to a *different target* than wp1044's market VA — specifically, the commodity-augmented terms-of-trade

target q_{AU} . Variant v6 tested this by replacing q with $q_{AU} = q + 0.18 \cdot \log(ToT)$ in the LR target and re-running the PV decomposition with free coefficients:

PV term	Option-2 v6 AU free-OLS	wp1044 imposed
PV(Δq_{AU_hat})	−0.11	(would be +1 if Option 2 held)
PV(Δq_{AU_bar})	+0.30	(would be +1)
PV($\Delta \log \hat{r}_{KB}$)	−0.27	$-\sigma = -0.54$
PV($\Delta \log \bar{r}_{KB}$)	+0.025	$-\sigma = -0.54$

The PV(Δq_{AU_hat}) ≈ 0 result tells us: **changing the target variable does not rescue PAC**. AU business investment does not respond to forward-looking expectations of EITHER market VA OR terms-of-trade with the structural FOC coefficient. The agent’s forward-looking weight is fundamentally different from wp1044’s framework in this block. Option 2 is empirically rejected and is not pursued further.

G.5 The Option 1 hybrid calibration adopted

Path-forward decision (§3.6 + §4.6.2 + §5.3): import the BI deep parameters from wp1044 Table 3.5.13.

```
% dynare/au_pac.mod lines 1746--1771 (Phase L2 P1c, 2026-05-26):
b0_ib = 0.096; // wp1044 Table 3.5.13 (overrides Phase V writeback 0.0181)
b1_ib = 0.33; // wp1044 Table 3.5.13 (overrides 0.0809)
b2_ib = 0.11; // wp1044 Table 3.5.13 (overrides 0; depth-2 PAC)
b3_ib = 0.69; // wp1044 Table 3.5.13 (overrides 0.3120; coef on Δdf gap)
omega_ib = 0.35; // already at wp1044 value
// sigma_ces already at 0.5366 (AU labour-FOC, matches wp1044's 0.50)
```

Bayesian estimation `au_pac_bayesian.mod` removes BI parameters from `estimated_params` (treats as calibrated). The wp1044 PAC structural form (coef = +1 on PV, coef = $-\sigma$ on user-cost PV) is preserved verbatim in the simulation; only the deep parameters are imported. AU-specific BI dynamics enter via:

- the AU-driven E-SAT VAR state propagating into BI’s PAC expectation;
- the AU-estimated other 4 PAC blocks feeding BI through the GDP identity / demand block;
- AU-sized shocks (`stderr eps_ib = 2.7211` AU-estimated);
- AU-specific commodity / mining-cycle shocks flowing through trade and the deflator block.

This is the standard small-open-economy modelling approach when local data identifies poorly; the deep parameters are borrowed from a larger / better-identified economy and the country-specific dynamics enter through the rest of the model. See §5.3 for the three explanatory hypotheses (different agent objective; mining-sector dominance; sample-period contamination), §4.6.2 for the

per-block paper presentation, and `PAC_BI_AU_EXPLORATION.md` for the complete saga.

G.6 Open research extensions

None of the 11 variants explored Hypothesis A or B of §5.3 directly. A future Phase L3 could (i) split AU BI into mining vs non-mining components via ABS Cat. 5625 and test wp1044 Eq 46 on each component separately; (ii) extend the equation with a Bernanke-Gertler-Gilchrist financial-accelerator shock; (iii) lengthen the AU sample by splicing pre-1993 ANA data; (iv) test wp1044 on a longer FR sample to verify Hypothesis C’s mechanical sample-length effect.

References

Models and methodology

Brayton, F., Laubach, T., and Reifschneider, D. (2014). “The FRB/US Model: A Tool for Macroeconomic Policy Analysis.” *FEDS Notes*, Board of Governors of the Federal Reserve System.

Lemoine, M., Turunen, H., Chahad, M., Lepetit, A., Zhutova, A., Aldama, P., Clerc, P., and Laffargue, J.-P. (2019). “The FR-BDF Model and an Assessment of Monetary Policy Transmission to the French Economy.” *Banque de France Working Paper* No. 736. (Cited throughout as “wp736” / “FR-BDF”).

Dubois, É., Aldama, P., Bonnet, X., Lepetit, A., Lemoine, M., Mogliani, M., and Zhutova, A. (2026). “FR-BDF Reloaded: An Update of the FR-BDF Model after a Decade of Use.” *Banque de France Working Paper* No. 1044. (Cited throughout as “wp1044”. Provides the supply-block re-calibration method (§4.2.3), the full set of five PAC equation specifications (Eqs 16, 30, 35, 37, 46) that this paper replicates on AU data via the Phase L2 partial replication of §3.5, and the BI Table 3.5.13 calibration imported via Option 1 of §3.6 / §4.6.2.) Source PDF at [references/FR-BDF-update.pdf](#).

Adjemian, S., Brayton, F., and Zimic, S. (2024). “Solving semi-structural models with Dynare.” Dynare reference manual and workflow notes for the `pac_expectation` / `cherrypick` / `aggregate` pipeline used to build `dynare/au_pac.mod` from per-block auxiliary `.mod` files (§3.4).

Dees, S., Henrot, M., and Vourc’h, A. (2022). “The financial accelerator in the non-financial corporate sector: Evidence and modelling for FR-BDF.” *Banque de France Working Paper*. Referenced in §8 (Conclusion) as a candidate v3.0 extension for the cost-of-capital channel.

Bové, M., Boutin-Dufresne, F., and Penalver, A. (2020). “Modelling household debt service in a semi-structural framework.” *Banque de France*. Referenced in §8 as a candidate v3.0 extension for the mortgage-rate channel.

Tinsley, P.A. (1993). “Fitting Both Data and Theories: Polynomial Adjustment Costs and Error-Correction Decision Rules.” *FEDS Working Paper* No. 93-21, Board of Governors of the Federal Reserve System.

Zimic, S. (2023). “SemiStructDynareBasics: A reference implementation of FRB/US-style PAC models in Dynare.” ECB internal documentation and public GitLab repository. <https://gitlab.com/srecko/SemiStructDynareBasics>

Galí, J., Smets, F., and Wouters, R. (2011). “Unemployment in an Estimated New Keynesian Model.” *NBER Macroeconomics Annual*, 26(1), 329-360.

Iskrev, N. (2010). “Local Identification in DSGE Models.” *Journal of Monetary Economics*, 57(2), 189–202.

Komunjer, I. and Ng, S. (2011). “Dynamic Identification of Dynamic Stochastic General Equilibrium Models.” *Econometrica*, 79(6), 1995–2032.

Adolfson, M., Laséen, S., Lindé, J., and Villani, M. (2007). “Bayesian Estimation of an Open-Economy DSGE Model with Incomplete Pass-Through.” *Journal of International Economics*, 72(2), 481–511.

Australian monetary transmission and data

Reserve Bank of Australia (2024). “Statement on Monetary Policy.” Various issues, 2010–2024.

Bishop, J. and Tulip, P. (2017). “Anchoring Inflation Expectations in a Low Inflation Environment.” *RBA Research Discussion Paper* 2017-08.

Beechey, M.J. and Wright, J.H. (2009). “The High-Frequency Impact of News on Long-Term Yields and Forward Rates: Is It Real?” *Journal of Monetary Economics*, 56(4), 535–544.

Mulqueeney, J., Ballantyne, A., and Hambur, J. (2025). “Monetary Policy Transmission through the Lens of the RBA’s Models.” *RBA Bulletin*, April 2025, pp. 9–18.

Beckers, B. (2020). “Credit Spreads, Monetary Policy and the Price Puzzle.” *RBA Research Discussion Paper* 2020-01.

Cusbert, T. and Kendall, E. (2018). “Meet MARTIN, the RBA’s New Macroeconomic Model.” *RBA Bulletin*, March 2018.

Hambur, J., Cassidy, N., and Pugsley, B. (2024). “How Has Australian Monetary Policy Affected the Economy in Recent Years?” *RBA Research Discussion Paper*.

Cassidy, N. and Doyle, M.-A. (2018). “Wage Growth Puzzles and the Phillips Curve in Australia.” *RBA Bulletin*, March 2018. Source for the post-2015 RBA wage Phillips-curve discussion (§4.4.1, §8 finding 4).

Bishop, J. and Cassidy, N. (2017). “Insights into Low Wage Growth in Australia.” *RBA Bulletin*, March 2017. Companion to Cassidy and Doyle (2018) on AU wage non-responsiveness.

Bishop, J. and Cassidy, N. (2018). “Trends in Potential Output Growth.” *RBA Bulletin*. Referenced in §4.2.3 for the trend-GDP-deceleration discussion.

Cusbert, T. (2017). “Estimating the NAIRU and the Unemployment Gap.” *RBA Bulletin*, June 2017. Trend potential-output reference, §4.2.3.

Ballantyne, A., Sharma, T., and Taylor, T. (2024). “Assessing Full Employment in Australia.” *RBA Bulletin*.

Data sources

Australian Bureau of Statistics (2024). “Australian National Accounts: National Income, Expenditure and Product.” Cat. No. 5206.0. Quarterly, seasonally adjusted.

Australian Bureau of Statistics (2024). “Australian National Accounts: Capital Stock and Capital Consumption (Tables 47, 48).” Cat. No. 5204.0. Annual factor-share series used for CES supply-side estimation.

Australian Bureau of Statistics (2024). “Residential Property Price Indexes: Eight Capital Cities.” Cat. No. 6416.0. Used for the housing-price gap (`ph_gap`) from 2003Q3 onward; spliced back to 1959Q3 using a long-history Australian house-price series with chained growth rates anchored at the 2003Q3 overlap.

Australian Bureau of Statistics (2024). “Wage Price Index, Australia.” Cat. No. 6345.0. Quarterly index, SA Private+Public All industries — used as the wage-inflation observable in §6.4.

Bank for International Settlements (2024). “Effective Exchange Rate Indices.” BIS REER / NEER series for AUD.

Fair Work Commission (2024). “Annual Wage Review.” Final decisions 2010–2024. Institutional source for the CPI-indexation channel in AU wage-setting (affects approximately 15% of the workforce — award rates plus award-floor enterprise agreements).

Reserve Bank of Australia (2024). “Statistical Tables: F1 (interest rates), F11 (mortgage rates), G3 (commodity prices).” Quarterly time series, 1976–2024.

Australian Bureau of Statistics (2025). “Private New Capital Expenditure and Expected Expenditure, Australia.” Cat. No. 5625.0. Quarterly chain-volume measures, seasonally adjusted, by industry. Mining vs non-mining split obtained via Series IDs A3515875V (mining buildings + structures) and A124798315W (non-mining); used to prepare the Phase L3 mining-vs-non-mining BI hypothesis test referenced in §5.3 and Appendix G. Downloaded artefacts in `data/abs_rba/abs_5625_*.xlsx`; documentation in `data/abs_rba/abs_5625_README.md`.